



# Basic Concepts for NetApp ONTAP 9.7

## SAN Services, System Manager Edition

May 2023 | LD00378 Version 2.2.1

## TABLE OF CONTENTS

<b>1</b>	<b>Introduction.....</b>	<b>3</b>
1.1	Why NetApp ONTAP?.....	3
1.2	Lab Objectives.....	4
1.3	Prerequisites.....	4
1.4	Accessing the Command Line.....	4
<b>2</b>	<b>Lab Environment.....</b>	<b>7</b>
<b>3</b>	<b>Lab Activities.....</b>	<b>9</b>
3.1	Clusters.....	9
3.1.1	Cluster Concepts.....	9
3.1.2	Connect to the Cluster with ONTAP System Manager.....	9
3.1.3	Advanced Drive Partitioning.....	12
3.1.4	Create a New Aggregate on Each Cluster Node.....	18
3.1.5	Networks.....	24
3.2	Create Storage for iSCSI.....	30
3.2.1	Create a Storage Virtual Machine for iSCSI.....	31
3.2.2	Create, Map, and Mount a Windows LUN.....	43
3.2.3	Create, Map, and Mount a Linux LUN.....	91
<b>4</b>	<b>References.....</b>	<b>110</b>
<b>5</b>	<b>Version History.....</b>	<b>111</b>
	<b>Legal Notices.....</b>	<b>112</b>

# 1 Introduction

This lab introduces the fundamentals of NetApp ONTAP<sup>®</sup>. Using a pre-created 2-node cluster, you will provision storage for access through iSCSI, and configure Windows 2019 and Red Hat Enterprise Linux 7.5 clients to access storage on the cluster over iSCSI.

There is a similar lab available that demonstrates how to configure these same clients to access storage on the cluster using NAS protocols instead, specifically using CIFS and NFS. The NAS and SAN lab guides are designed so they can both be completed within the same running lab instance. If that option interests you, then for the best user experience we recommend that you complete the NAS guide first, followed by the SAN guide, as that ordering will most closely match the workflow examples you will see documented in the lab guides.

## 1.1 Why NetApp ONTAP?

One of the key ways to understand the benefits of ONTAP is to consider server virtualization. Before server virtualization, system administrators frequently deployed applications on dedicated servers in order to maximize application performance, and to avoid the instabilities often encountered when combining multiple applications on the same operating system instance. While this design approach was effective, it also had the following drawbacks:

- It did not scale well — adding new servers for every new application was expensive.
- It was inefficient — most servers are significantly under-utilized, and businesses are not extracting the full benefit of their hardware investment.
- It was inflexible — re-allocating standalone server resources for other purposes is time consuming, staff intensive, and highly disruptive.

Server virtualization directly addresses these limitations by decoupling the application instance from the underlying physical hardware. Multiple virtual servers can share a pool of physical hardware, allowing businesses to consolidate their server workloads to a smaller set of more effectively utilized physical servers. Additionally, the ability to transparently migrate running virtual machines across a pool of physical servers reduces the impact of downtime due to scheduled maintenance activities.

NetApp ONTAP brings these same benefits, and many others, to storage systems. As with server virtualization, ONTAP enables you to combine multiple physical storage controllers into a single logical cluster that can non-disruptively service multiple storage workload needs. With ONTAP you can:

- Combine different types and models of NetApp storage controllers (known as nodes) into a shared physical storage resource pool (referred to as a cluster).
- Support multiple data access protocols (CIFS, NFS, Fibre Channel, iSCSI, NVMeFC, etc.) concurrently on the same storage cluster.
- Consolidate various storage workloads to the cluster. Each workload can be assigned its own Storage Virtual Machine (SVM), which is essentially a dedicated virtual storage controller, and its own data volumes, LUNs, CIFS shares, and NFS exports.
- Support multi-tenancy with delegated administration of SVMs. Tenants can be different companies, business units, or even individual application owners, each with their own distinct administrators whose admin rights are limited to just the assigned SVM.
- Use Quality of Service (QoS) capabilities to manage resource utilization between storage workloads.
- Non-disruptively migrate live data volumes and client connections from one cluster node to another.
- Non-disruptively scale the cluster out by adding nodes. Nodes can likewise be non-disruptively removed from the cluster, meaning that you can non-disruptively scale a cluster up and down during hardware refresh cycles.
- Leverage multiple nodes in the cluster to simultaneously service a given SVM's storage workloads. This means that businesses can scale out their SVMs beyond the bounds of a single physical node in response to growing storage and performance requirements, all non-disruptively.
- Apply software and firmware updates, and configuration changes without downtime.


## 1.2 Lab Objectives

This lab explores fundamental concepts of ONTAP by walking you through the initial setup of a cluster, with a focus on support for SAN protocols (iSCSI).

Here is a summary of the exercises in this lab, along with their Estimated Completion Times (ECT):

- Clusters (Required, ECT = 20 minutes).
  - Explore a cluster.
  - View Advanced Drive Partitioning.
  - Create a data aggregate.
  - Create a Subnet.
- Storage Virtual Machines for iSCSI (Optional, ECT = 90 minutes including all optional subsections)
  - Create a Storage Virtual Machine.
  - For Windows (Optional, ECT = 40 minutes).
    - Create a Windows LUN and map the LUN to an igroup.
    - Configure a Windows client for iSCSI and MPIO and mount the LUN.
  - For Linux (Optional, ECT = 40 minutes).
    - Create a Linux LUN and map the LUN to an igroup.
    - Configure a Linux client for iSCSI and multipath and mount the LUN.

There are two versions of this lab guide, one uses System Manager (NetApp's graphical administration interface) to complete these tasks, while the other uses the ONTAP Command Line Interface (CLI). The lab end state produced by either method is exactly the same, so use the method you are the most comfortable with by selecting the appropriate lab guide.

 **Note:** The System Manager version of this lab demonstrates the “classic” System Manager interface that is standard in ONTAP 9.7 and earlier ONTAP versions. It does not discuss the “new experience” System Manager interface that is previewed in ONTAP 9.7 and that becomes standard in ONTAP 9.8.

## 1.3 Prerequisites

This lab introduces NetApp ONTAP, and makes no assumptions that the user has previous experience with ONTAP. The lab does assume some basic familiarity with storage system related concepts such as RAID, CIFS, NFS, LUNs, and DNS.

This lab includes steps for mapping shares and/or mounting LUNs on a Windows client. These steps assume that the lab user has a basic familiarity with Microsoft Windows.

This lab also includes steps for mounting NFS volumes and/or LUNs on a Linux client. All steps are performed from the Linux command line, and assumes a basic working knowledge of the Linux command line. A basic working knowledge of a text editor such as vi may be useful, but is not required.

## 1.4 Accessing the Command Line

PuTTY is the terminal emulation program used in the lab to log into Linux hosts and storage controllers in order to run command line commands.

1. The launch icon for the PuTTY application is pinned to the task bar on the Windows host Jumphost as shown in the following screen shot; just double-click on the icon to launch it.



**Tip:** If you already have a PuTTY session open and you want to start another (even to a different host), right-click the PuTTY icon and select **PuTTY** from the context menu.



**Figure 1-1:**

Once PuTTY launches you need to select the lab host you want to connect to. The following example shows how to connect to the ONTAP cluster named “cluster1”.

2. By default PuTTY should launch into the “Basic options for your PuTTY session” display as shown in the screen shot. If you accidentally navigate away from this view just click on the **Session** category item to return to this view.
3. Use the scrollbar in the “Saved Sessions” box to navigate down to the desired host and double-click it to open the connection. A terminal window will open and you will be prompted to log into the host. You can find the

correct username and password for the host in the [Lab Host Credentials](#) table found in the “Lab Environment” section of this guide.

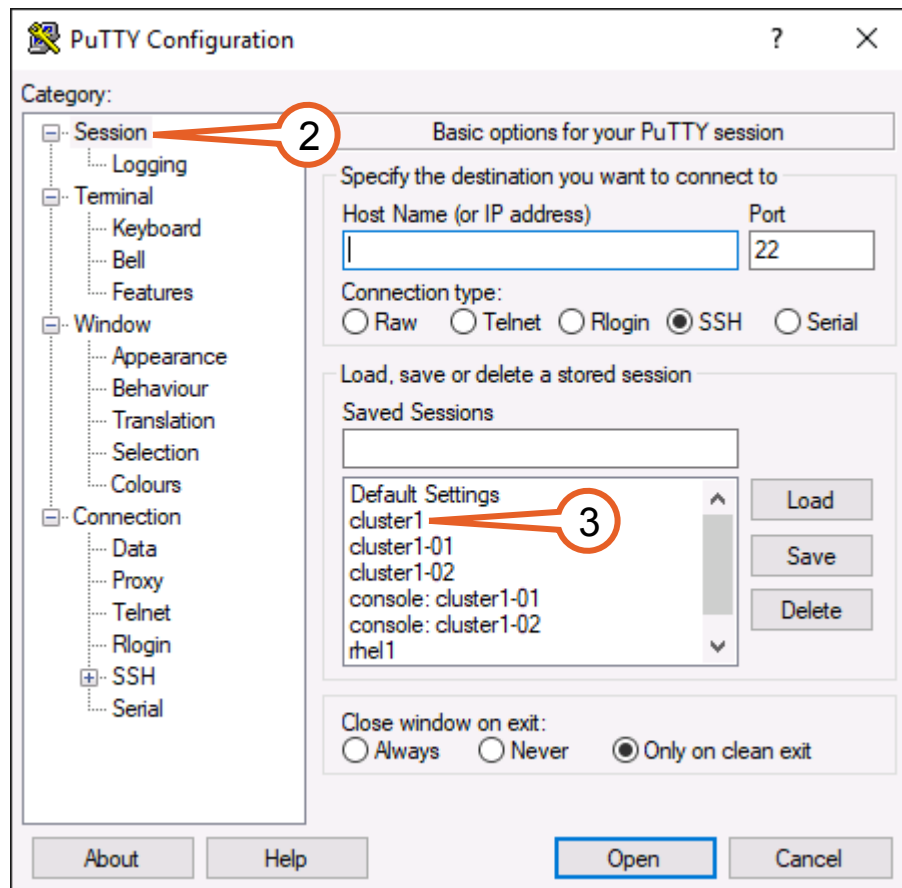


Figure 1-2:

If you are new to the ONTAP CLI, the length of the commands can seem a little intimidating. However, the commands are actually quite easy to use if you remember the following three tips:

- Make liberal use of the **Tab** key while entering commands, as the ONTAP command shell supports tab completion. If you hit the **Tab** key while entering a portion of a command word, the command shell will examine the context and try to complete the rest of the word for you. If there is insufficient context to make a single match, it will display a list of all the potential matches. Tab completion also usually works with command *argument values*, but there are some cases where there is simply not enough context for it to know what you want, in which case you will just need to type in the *argument value*.
- You can recall your previously entered commands by repeatedly pressing the **up-arrow** key, and you can then navigate up and down the list using the **up-arrow** and **down-arrow** keys. When you find a command you want to modify, you can use the **left-arrow**, **right-arrow**, and **Delete** keys to navigate around in a selected command to edit it.
- Entering a question mark character (?) causes the CLI to print contextual help information. You can use this character on a line by itself, or while entering a command.

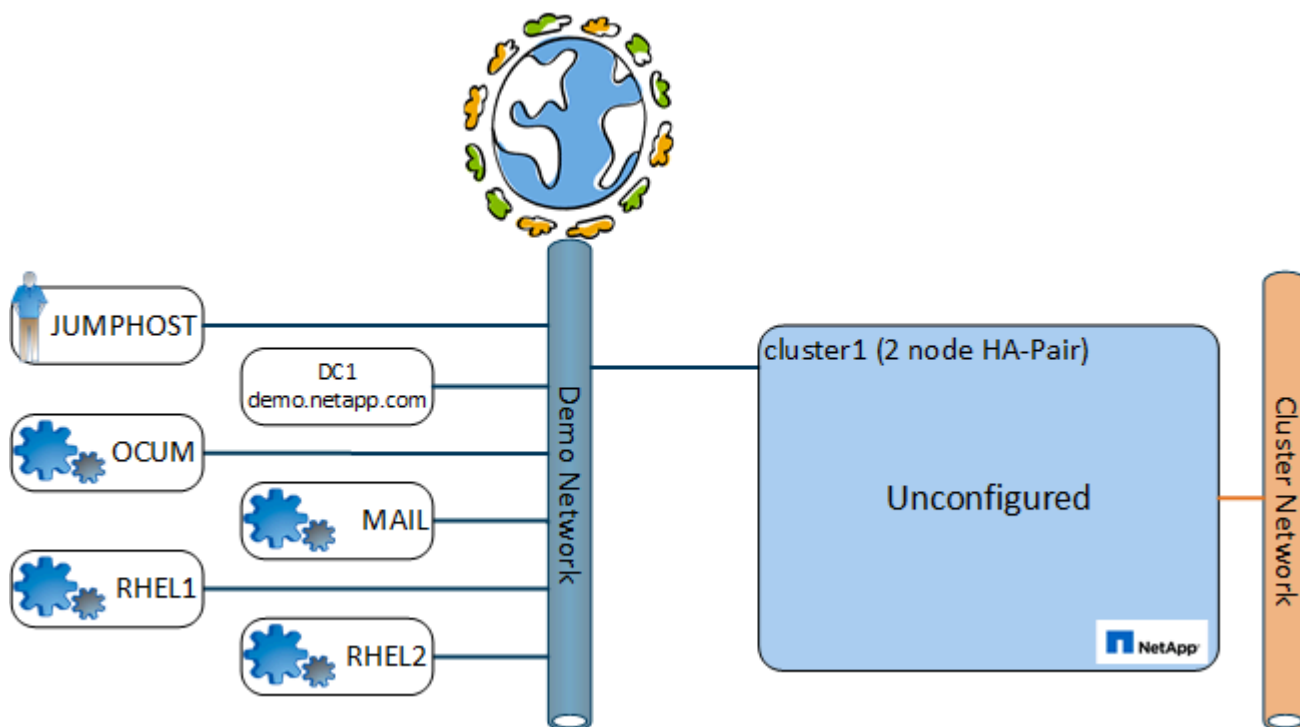
The ONTAP command line supports additional usability features that make the command line easier to use. If you are interested in learning more about this topic, refer to the “Hands-On Lab for Advanced Features of ONTAP” lab, which contains an entire section dedicated to this subject.



**Note:** In this lab, ONTAP automatically terminates ssh sessions to cluster1 after 2 hours of idle time. If you would prefer your PuTTY sessions to cluster1 to stay open indefinitely, issue the `system timeout modify 0` command at the ONTAP CLI.

## 2 Lab Environment

The following figure contains a diagram of the environment for this lab.



**Figure 2-1:**

All of the servers and storage controllers presented in this lab are virtual devices, and the networks that interconnect them are exclusive to your lab session. While we encourage you to follow the demonstration steps outlined in this lab guide, you are free to deviate from this guide and experiment with other ONTAP features that interest you. While the virtual storage controllers (vsims) used in this lab offer nearly all of the same functionality as physical storage controllers, they are not capable of providing the same performance as a physical controller, which is why these labs are not suitable for performance testing.

Table 1 provides a list of the servers and storage controller nodes in the lab, along with their IP address.

*Table 1: Table 1: Lab Host Credentials*

Hostname	Description	IP Address(es)	Username	Password
JUMPHOST	Windows 2019 Remote Access host	192.168.0.5	DEMO\Administrator	Netapp1!
RHEL1	Red Hat 7.5 x64 Linux host	192.168.0.61	root	Netapp1!
RHEL2	Red Hat 7.5 x64 Linux host	192.168.0.62	root	Netapp1!
DC1	Windows 2019 Active Directory Server	192.168.0.253	DEMO\Administrator	Netapp1!
cluster1	ONTAP 9.5 cluster	192.168.0.101	admin	Netapp1!
cluster1-01	ONTAP 9.5 cluster node	192.168.0.111	admin	Netapp1!
cluster1-02	ONTAP 9.5 cluster node	192.168.0.112	admin	Netapp1!
MAIL	E-Mail server/web email client	192.168.0.89	admin	Netapp1!

Hostname	Description	IP Address(es)	Username	Password
OCUM	OnCommand Unified Manager v9.5	192.168.0.71	DEMOAdministrator	Netapp1!

Table 2 lists the NetApp software that is pre-installed on the various hosts in this lab.

*Table 2: Table 2: Preinstalled NetApp Software*

Hostname	Description
JUMPHOST	Windows Unified Host Utility Kit v7.1.0, NetApp PowerShell Toolkit v4.7.0
RHEL1, RHEL2	Linux Unified Host Utilities Kit v7.1



## 3 Lab Activities

### 3.1 Clusters

Expected Completion Time: 20 Minutes

#### 3.1.1 Cluster Concepts

A cluster is a group of physical storage controllers, or nodes, that are joined together for the purpose of serving data to end users. The nodes in a cluster can pool their resources together so that the cluster can distribute its work across the member nodes. Communication and data transfer between member nodes (such as when a client accesses data on a node other than the one actually hosting the data) takes place over a high-speed cluster-interconnect network (10 to 100 Gb Ethernet) to which all the nodes are connected, while management and client data traffic passes over separate management and data networks configured on the member nodes.

Clusters typically consist of one, or more, NetApp storage controller High Availability (HA) pairs. Both controllers in an HA pair actively host and serve data, but they are also capable of taking over their partner's responsibilities in the event of a service disruption by virtue of their redundant cable paths to each other's disk storage. Having multiple HA pairs in a cluster allows the cluster to scale out to handle greater workloads, and to support non-disruptive migrations of volumes and client connections to other nodes in the cluster resource pool. This means that cluster expansion and technology refreshes can take place while the cluster remains fully online, and serving data.

Since clusters are almost always comprised of one or more HA pairs, a cluster almost always contains an even number of controller nodes. There is one exception to this rule, the "single node cluster", which is a special cluster configuration that supports small storage deployments using a single physical controller head. The primary difference between single node and standard clusters, besides the number of nodes, is that a single node cluster does not have a cluster network. Single node clusters can be converted into traditional multi-node clusters, at which point they become subject to all the standard cluster requirements like the need to utilize an even number of nodes consisting of HA pairs. This lab does not contain a single node cluster, so does not discuss them further.

ONTAP 9 clusters that only serve NFS and CIFS can scale up to a maximum of 24 nodes, although the node limit can be lower depending on the model of FAS controller and the specific ONTAP version in use. ONTAP 9 clusters that also host iSCSI and FC can scale up to a maximum of 12 nodes, but once again the limit may be lower depending on the FAS controller model.

This lab utilizes simulated NetApp storage controllers rather than physical FAS controllers. The simulated controller, also known as a "VSIM", is a virtual machine that simulates the functionality of a physical controller without the need for dedicated controller hardware. The vsim is not designed for performance testing, but does offer much of the same functionality as a physical FAS controller, including the ability to generate I/O to disks. This makes the vsim a powerful tool to explore and experiment with ONTAP product features. The vsim is limited when a feature requires a specific physical capability that the vsim does not support. For example, vsims do not support Fibre Channel connections, which is why this lab uses iSCSI to demonstrate block storage functionality.

This lab starts with a pre-created, minimally configured cluster. The pre-created cluster already includes ONTAP licenses, the cluster's basic network configuration, and a pair of pre-configured HA controllers. In this next section you will create the aggregates that are used by the SVMs that you will create in later sections of the lab. You will also take a look at the Advanced Drive Partitioning feature.

#### 3.1.2 Connect to the Cluster with ONTAP System Manager

ONTAP System Manager is NetApp's browser-based management tool to configure and manage NetApp storage systems and clusters. After performing ONTAP first-boot setup, point your web browser to the cluster management address to access System Manager.

On the Jumpstart, the Windows 2019 Server desktop you see when you first connect to the lab, open the web browser of your choice. This lab guide uses Chrome, but you can use Firefox if you prefer that browser instead. In the lab, both browsers already have System Manager set as the home page.

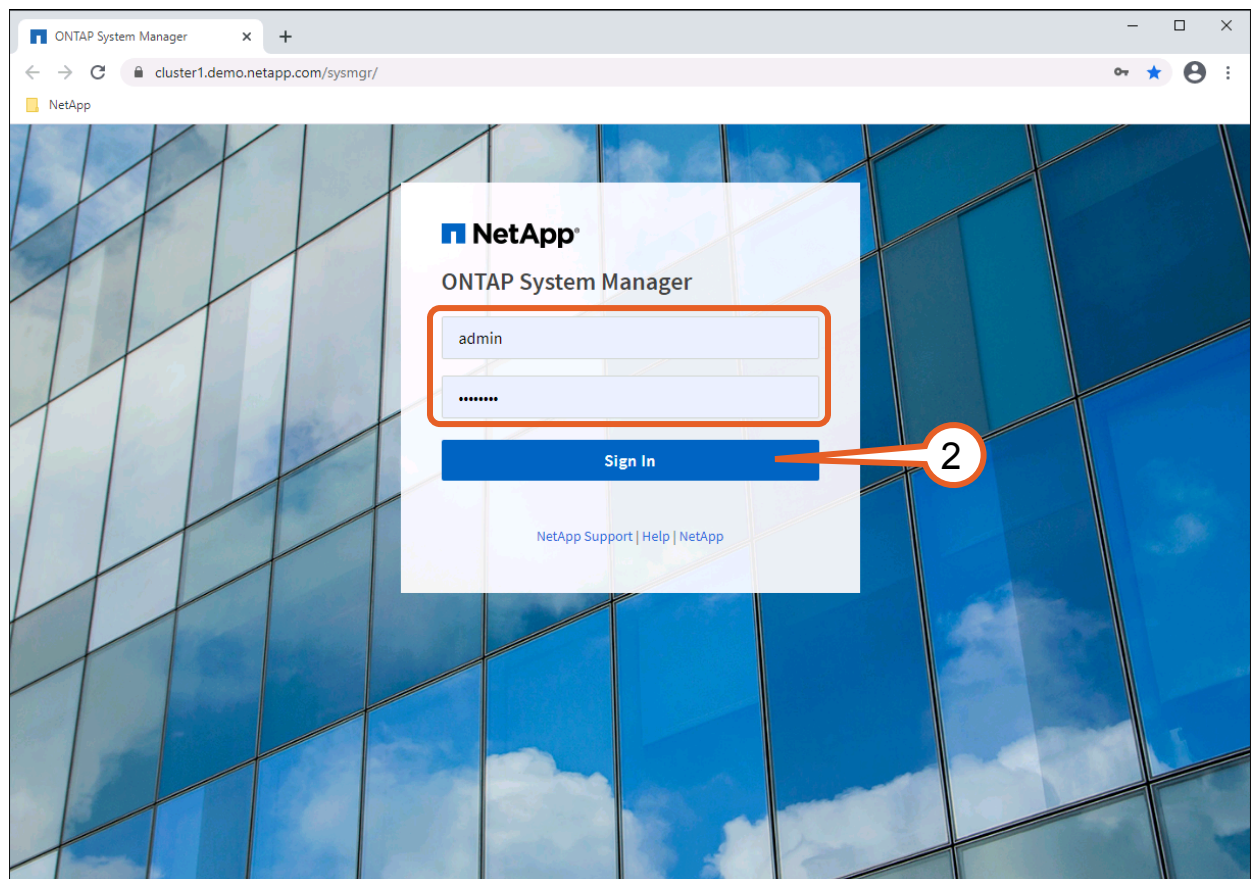
1. Launch **Chrome** to open System Manager.



**Figure 3-1:**

The ONTAP System Manager Login window opens.

2. Enter the User Name as **admin**, and the Password as **netapp1!**, and click **Sign In**.



**Figure 3-2:**

System Manager logs in to cluster1 and presents the “classic” System Manager interface that has been the default in recent releases up through ONTAP 9.7.

As ONTAP 9 has evolved, System Manager's user interface (UI) has undergone some fundamental redesign in order to improve usability. If you are unfamiliar with System Manager, or have used a prior version, here is a brief overview of the current ONTAP System Manager interface.

3. ONTAP 9.7 offers a preview of the new System Manager interface that will become the default in ONTAP 9.8, but since this new interface is not yet complete in 9.7 this lab guide only covers the “classic” System Manager interface.



**Tip:** If you wish to preview the new interface on your own, we recommend you do so only after completing the exercises in this lab guide so that any actions you might take in the new interface won't interfere with the workflows documented here.

Once you are ready to try out the new interface, you can click the green **Switch to the new experience** button in the upper left corner of the System Manager page. This link will re-direct you to the URL for the new System Manager interface, where you will be prompted to log in again using the same credentials you used for the classic interface. If you subsequently wish to return to the classic interface you can do so at any time by clicking the **(Return to classic version)** link located at the top of the “new experience” interface page next to the “ONTAP System Manager” title bar.

4. The **Dashboard** is the first page you see when you log into System Manager, and it displays summary information for the whole cluster. You can return to this view at any time by clicking **Dashboard** in the navigation pane.
5. **Applications & Tiers** provides access to templates you can use to provision storage configurations for a number of common applications, and also provides access to ONTAP's storage tiering features (i.e., FabricPool).
6. **Storage** is where you create and manage aggregates, Storage Virtual Machines, volumes, LUNs, and CIFS/NFS shares.
7. **Network** provides access to all the network interfaces for the cluster and the storage virtual machines.
8. **Protection** allows you to manage settings for SnapMirror and SnapVault relationships.
9. **Events & Jobs** provides access to ONTAP's system logs and alerts, and is where you interface with processes running in ONTAP's background.
10. **Configuration** is where you find cluster controls for activities like ONTAP upgrades, installing licenses, adding/removing nodes, High Availability (HA), cluster peering, and so on.
11. The **Plus Sign** button just above the far right upper right of the main pane list allows you to quickly launch wizards for many common administrative tasks such as creating an aggregate, creating a subnet, and creating/resizing/moving a volume.

12. The main pane of the dashboard provides real-time details on cluster alert status, health, performance, node status, and application/client metrics. Some of these widgets will not show meaningful information until after you configure the cluster in later exercises.

Take some time to expand and browse these tabs to familiarize yourself with their contents.

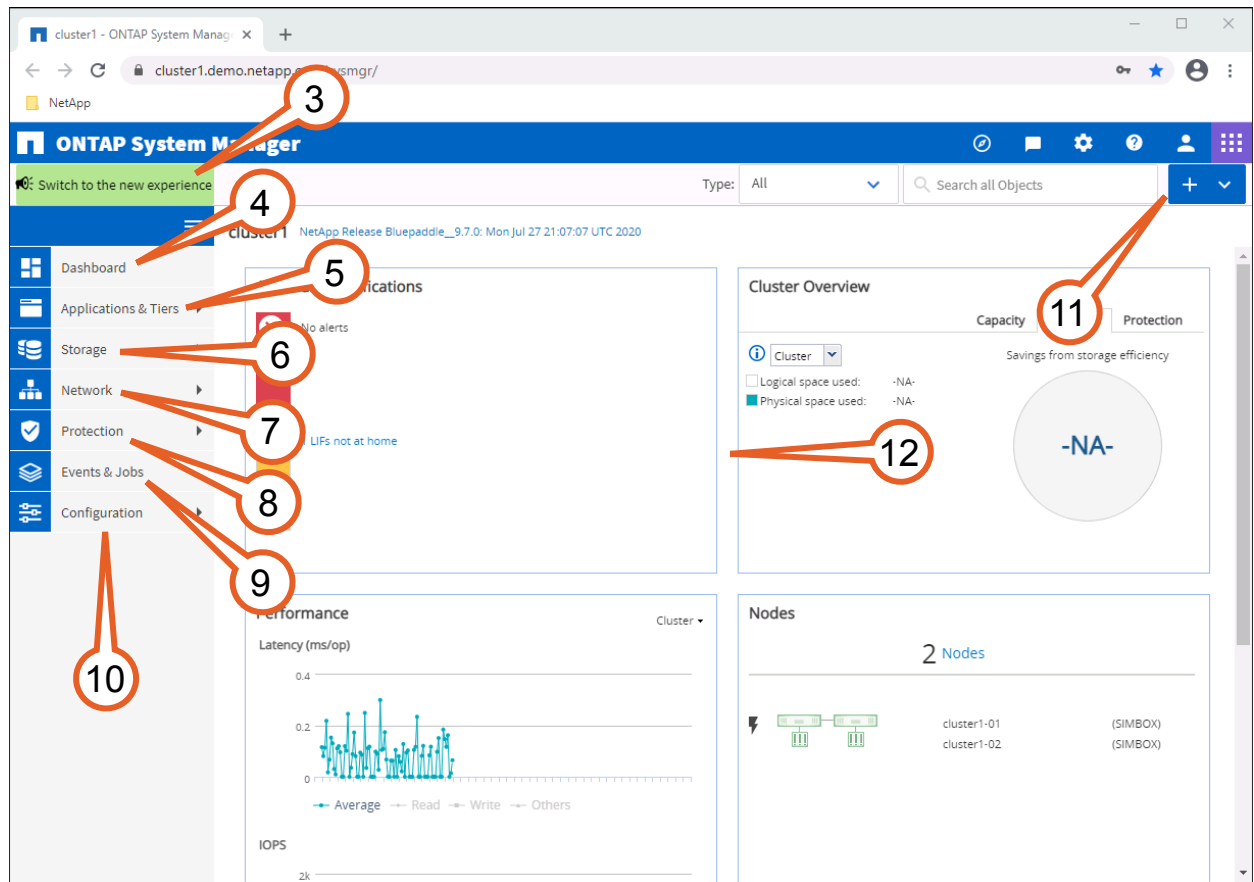


Figure 3-3:



**Note:** As you use System Manager in this lab, you may encounter situations where buttons at the bottom of a System Manager pane are beyond the viewing size of the window, and no scroll bar exists to allow you to scroll down to see them. If this occurs you have two options; either increase the size of the remote desktop browser window, or in the System Manager window, use the **tab** key to cycle through all the various fields and buttons, which eventually forces the window to scroll down to the non-visible items.

### 3.1.3 Advanced Drive Partitioning

#### 3.1.3.1 Advanced Drive Partitioning Concepts

Disks, whether Hard Disk Drives (HDD) or Solid State Disks (SSD), are the fundamental unit of physical storage in ONTAP, and are tied to a specific cluster node by virtue of their physical connectivity (i.e., cabling) to a given controller head.

ONTAP manages disks in groups called aggregates. An aggregate defines the RAID properties for a group of disks that are all physically attached to the same node. A given disk can only be a member of a single aggregate.

By default each cluster node has one aggregate known as the root aggregate, which is a group of the node's local disks that host the node's ONTAP operating system. A node's root aggregate is automatically created during ONTAP installation in a minimal RAID-DP configuration. This means it is initially comprised of 3 disks (1

data, 2 parity), and it has a name that begins the string `aggr0`. For example, in this lab the root aggregate of the node `cluster1-01` is named `aggr0_cluster1_01`, and the root aggregate of the node `cluster1-02` is named `aggr0_cluster1_02`.

On higher end FAS systems that have many disks, the requirement to dedicate 3 disks for each controller's root aggregate is not a burden, but for FAS systems that only have 24 or 12 disks, or for AFF systems with expensive high-capacity SSDs, this root aggregate disk overhead requirement significantly reduces the disks available for storing user data. To improve usable capacity, NetApp introduced Advanced Drive Partitioning (ADP), which divides the disks on nodes that have this feature enabled into two partitions; a small root partition, and a much larger data partition. ONTAP allocates the root partitions to the node root aggregate, and the data partitions for data aggregates. Each partition behaves like a virtual disk, so in terms of RAID, ONTAP treats these partitions just like physical disks when creating aggregates. The key benefit is that a much higher percentage of the node's overall disk capacity is now available to host user data.

For a FAS controller that uses Advanced Drive Partitioning, ONTAP dynamically determines the size of the root and data disk partitions at system installation time based on the quantity and size of the available disks assigned to each node.

ONTAP only supports HDD ADP for entry-level FAS controllers, and only for HDDs installed in their internal shelf on those models. ADP can only be enabled at system installation time, and there is no way to convert an existing system to use ADP other than to completely evacuate the affected HDDs and re-install ONTAP.

All-Flash FAS (AFF) supports ADP V2 that utilizes SSDs instead of HDDs, and divides the SSD into three partitions; a small root partition, plus two equal-sized data partitions. ADP V2 supports entry-level, mid-range, and high-end AFF platforms.

### 3.1.3.2 Advanced Drive Partitioning Exercise

In this lab activity, you use System Manager to determine if a cluster node is utilizing Advanced Drive Partitioning. System Manager provides a basic view into this information, but some details are only available through the ONTAP CLI.

1. In System Manager, navigate to **Storage > Aggregates & Disks > Disks** tab.
2. Click the **Summary** tab if it is not already selected.
3. The graphic in the middle of the page indicates that there are a total of 24 SSDs in the cluster. If you hover the mouse over the shaded portion of the bar chart you will see that the two cluster node root aggregates together consume about 70 GB of disk capacity. If you hover the mouse over the unshaded portion of the bar chart you will see that there is nearly 600 GB of spare disk capacity available for the creation of data aggregates.

4. In the “Spare Disks” pane at the bottom of the window, the spare disks are broken out by cluster node. There are 12 disks on each node, and each spare is 24.8 GB in size.

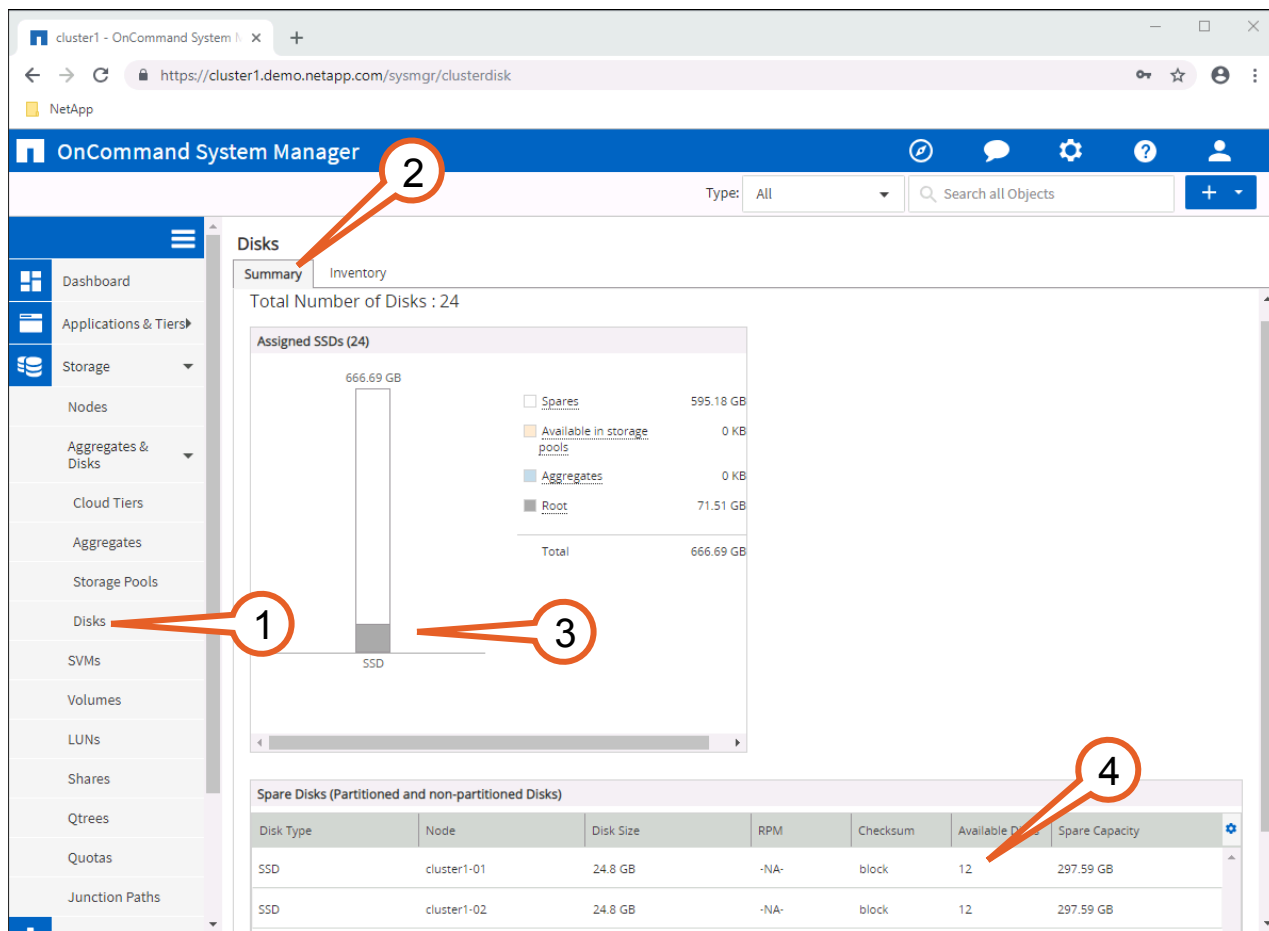


Figure 3-4:

- Click on the **Inventory** tab at the top of the right pane .

The screenshot shows the OnCommand System Manager web interface. The 'Disks' section is active, and the 'Inventory' tab is selected. A red circle with the number 5 highlights the 'Inventory' tab. The table below lists several disks with their respective attributes.

Name	Container Type	Home Owner	Current Owner	Type	RPM	Effective Size	Physical Space
VMw-1.36	Shared	cluster1-02	cluster1-02	SSD	-NA-	28.44 GB	28.44 GB
VMw-1.38	Shared	cluster1-01	cluster1-01	SSD	-NA-	28.44 GB	28.44 GB
VMw-1.26	Shared	cluster1-02	cluster1-02	SSD	-NA-	28.44 GB	28.44 GB
VMw-1.33	Shared	cluster1-02	cluster1-02	SSD	-NA-	28.44 GB	28.44 GB
VMw-1.40	Shared	cluster1-01	cluster1-01	SSD	-NA-	28.44 GB	28.44 GB
VMw-1.45	Shared	cluster1-01	cluster1-01	SSD	-NA-	28.44 GB	28.44 GB

Details:

- Aggregate: -NA-
- Vendor ID: VMware
- Zeroing (%): Zeroed
- Serial No: 6000c2997e7013063083edfb01012530
- Broken Details: -NA-

**Figure 3-5:**

System Manager's main pane now shows a list of the disks available across all the nodes in the cluster, which nodes own those disks, their effective and physical capacity, and so on. If you look at the "Container Type" column you see that the disks in your lab all show a value of "shared"; this value indicates that the physical disk is partitioned. For disks that are not partitioned you would typically see values like "spare", "data", "parity", and "dparity".

In this lab each cluster node has twelve 32 GB SSDs attached. Each SSD reports having 28.8 GB of capacity, which is the capacity remaining after formatting/overhead.

- Navigate to **Storage > Aggregates & Disks > Aggregates**
- Expand the entry for "aggr0\_cluster1\_01".

8. The “aggr0\_cluster1\_01” aggregate is comprised of ten disks. Click on the **10** link to see more details.

**Figure 3-6:**

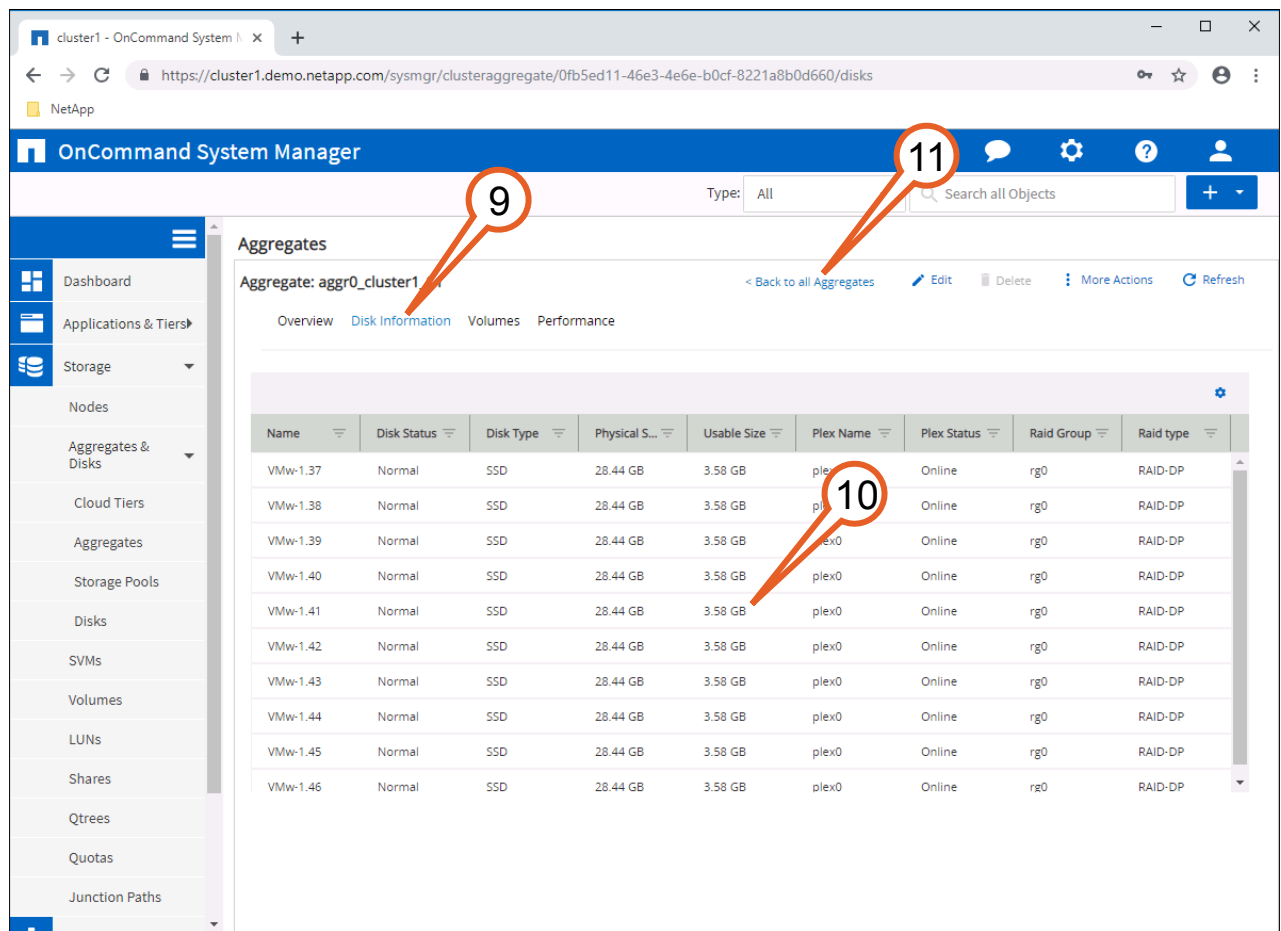
System Manager displays the “Aggregate: aggr0\_cluster1\_01” page.

9. Click the **Disk Information** tab.
10. The main pane displays a list of the disks that are in the aggregate. The “Usable Space” column shows that all the disks have a capacity of 3.58 GB, which represents the size of the root partition established on each disk by ADP.

**Note:** Recall that the cluster1-01 node has twelve local disks. The root partitions of ten of those disks form cluster1-01’s root aggregate (aggr0\_cluster1\_01). The other two disks are also partitioned, but are configured as spares as protection against potential drive failure.



11. Click **Back to all Aggregates** to return to the main aggregates window.



**Figure 3-7:**

System Manager displays the “Aggregates” page.

The ONTAP CLI includes a diagnostic level command that provides a more comprehensive single view of a system's partitioned disks.

12. If you do not already have a PuTTY session established to cluster1, launch PuTTY as described in the “[Accessing the Command Line](#)” section at the beginning of this guide, and connect to the host cluster1 with username **admin** and password **Netapp1!**.
13. The following commands shows all the disk partitions on the disks that belong to the node cluster1-01.

```
cluster1::> set -priv diag
```

```
Warning: These diagnostic commands are for use by NetApp personnel only.
Do you want to continue? {y|n}: y
```

```
cluster1::> disk partition show -owner-node-name cluster1-01
```

Partition	Usable Size	Container Type	Container Name	Owner
VMw-1.37.P1	12.41GB	spare	Pool0	cluster1-01
VMw-1.37.P2	12.41GB	spare	Pool0	cluster1-01
VMw-1.37.P3	3.58GB	aggregate	/aggr0_cluster1_01/plex0/rg0	cluster1-01
VMw-1.38.P1	12.41GB	spare	Pool0	cluster1-01
VMw-1.38.P2	12.41GB	spare	Pool0	cluster1-01
VMw-1.38.P3	3.58GB	aggregate	/aggr0_cluster1_01/plex0/rg0	cluster1-01
VMw-1.39.P1	12.41GB	spare	Pool0	cluster1-01
VMw-1.39.P2	12.41GB	spare	Pool0	cluster1-01
VMw-1.39.P3	3.58GB	aggregate	/aggr0_cluster1_01/plex0/rg0	cluster1-01

```

VMw-1.40.P1          12.41GB spare      Pool0      cluster1-01
VMw-1.40.P2          12.41GB spare      Pool0      cluster1-01
VMw-1.40.P3          3.58GB aggregate /aggr0_cluster1_01/plex0/rg0 cluster1-01
VMw-1.41.P1          12.41GB spare      Pool0      cluster1-01
VMw-1.41.P2          12.41GB spare      Pool0      cluster1-01
VMw-1.41.P3          3.58GB aggregate /aggr0_cluster1_01/plex0/rg0 cluster1-01
VMw-1.42.P1          12.41GB spare      Pool0      cluster1-01
VMw-1.42.P2          12.41GB spare      Pool0      cluster1-01
VMw-1.42.P3          3.58GB aggregate /aggr0_cluster1_01/plex0/rg0 cluster1-01
VMw-1.43.P1          12.41GB spare      Pool0      cluster1-01
VMw-1.43.P2          12.41GB spare      Pool0      cluster1-01
VMw-1.43.P3          3.58GB aggregate /aggr0_cluster1_01/plex0/rg0 cluster1-01
VMw-1.44.P1          12.41GB spare      Pool0      cluster1-01
VMw-1.44.P2          12.41GB spare      Pool0      cluster1-01
VMw-1.44.P3          3.58GB aggregate /aggr0_cluster1_01/plex0/rg0 cluster1-01
VMw-1.45.P1          12.41GB spare      Pool0      cluster1-01
VMw-1.45.P2          12.41GB spare      Pool0      cluster1-01
VMw-1.45.P3          3.58GB aggregate /aggr0_cluster1_01/plex0/rg0 cluster1-01
VMw-1.46.P1          12.41GB spare      Pool0      cluster1-01
VMw-1.46.P2          12.41GB spare      Pool0      cluster1-01
VMw-1.46.P3          3.58GB aggregate /aggr0_cluster1_01/plex0/rg0 cluster1-01
VMw-1.47.P1          12.41GB spare      Pool0      cluster1-01
VMw-1.47.P2          12.41GB spare      Pool0      cluster1-01
VMw-1.47.P3          3.58GB spare      Pool0      cluster1-01
VMw-1.48.P1          12.41GB spare      Pool0      cluster1-01
VMw-1.48.P2          12.41GB spare      Pool0      cluster1-01
VMw-1.48.P3          3.58GB spare      Pool0      cluster1-01
36 entries were displayed.

cluster1::*> set -priv admin

cluster1::>

```

## 3.1.4 Create a New Aggregate on Each Cluster Node

### 3.1.4.1 Aggregate Concepts

An aggregate is a group of disks on a node. A node can host multiple aggregates depending on the data sizing, performance, and isolation needs of the storage workloads that it will be hosting. When you create a Storage Virtual Machine (SVM) you assign it to use one or more specific aggregates to host the SVM's volumes. You can assign multiple SVMs to use the same aggregate, which offers greater flexibility in managing storage space, whereas dedicating an aggregate to just a single SVM provides greater workload isolation.

Aggregates utilize RAID to span storage across multiple disks. Although ONTAP supports several different RAID configuration options, NetApp recommends using RAID-DP, which is a double-parity RAID-6 implementation that prevents data loss when two disks fail in the RAID group.

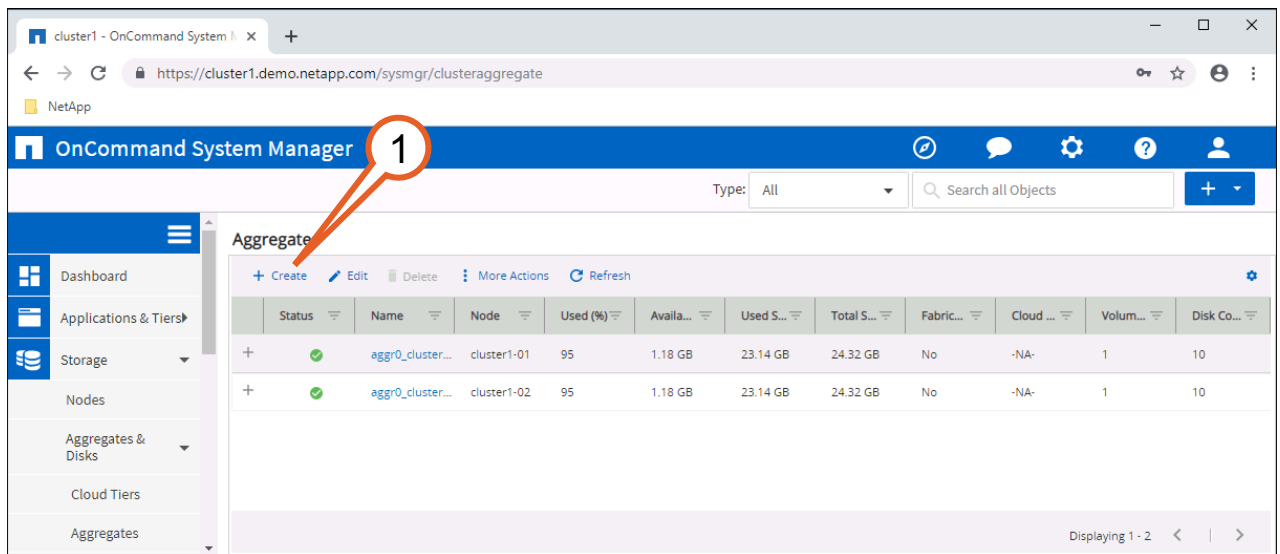
The only aggregates that exist on a newly created cluster are the node root aggregates. The root aggregate should not be used to host user data, so in this section you will create a new aggregate on each of the nodes in cluster1 so they can host any storage virtual machines, volumes, and LUNs that you create later in this lab.

### 3.1.4.2 Create Aggregate Exercise

In this lab activity, you create a data aggregate on each node in the cluster.

If you completed the preceding exercise then System Manager should still be displaying the contents of the Aggregates page. If you skipped that exercise then, starting from the Dashboard view, you can navigate to the Aggregates view by going to **Storage > Aggregates & Disks > Aggregates**.

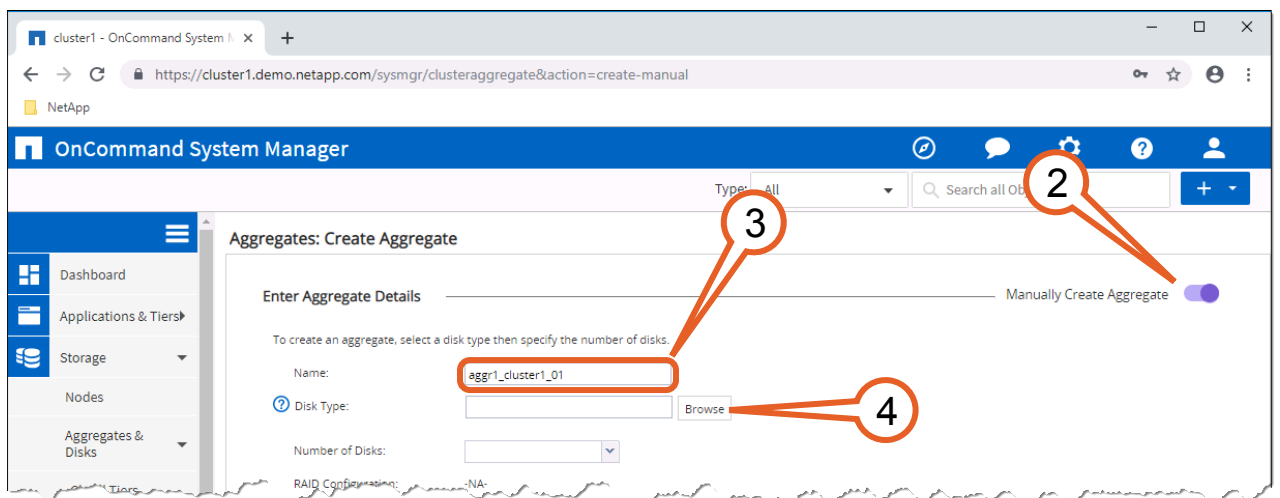
1. Click **Create** to launch the Aggregates: Create Aggregate Wizard.



**Figure 3-8:**

The “Aggregates: Create Aggregate” page opens.

2. ONTAP will offer to automatically create a maximally sized aggregate on each node. For the purpose of this lab a minimally sized aggregate is sufficient, but you must create it manually, so activate the **Manually Create Aggregate** toggle switch.
3. In the “Name” field enter `aggr1_cluster1_01`.
4. Click **Browse**.

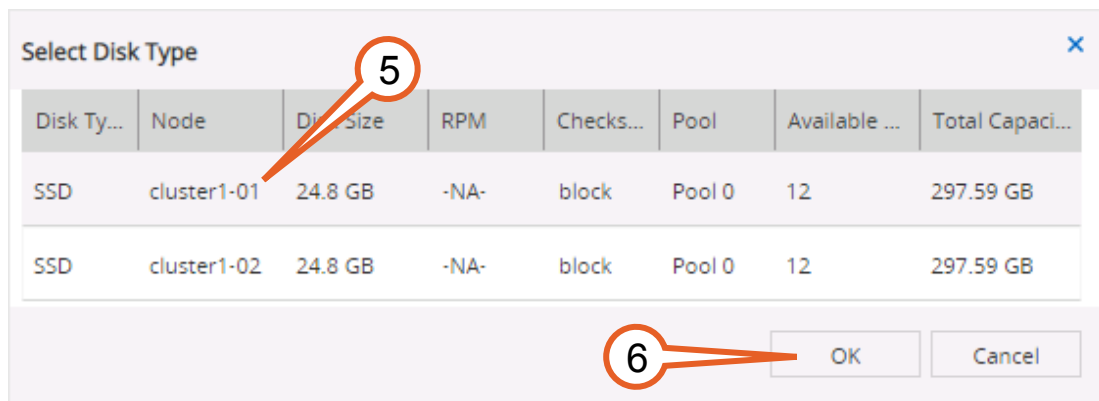


**Figure 3-9:**

The “Select Disk Type” dialog opens.

5. Select the Disk Type entry for the node **cluster1-01**.


6. Click **OK**.



**Figure 3-10:**

The “Select Disk Type” dialog closes, and focus returns to the “Aggregates: Create Aggregate” page.

7. The “Disk Type” should now display as “SSD”.
8. Set the “Number of Disks” to 5.
9. Note that the RAID configuration defaults to RAID-DP, which offers protection against double-disk failures. The default value is acceptable for this lab.
10. Click **Submit** to create the new aggregate, and to close the wizard.



**Figure 3-11:**

The “Create Aggregate” window closes, and focus returns to the “Aggregates” page. You will see a green box display at the top of the page stating “SUCCESS: The aggregate has been created”. You can

close the green box to dismiss the message. The newly created aggregate will also be visible in the Aggregates list.

11. Expand the entry for the aggregate **aggr1\_cluster1\_01**.
12. The expanded “Overview” section presents a summary of this aggregate's configuration. If you scroll down the page you can also view graphical metrics for the aggregates space allocation, storage efficiency, and performance characteristics (IOPS and Throughput). Since this is a newly created aggregate, there is little data for these metrics to report, but that will change when you start populating volumes and workloads on this aggregate.

The screenshot shows the OnCommand System Manager web interface. The left sidebar contains a navigation menu with options like Dashboard, Applications & Tiers, Storage, Nodes, Aggregates & Disks, Cloud Tiers, Aggregates, Storage Pools, Disks, SVMs, Volumes, LUNs, Shares, Qtrees, Quotas, and Junction Paths. The main content area displays a table of Aggregates. The table has columns for Status, Node, Used, Availa..., Used..., Total..., Fabric..., Cloud..., Volum..., and Disk Count. The aggregate 'aggr1\_cluster1\_01' is highlighted. Below the table, the 'OVERVIEW' section for 'aggr1\_cluster1\_01' is shown, displaying details such as Status (online), Node (cluster1-01), RAID Configuration (raid\_dp), RAID Status (normal), Root (No), Number of Volumes (0), and Allocated Disks (5). The bottom of the page shows tabs for SPACE ALLOCATION, SPACE SAVINGS, and PERFORMANCE.

Status	Node	Used (...)	Availa...	Used ...	Total ...	Fabric...	Cloud ...	Volum...	Disk Count	
+	aggr0_cluster1_01	cluster1-01	95	1.18 GB	23.14 GB	24.32 GB	No	-NA-	1	10
+	aggr0_cluster1_02	cluster1-02	95	1.18 GB	23.14 GB	24.32 GB	No	-NA-	1	10
-	aggr1_cluster1_01	cluster1-01	0	66.9 GB	172 KB	66.9 GB	No	-NA-	0	5

**OVERVIEW**

Status: online

Node: cluster1-01

RAID Configuration: raid\_dp (Data RAID group size of 24 d...)

RAID Status: normal

Root: No

Number of Volumes: 0 volumes

Allocated Disks: 5 Disks

SPACE ALLOCATION | SPACE SAVINGS | PERFORMANCE

**Figure 3-12:**

Now repeat the process to create a new aggregate on the node “cluster1-02”.

13. Click **Create** again.

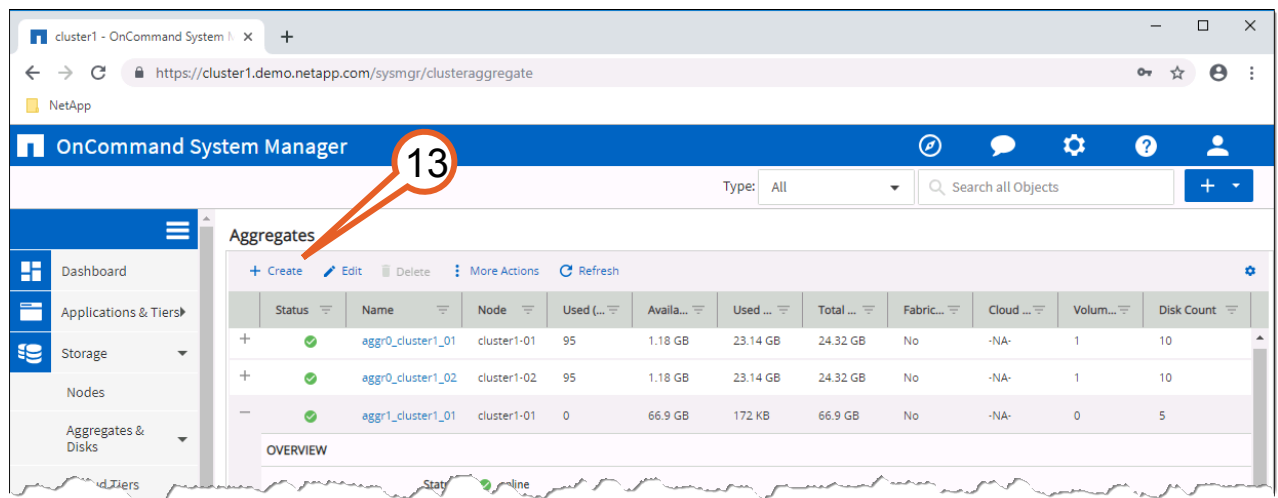


Figure 3-13:

The “Aggregates: Create Aggregate” page opens.

14. Activate the **Manually Create Aggregate** toggle switch.

15. Specify the fields on this page as follows:

- “Name”: **aggr1\_cluster1\_02**
- “Disk Type”: Use the **Browse** button to select the **SSD** disk type for the disks on node **cluster1-02**.
- “Number of Disks”: **5**

16. Click **Submit** to create the new aggregate.

The screenshot shows the 'Aggregates: Create Aggregate' page in the OnCommand System Manager. The page has a sidebar with navigation options like Dashboard, Applications & Tiers, Storage, Nodes, Aggregates & Disks, Cloud Tiers, Aggregates, Storage Pools, Disks, SVMs, Volumes, LUNs, Shares, Qtrees, Quotas, and Junction Paths. The main content area is titled 'Enter Aggregate Details' and includes a toggle for 'Manually Create Aggregate'. The form fields are as follows:

- Name:  (Callout 15)
- Disk Type:  (Callout 14 points to the search bar above)
- Number of Disks:  (Callout 15)
- RAID Configuration: RAID-DP: RAID group size of 24 disks
- New Usable Capacity: 66.9 GB (Estimated)
- SnapLock Type:

At the bottom right, there are 'Cancel' and 'Submit' buttons. Callout 16 points to the 'Submit' button.

**Figure 3-14:**

The “Aggregates: Create Aggregate” page closes, and focus returns to the “Aggregates” page where you will once again see a green box display at the top of the page stating “SUCCESS: The aggregate has been created”. Dismiss the message.

17. The new aggregate, “aggr1\_cluster1\_02” now appears in the cluster's aggregate list.

cluster1 - OnCommand System Manager

+

← → ↺

https://cluster1.demo.netapp.com/sysmgr/clusteraggregate&action=create-manual

🔒 ☆ 👤 ⋮

NetApp

OnCommand System Manager

🔍

🗨

⚙

?

👤

Type: All 🔽 🔍 Search all Objects + 🔽

☰

Dashboard

Applications & Tiers

Storage

Nodes

Aggregates & Disks

Cloud Tiers

Aggregates

Aggregates

+ Create ✎ Edit 🗑 Delete ⋮ More Actions ↺ Refresh ⚙

	Status	Name	Node	Used	Availa	Used	Total	Fabric	Cloud	Volum	Disk Count
+	✔	aggr0_cluster1_01	cluster1-01	17	1.18 GB	23.14 GB	24.32 GB	No	-NA-	1	10
+	✔	aggr0_cluster1_02	cluster1-01	95	1.18 GB	23.14 GB	24.32 GB	No	-NA-	1	10
+	✔	aggr1_cluster1_01	cluster1-01	0	66.9 GB	172 KB	66.9 GB	No	-NA-	0	5
+	✔	aggr1_cluster1_02	cluster1-02	0	66.9 GB	172 KB	66.9 GB	No	-NA-	0	5

Displaying 1 - 4 < | >

**Figure 3-15:**

## 3.1.5 Networks

### 3.1.5.1 Network Concepts

This section discusses the network components that ONTAP provides to manage your cluster.

Ports are the physical Ethernet and Fibre Channel connections on each node, the interface groups (ifgrps) you can create to aggregate those connections, and the VLANs you can use to subdivide them.

A logical interface (LIF) is essentially an IP address that is associated with a port, and has a number of associated characteristics such as an assigned home node, an assigned physical home port, a list of physical ports it can fail over to, an assigned SVM, a role, a routing group, and so on. A given LIF can only be assigned to a single SVM, and since LIFs are mapped to physical network ports on cluster nodes this means that an SVM runs, in part, on all nodes that are hosting its LIFs.

Routing tables in ONTAP are defined for each Storage Virtual Machine. Since each SVM has its own routing table, changes to one SVM's routing table do not have impact on any other SVM's routing table.

IPspaces allow you to configure an ONTAP cluster to logically separate one IP network from another, even if those two networks are using the same IP address range. IPspaces are a multi-tenancy feature that allow storage service providers to share a cluster between different companies while still separating storage traffic for privacy and security. Every cluster includes a default IPspace to which ONTAP automatically assigns new SVMs, and that default IPspace is usually sufficient for NetApp customers who deploy a cluster within a single company or organization that uses a non-conflicting IP address range.

Broadcast Domains are collections of ports that all have access to the same layer 2 networks, both physical and virtual (i.e., VLANs). Every IPspace has its own set of Broadcast Domains, and ONTAP provides a default broadcast domain to go along with the default IPspace. Broadcast domains are used by ONTAP to determine what ports an SVM can use for its LIFs.

Subnets in ONTAP are a convenience feature, intended to make LIF creation and management easier for ONTAP administrators. A subnet is a pool of IP addresses that you can specify by name when creating a LIF. ONTAP will automatically assign an available IP address from the pool to the LIF, along with a subnet mask and a gateway. A subnet is scoped to a specific broadcast domain, so all the subnet's addresses belong to the same layer 3 network. ONTAP manages the pool automatically as you create or delete LIFs, and if you manually configure a LIF with an address from the pool, it will detect that the address is in use and mark it as such in the pool.

DNS Zones allow an SVM to manage DNS name resolution for its own LIFs. Since multiple LIFs can share the same DNS name, this allows the SVM to load balance traffic by IP address across a set of LIFs. To use DNS Zones you must configure your DNS server to delegate DNS authority for the subdomain to the SVM.

### 3.1.5.2 Create Subnet Exercise

In this lab activity, you use System Manager to create a Subnet that you will leverage in later exercises to assign IP addresses while provisioning SVMs and LIFs. You will not create IPspaces or Broadcast Domains, as the system defaults are sufficient for this lab.

1. In System Manager, navigate to **Network > Broadcast Domains**.



2. In the “Broadcast Domains” pane, select the **Default** broadcast domain.

The screenshot shows the OnCommand System Manager interface. On the left, the 'Network' menu is expanded, and 'Broadcast Domains' is selected, indicated by a red circle with the number 1. The main pane displays the 'Broadcast Domains' table, where the 'Default' row is selected, indicated by a red circle with the number 2. Below the table, the 'Port Details' section shows a table of ports (e0c, e0e, e0f, e0g) on both cluster nodes (cluster1-01 and cluster1-02), all marked with green checkmarks.

Broadcast Domain	IPspace	Combined Port Updat...	VIP Broadcast Domain
Cluster	Cluster	complete	No
Default	Default	complete	No

Port Name	cluster1-01	cluster1-02
e0c	✓	✓
e0e	✓	✓
e0f	✓	✓
e0g	✓	✓

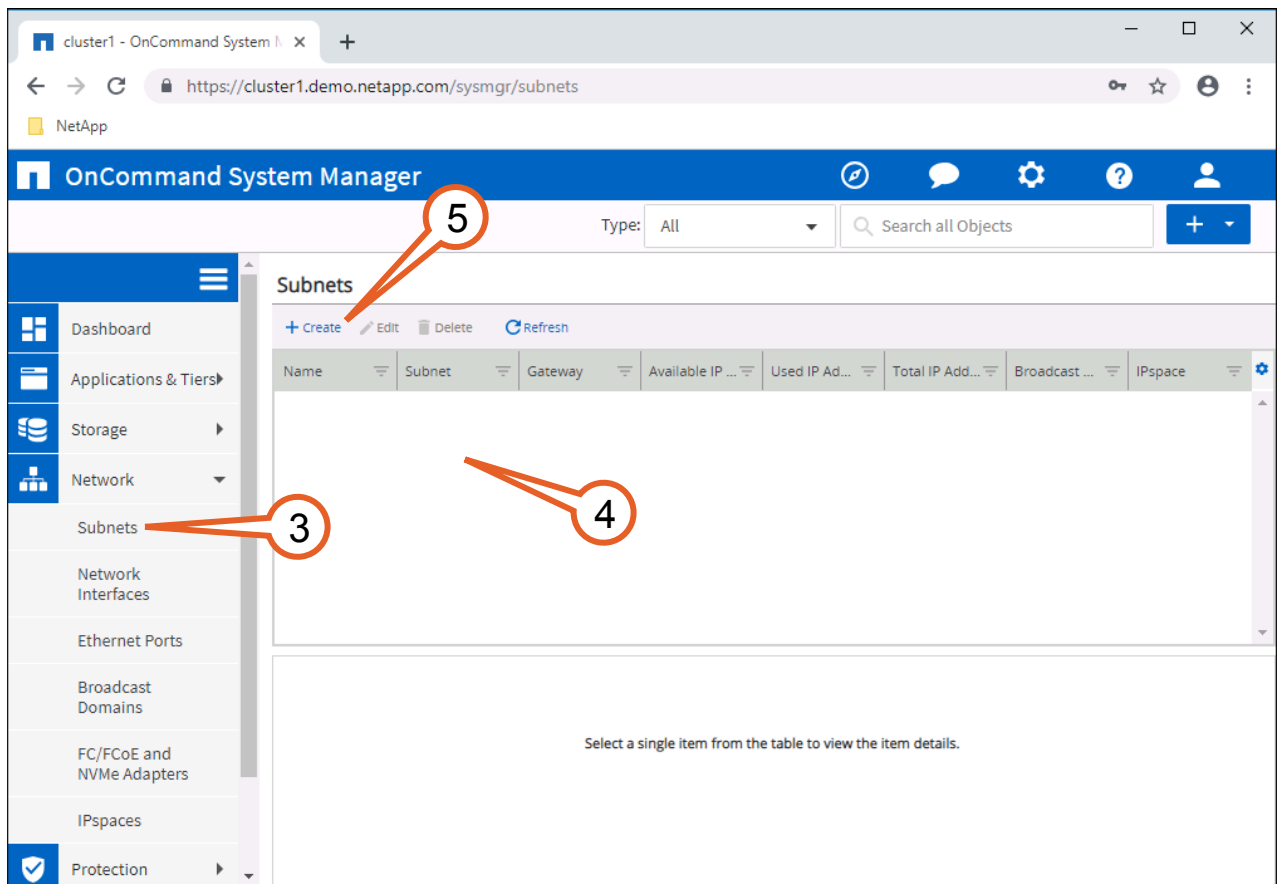
**Figure 3-16:**

Review the “Port Details” section at the bottom of the Network pane, and note that the e0c – e0g ports on both cluster nodes are all part of this broadcast domain. These are the network ports that you will use in this lab.

Now create a new Subnet for this lab.

3. Navigate to **Network > Subnets**.
4. There are currently no subnets listed in the “Subnets” pane. Unlike Broadcast Domains and IPspaces, ONTAP does not provide a default Subnet.

5. Click **+ Create**.



**Figure 3-17:**

The “Create Subnet” dialog opens.

6. Set the fields in this window as follows:

- “Name”: **Demo**
- “Subnet IP/Subnet mask”: **192.168.0.0/24**
- “IP Addresses”: The values you enter in the “IP address” field depend on which lab guide(s) you intend to complete in this lab instance.



**Important:** It is important that you choose the right values here so that the values in your lab will correctly match up with the values used in this lab guide.

- If you plan to complete only the NAS lab guide, or both the NAS and SAN lab guides, then enter **192.168.0.131-192.168.0.139**.
- If you plan to complete just the SAN lab guide then enter **192.168.0.133-192.168.0.139**.
- “Gateway”: **192.168.0.1**

7. Click **Browse**.

Create Subnet

You can create a subnet to provide a logical subdivision of an IP network to pre-allocate the IP addresses and divide space efficiently.

Name: Demo

Subnet IP/Subnet mask: 192.168.0.0/24

IP Addresses: 192.168.0.131-192.168.0.139 (Optional)

Gateway: 192.168.0.1 (Optional)

Broadcast Domain: Browse... (Optional)

Show ports on this domain

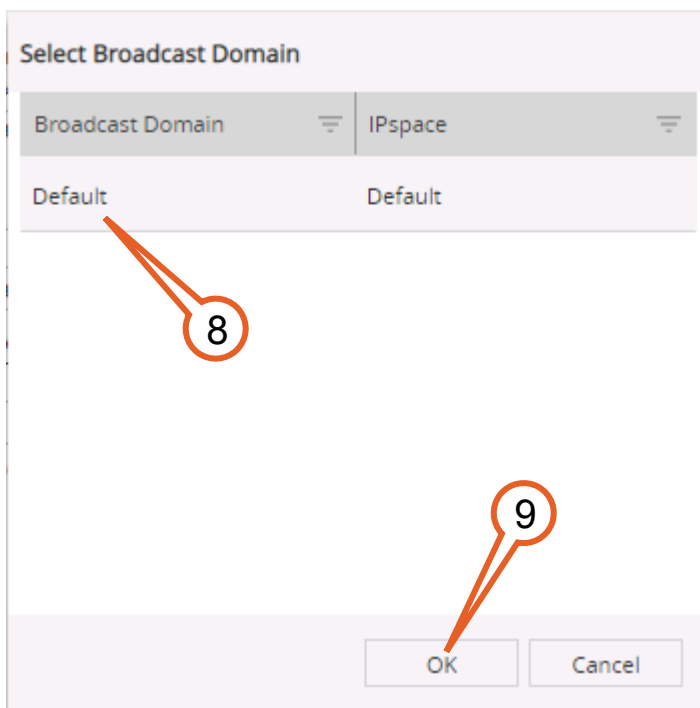
Create Cancel

**Figure 3-18:**

The “Select Broadcast Domain” dialog opens.

8. Select the **Default** entry from the list.

9. Click **OK**.



**Figure 3-19:**

The “Select Broadcast Domain” dialog closes, and focus returns to the “Create Subnet” dialog.

10. The values in your “Create Subnet” dialog should now match those shown in the following screenshot. The only possible exception would be the value you entered in the “IP Addresses” field, since that value depends on the range you entered to support the lab guide(s) (NAS vs. SAN) that you are using.
11. If it is not already displayed, click the **Show ports on this domain** link under the Broadcast Domain textbox to see the list of ports that this broadcast domain includes.

12. Click **Create**.

**Create Subnet**

You can create a subnet to provide a logical subdivision of an IP network to pre-allocate the IP addresses and divide space efficiently.

Name:

? Subnet IP/Subnet mask:

? IP Addresses:  (Optional)  
IP addresses can be a range or a list of comma separated entries or a mix of both.  
For example: IPv4 - 192.168.1.2,192.168.1.5-192.168.1.9 etc. or IPv6 - f6::c3,f6::c5-f6::c9 etc.

Gateway:  (Optional)  
! If you specify the gateway, a default route to that gateway is added to the associated SVM when a LIF is created using this subnet.

? Broadcast Domain:

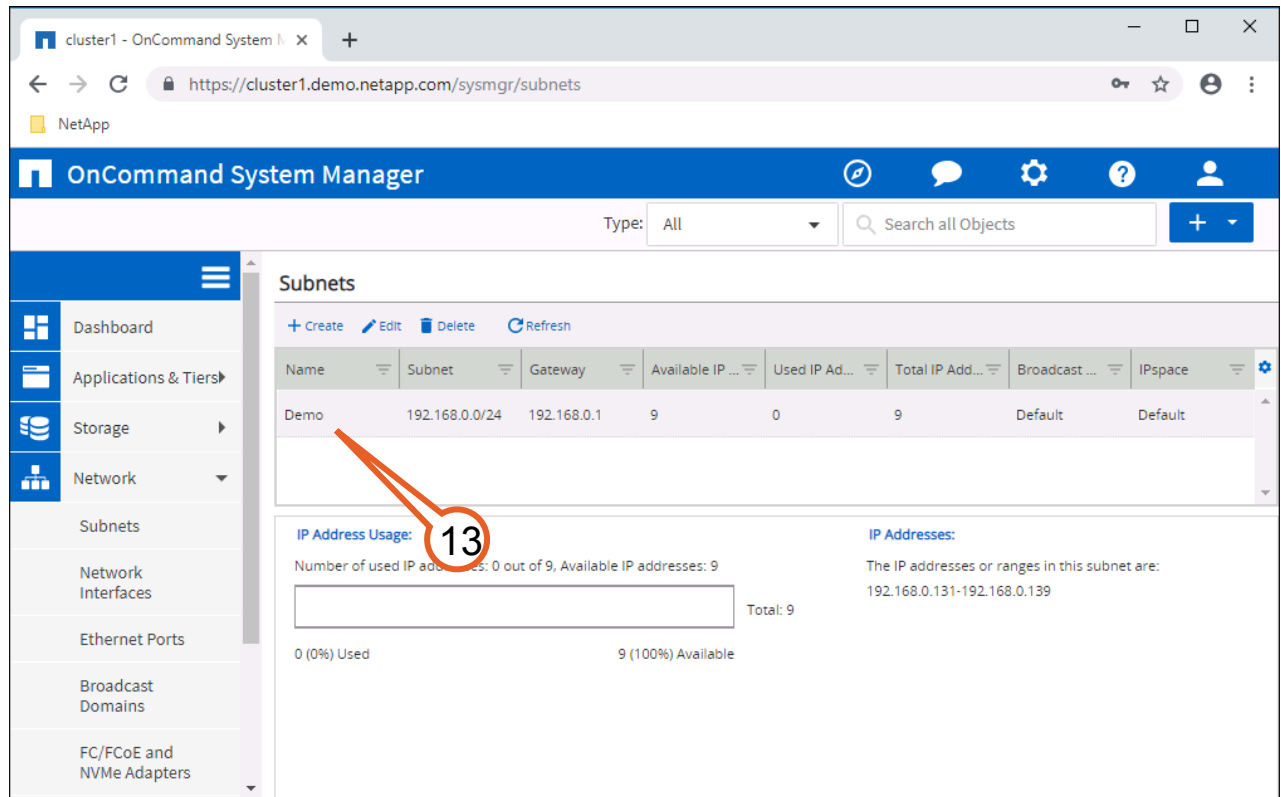
▼ Show ports on this domain

Port Name	cluster1-01	cluster1-02
e0c	↑	↑
e0d	↑	↑
e0e	↑	↑

**Figure 3-20:**

The “Create Subnet” dialog closes, and focus returns to the “Subnets” tab in System Manager.

13. Notice that the main pane of the “Subnets” tab now includes an entry for your newly created subnet, and that the lower portion of the pane includes metrics that track the consumption of the IP addresses that belong to this subnet.



**Figure 3-21:**

Feel free to explore the contents of the other tabs available in the navigation pane under the “Network” section. Here is a brief summary of the information available on those tabs.

- The “Network Interfaces” tab displays a list of all of the LIFs on your cluster.
- The “Ethernet Ports” tab displays the physical NICs on your controller. These include the NICs that you previously saw listed as belonging to the default broadcast domain. The other NICs you see listed on the Ethernet Ports tab include the node's cluster network NICs.
- The “FC/FCoE and NVMe Adapters” tab lists all the WWPNs for all the controller’s NICs in the event they will be used for iSCSI or FCoE connections. The NetApp controllers you are using in this lab do not include FC adapters, and this lab does not make use of FCoE.
- The “IPspaces” tab lists the network spaces used to segregate networks for tenants in multi-tenant systems. This allows different tenants to use the same IP address ranges without mixing traffic. A multi-node cluster contains two IPspaces by default; the “Cluster” IPspace segregates IP addresses for the cluster NICs, while the “Default” IPspace is used for everything else.

## 3.2 Create Storage for iSCSI

Expected Completion Time: 50 Minutes

The 50 minute time estimate assumes you complete only one of the Windows or Linux “Create, Map, and Mount a LUN” sections. You are welcome to complete both of those sections, but plan it taking approximately 90 minutes to complete the entire exercise.

In this section you create a new SVM, and configure it for iSCSI. The configuration steps for iSCSI and FC are similar, so the information provided here is very applicable for FC deployments too. After you create the SVM and configure it for iSCSI, you will create a LUN for Windows and/or a LUN for Linux, and then mount the LUN(s) on their respective hosts.

In this lab the cluster contains two nodes connected to a single storage network. You will still configure a total of 4 SAN LIFs, because it is common to see implementations with 2 paths per cluster node for redundancy.

## 3.2.1 Create a Storage Virtual Machine for iSCSI

### 3.2.1.1 General SVM Concepts

Storage Virtual Machines (SVMs), previously known as Vservers, are the logical storage servers within a cluster that serve data out to storage clients. A single cluster can host hundreds of SVMs, with each SVM managing its own set of volumes (FlexVols), Logical Network Interfaces (LIFs), storage access protocols (e.g., NFS/CIFS/iSCSI/FC/FCoE), and for NAS clients, its own namespace.

A single SVM can host any combination of the supported storage protocols. You explicitly configure which storage protocols you want a given SVM to support at the time you create the SVM. You can later add or remove protocols as desired.

The ability to support many SVMs in a single cluster is a key feature in ONTAP, and NetApp encourages customers to embrace this feature in order to take full advantage of a cluster's capabilities. A cluster that utilizes multiple SVMs will scale significantly better than a cluster that only has a single SVM, and provides administrators with much greater flexibility when managing workloads and access controls. The number of ONTAP storage controller nodes in a cluster determines the total number of SVMs that can run in the cluster. Each storage controller node can host a maximum of 125 SVMs, so you can calculate the cluster's effective SVM limit by multiplying the number of nodes by 125. As cluster size gets larger, the maximum number of SVMs a cluster can support approaches 1000 SVMs.

An SVM's assigned aggregates and LIFs determine which cluster nodes handle processing for that SVM. As you saw earlier, an aggregate is directly connected to the specific node hosting its disks, which means that an SVM runs in part on any nodes whose aggregates are hosting volumes for the SVM. An SVM also has a direct relationship to any nodes that are hosting its LIFs. LIFs are essentially an IP address with a number of associated characteristics, such as an assigned home node, an assigned physical home port, a list of physical ports it can fail over to, an assigned SVM, a role, a routing group, and so on. You can only assign a given LIF to a single SVM, and since LIFs map to physical network ports on cluster nodes, this means that an SVM runs in part on all the nodes that are hosting its LIFs.

The most efficient client access path to a volume's data is through a LIF that is mapped to a physical network port located on the same node as the aggregate that hosts the volume's storage. However, clients can also access volume data through LIFs bound to physical network ports on other nodes in the cluster; in these cases ONTAP uses the high speed cluster network to bridge communication between the node hosting the LIF and the node hosting the volume.

When you configure an SVM with multiple data LIFs, clients can use any of those LIFs to access volumes hosted by the SVM. The specific LIF IP address (and by extension, which LIF) a client will choose is a function of the storage protocol and of name resolution.

### 3.2.1.2 SAN-Specific SVM Concepts

NetApp supports configuring an SVM to serve data over both SAN and NAS protocols, but it is common to see customers use separate SVMs for each in order to separate administrative responsibilities, or for architectural and operational clarity. For example, SAN protocols do not support LIF failover, so you cannot use NAS LIFs to support SAN protocols. You must instead create dedicated LIFs just for SAN. Implementing separate SVMs for SAN and NAS can, in this example, simplify the operational complexity of each SVM's configuration, making each easier to understand and manage. But ultimately, whether to mix or separate is a customer decision and not a NetApp recommendation.

Since SAN LIFs do not support migration to different nodes, an SVM must have dedicated SAN LIFs on every node that you want to service SAN requests, and you must utilize MPIO and ALUA to manage the controller's

available paths to the LUNs. In the event of a path disruption, MPIO and ALUA will be compensated by re-routing the LUN communication over an alternate controller path (i.e., over a different SAN LIF).

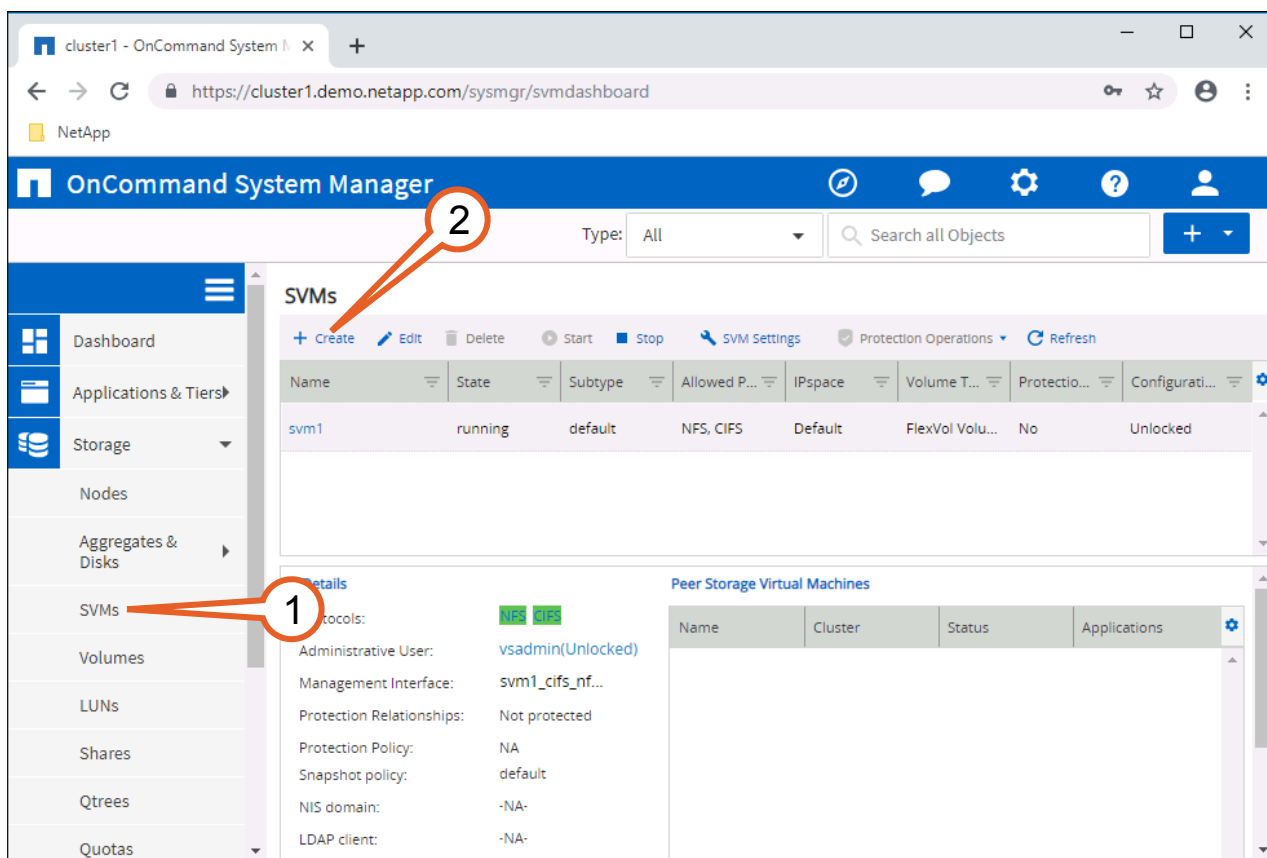
NetApp best practice is to configure at least one SAN LIF per storage fabric/network on each node in the cluster so that all nodes can provide a path to the LUNs. In large clusters where this would result in the presentation of a large number of paths for a given LUN we recommend that you use portsets to limit the LUN to seeing no more than 8 LIFs.

ONTAP also supports Selective LUN Mapping (SLM) to provide further assistance in managing fabric paths. SLM limits LUN path access to just the node that owns the LUN and its HA partner. ONTAP automatically applies SLM to all new LUN map operations.

### 3.2.1.3 Create an SVM for iSCSI Exercise

In this exercise you will create a new SVM named “svmluns” on the cluster. You will then configure it for iSCSI, and create four data LIFs to support LUN access to the SVM (two on each cluster node).

1. In the left pane of System Manager, navigate to **Storage > SVMs**.
2. Click **Create** to launch the Storage Virtual Machine Setup wizard.



**Figure 3-22:**

The “Storage Virtual machine (SVM) Setup” dialog opens.



3. Set the fields in this window as follows:

- “SVM Name”: **svmluns**
- “Data Protocols”: check the **iSCSI** check box.



**Tip:** The list of available Data Protocols is dependant upon what protocols are licensed on your cluster. If a given protocol is not listed, it is because you are not licensed for it. (In this lab the cluster is fully licensed for all features.)

- “Security Style”: **UNIX**.
- “Root Aggregate”: **aggr1\_cluster1\_01**.



**Tip:** If you completed the NAS version of this lab, you will note that this is the same aggregate you used to hold the volumes for svm1. Multiple SVMs can share the same aggregate.

The default values for IPspace, Default Language, and Security Style are already populated for you by the wizard, as is the DNS configuration.

4. Click **Submit & Continue**.

Storage Virtual Machine (SVM) Setup

1 Enter SVM basic details

### SVM Details

Specify a unique name and the data protocols for the SVM

SVM Name:

IPspace:

Data Protocols: ☐ CIFS ☐ NFS ☒ iSCSI ☐ FC/FCoE

Default Language:

The language of the SVM specifies the default language encoding setting for the SVM and its volumes. Using a setting that incorporates UTF-8 character encoding is recommended.

Security Style:

Root Aggregate:

### DNS Configuration

Specify the DNS domain and name servers. DNS details are required to configure the CIFS protocol.

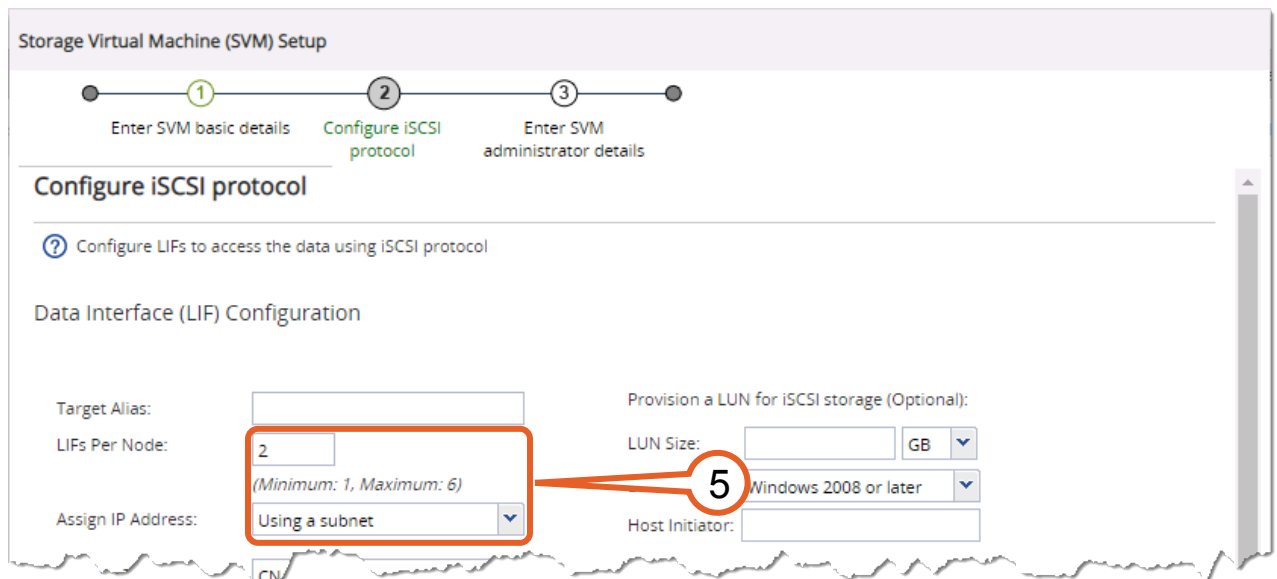
Search Domains:

Name Servers:

**Figure 3-23:**

The Configure iSCSI Protocol stage of the wizard opens.

5. Set the fields in this window as follows:
- “LIFs Per Node”: 2
  - “Assign IP Address”: **Using a subnet**.

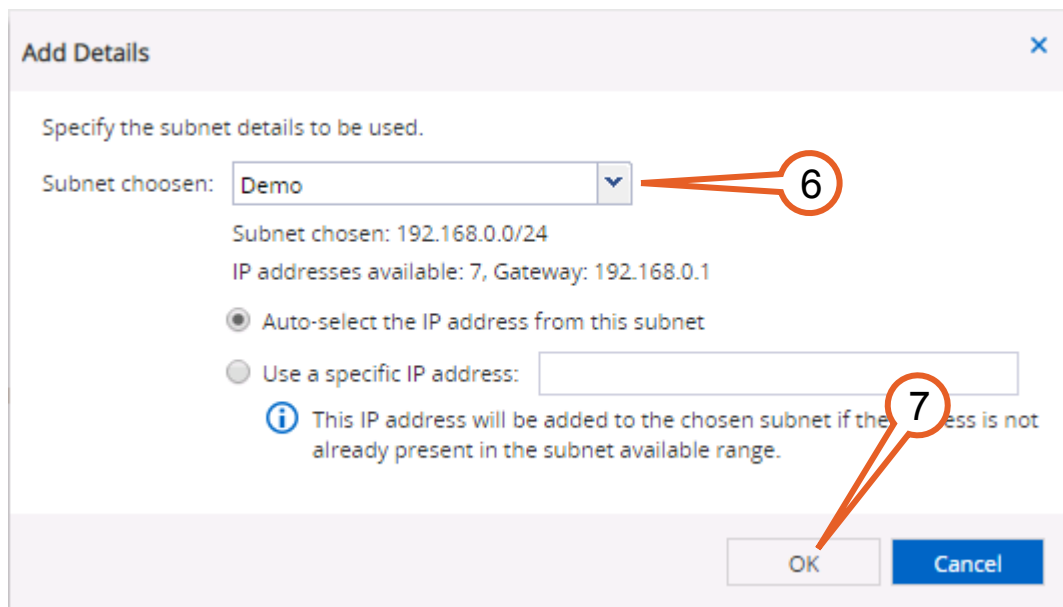


The image shows the 'Storage Virtual Machine (SVM) Setup' wizard, specifically the 'Configure iSCSI protocol' stage. A progress bar at the top indicates three steps: 1. Enter SVM basic details, 2. Configure iSCSI protocol (current), and 3. Enter SVM administrator details. The 'Data Interface (LIF) Configuration' section contains several fields. A red box highlights the 'LIFs Per Node' field, which is set to '2' (with a note '(Minimum: 1, Maximum: 6)'). Below it, the 'Assign IP Address' dropdown is set to 'Using a subnet'. A red callout bubble with the number '5' points to the 'Assign IP Address' dropdown. To the right, there are optional fields for 'Provision a LUN for iSCSI storage', including 'LUN Size' (set to 'GB'), 'Host Initiator' (set to 'Windows 2008 or later'), and a 'Host Initiator' text field.

**Figure 3-24:**

The “Add Details” dialog opens.

6. The **Auto-select the IP address from this subnet** radio button is already selected, which is what you want for this exercise.
7. Click **OK**.



The image shows the 'Add Details' dialog box. It prompts the user to 'Specify the subnet details to be used.' The 'Subnet chosen:' dropdown is set to 'Demo'. Below this, it displays 'Subnet chosen: 192.168.0.0/24' and 'IP addresses available: 7, Gateway: 192.168.0.1'. There are two radio buttons: 'Auto-select the IP address from this subnet' (which is selected) and 'Use a specific IP address:'. A red callout bubble with the number '6' points to the 'Subnet chosen:' dropdown. Another red callout bubble with the number '7' points to the 'OK' button at the bottom right of the dialog.

**Figure 3-25:**

The “Add Details” dialog closes, and focus returns to the “Configure iSCSI Protocol” stage in the “Storage Virtual Machine (SVM) Setup” wizard.

8. Continue setting the fields in this window as follows:
  - “LIFs Per Node”: 2
  - “Assign IP Address”: **Using a subnet.**
  - “Adapter Type”: **NIC.**
9. The “Provision a LUN for iSCSI Storage (Optional)” section shows how to quickly create a LUN when first creating an SVM. This lab guide does not use that method, but instead shows you in a later exercise the more common activity of adding a new volume and LUN to an existing SVM.
10. Check the **Review or modify LIF configuration (Advanced Settings)** check box.

**Figure 3-26:**

Once you check the **Review or modify LIF configuration** check box, the “Configure iSCSI Protocol” window changes to include a list of the LIFs that the wizard plans to create, and some of the previously editable fields in the window now become read-only.

11. Note the LIF interface names and home ports that the wizard has chosen to create.
12. Since this lab utilizes a cluster that only has two nodes, and those nodes are configured as an HA pair, there is no need to create a portset as ONTAP's automatically configured Selective LUN Mapping is more than sufficient for this lab. In other words, leave “Number of portsets” at 0.

13. Click **Submit & Continue**.

Adapter Type: NIC

☒ Review or modify LIF configuration (Advanced Settings)

Number of portsets: 0 (Minimum: 0, Maximum: 1) 12

i Double-click row to edit

Node Name	Interface Name	Home P...	Portset
cluster1-01	cluster1-01_iscsi_lif_1	e0d	
cluster1-01	cluster1-01_iscsi_lif_2	e0e	
cluster1-02	cluster1-02_iscsi_lif_1	e0d	
cluster1-02	cluster1-02_iscsi_lif_2	e0e	

11

13

Skip Submit & Continue Cancel

**Figure 3-27:**

The wizard advances to the SVM Administration stage. Unlike data LIFs for NAS protocols, which automatically support both data and management functionality, iSCSI LIFs only support data protocols, so you must create a dedicated management LIF for this new SVM.

14. Set the fields in the window as follows:

- “Password”: **netapp123**
- “Confirm Password”: **netapp123**

15. Set the “Assign IP Address” dropdown to **Using a subnet**.

Storage Virtual Machine (SVM) Setup

1 Enter SVM basic details 2 Configure iSCSI protocol 3 Enter SVM administrator details

**SVM Administration (optional)**

Specify the following details to enable host side applications such as SnapDrive and SnapManager

To enable the SVM administrator to create volumes, you must assign aggregates to the SVM by using Edit SVM dialog

**Administrator Details**

Username: vsadmin

Password: .....

Confirm Password: .....

**Management Interface (LIF) Configuration for SVM**

☒ Create a new LIF for SVM management

For CIFS and NFS protocols, data LIFs have management access enabled by default; therefore, create a new management LIF only if required. However, for iSCSI and FC/FCoE protocols, create a dedicated management LIF because data LIFs cannot be used for SVM management.

Assign IP Address: Using a subnet

**Figure 3-28:**

The “Add Details” dialog opens.

16. The default values are all suitable, so click **OK**.

**Add Details**

Specify the subnet details to be used.

Subnet chosen: Demo

Subnet chosen: 192.168.0.0/24

IP addresses available: 3, Gateway: 192.168.0.1

☒ Auto-select the IP address from this subnet

☐ Use a specific IP address:

*This IP address will be added to the chosen subnet if the address is not already present in the subnet available range.*

OK Cancel

**Figure 3-29:**

The “Add Details” dialog closes, and focus returns to the “SVM Administration” stage of the “Storage Virtual Machine (SVM) Setup” wizard.

17. Click the **Browse** button next to the “Port:” text box.

Storage Virtual Machine (SVM) Setup

1 Enter SVM basic details 2 Configure iSCSI protocol 3 Enter SVM administrator details

**SVM Administration (optional)**

? Specify the following details to enable host side applications such as SnapDrive and SnapManager

? To enable the SVM administrator to create volumes, you must assign aggregates to the SVM by using Edit SVM dialog

**Administrator Details**

Username: vsadmin

Password: .....

Confirm Password: .....

**Management Interface (LIF) Configuration for SVM**

☒ Create a new LIF for SVM management

For CIFS and NFS protocols, data LIFs have management access enabled by default; therefore, create a new management LIF only if required. However, for iSCSI and FC/FCoE protocols, create a dedicated management LIF because data LIFs cannot be used for SVM management.

Assign IP Address: Using a subnet

Subnet: Demo Change

auto-assign mode

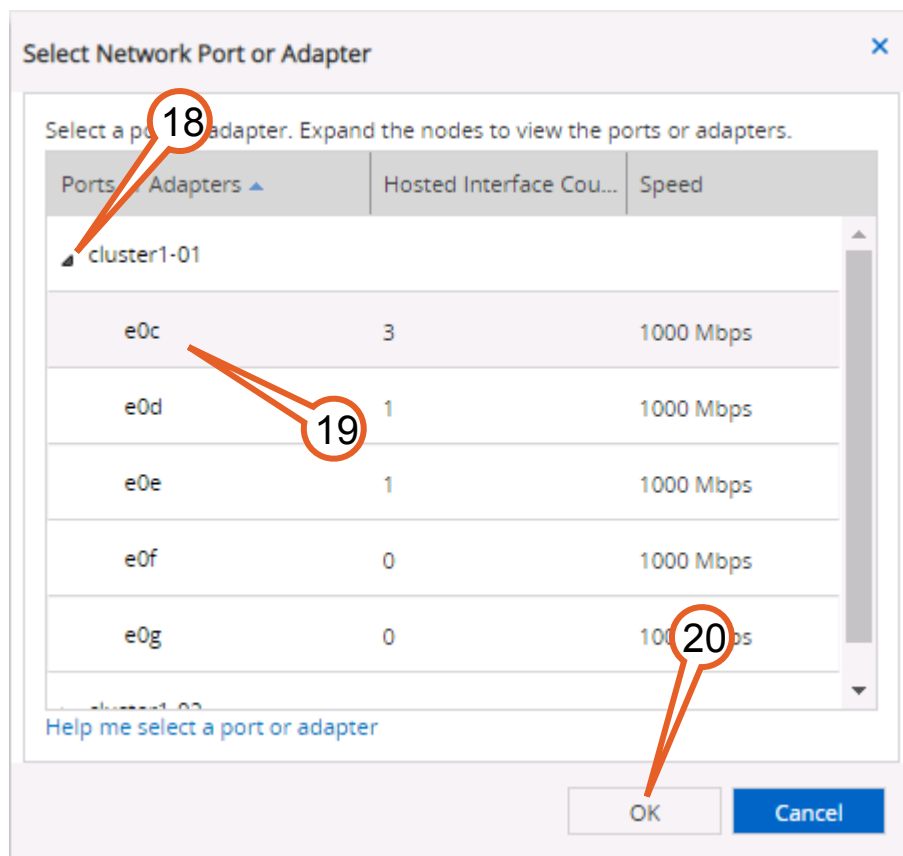
? Port: Browse...

**Figure 3-30:**

The “Select Network Port or Adapter” dialog opens.

18. Expand the entry for **cluster1-01**.
19. Select port **e0c**.

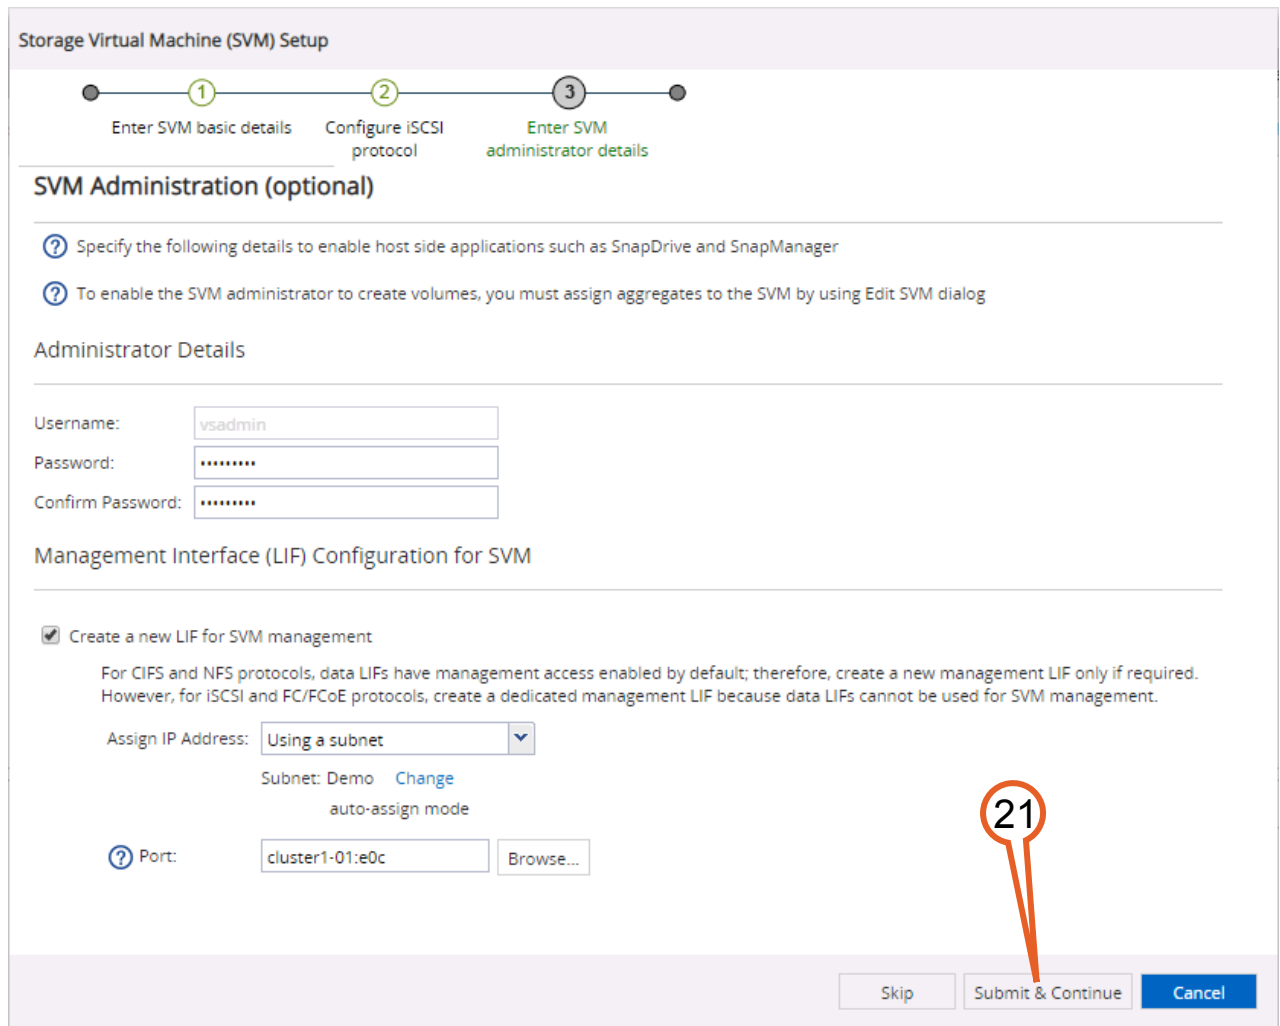
20. Click **OK**.



**Figure 3-31:**

The “Select Network Port or Adapter” dialog closes, and focus returns to the “SVM Administration” stage of the “Storage Virtual Machine (SVM) Setup” wizard.

21. Click **Submit & Continue**.

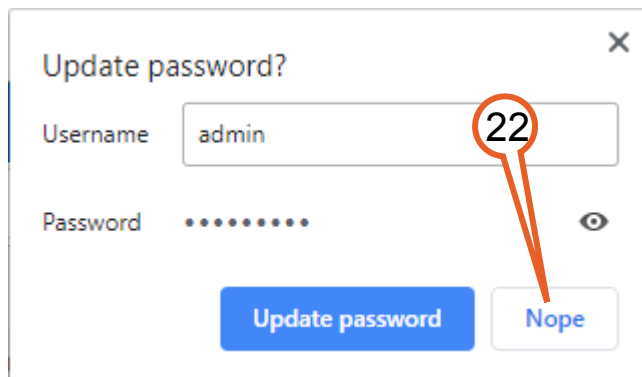


The screenshot shows the 'Storage Virtual Machine (SVM) Setup' wizard. At the top, a progress bar indicates three steps: 1. Enter SVM basic details, 2. Configure iSCSI protocol, and 3. Enter SVM administrator details. The third step is currently active. Below the progress bar, the section 'SVM Administration (optional)' contains two informational messages. The 'Administrator Details' section has input fields for Username (vsadmin), Password (masked with dots), and Confirm Password (masked with dots). The 'Management Interface (LIF) Configuration for SVM' section has a checked checkbox 'Create a new LIF for SVM management' and explanatory text. Below this, there is a dropdown for 'Assign IP Address' set to 'Using a subnet', a 'Subnet' field set to 'Demo' with a 'Change' link, and a 'Port' field set to 'cluster1-01:e0c' with a 'Browse...' button. At the bottom right, there are three buttons: 'Skip', 'Submit & Continue', and 'Cancel'. An orange callout bubble with the number '21' points to the 'Submit & Continue' button.

**Figure 3-32:**

The wizard advances to the “New Storage Virtual Machine (SVM) Summary” stage, and the web browser opens the “Update Password” dialog.

22. In the “Update Password” dialog, click **Nope**.



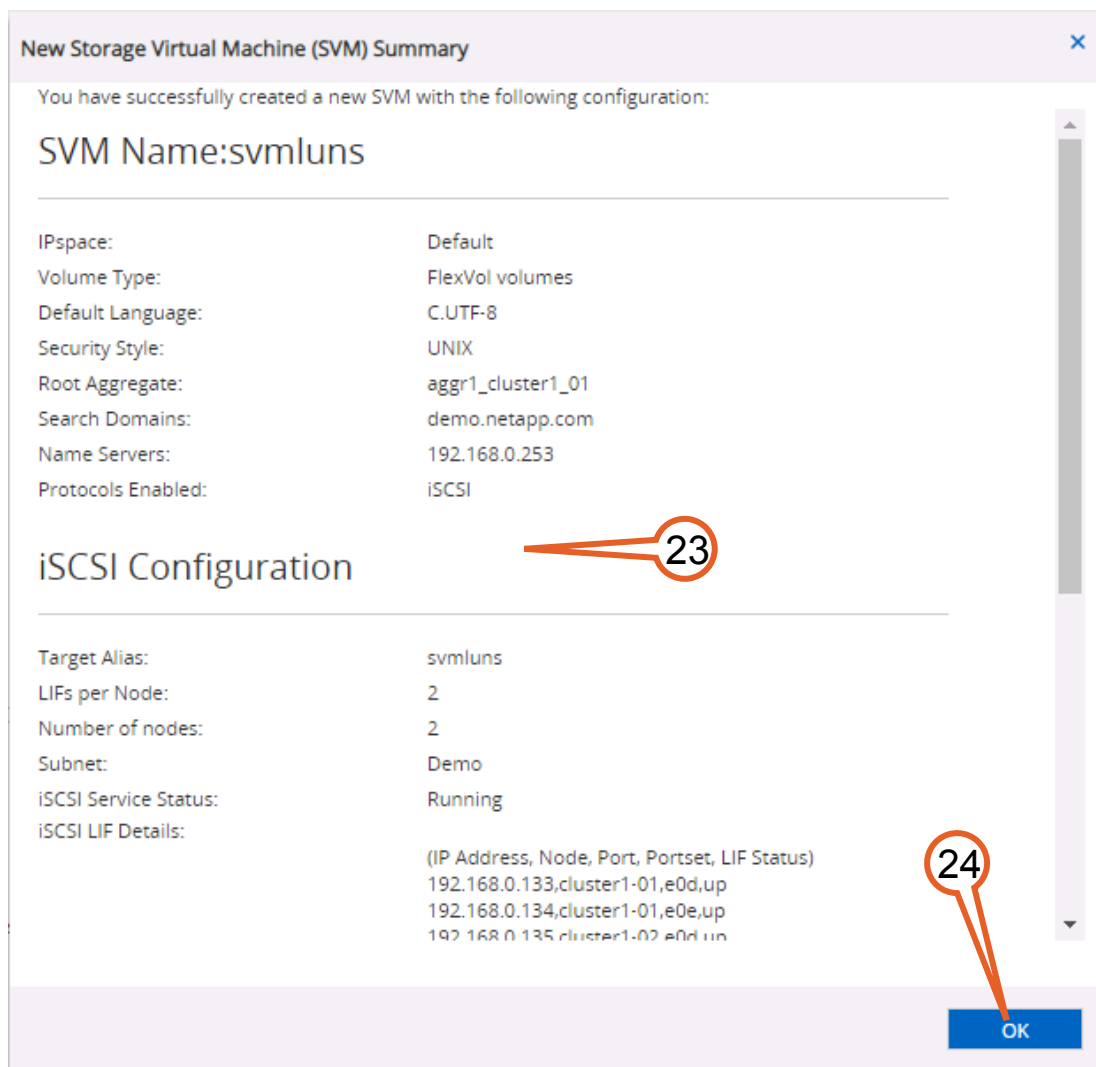
The screenshot shows a dialog box titled 'Update password?'. It has a close button (X) in the top right corner. The 'Username' field contains 'admin'. The 'Password' field is masked with dots and has an eye icon to its right. At the bottom, there are two buttons: 'Update password' and 'Nope'. An orange callout bubble with the number '22' points to the 'Nope' button.

**Figure 3-33:**

The “Update Password” dialog closes.



23. Review the contents of the “New Storage Virtual Machine (SVM) Summary” stage of the wizard, taking note of the names, IP addresses, and port assignments for the 4 iSCSI LIFs, and the management LIF that the wizard created for you.
24. Click **OK** to close the window.

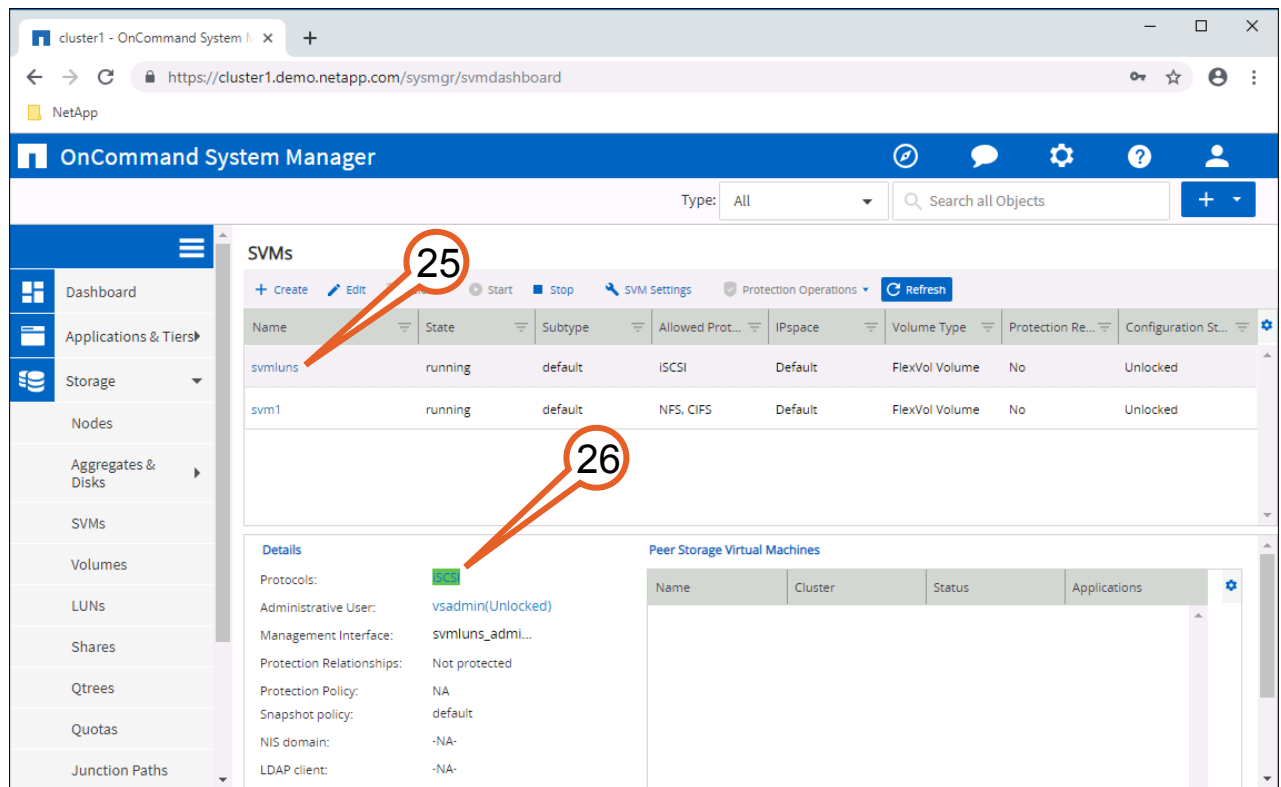


**Figure 3-34:**

The “New Storage Virtual Machine (SVM) Summary” window closes, and focus returns to System Manager, which now displays a summary view for the new “svmluns” SVM.

25. Select the **svmluns** entry if it is not already selected.

26. In the “Details” pane, observe that the “Protocols” field lists iSCSI in a green background, indicating that iSCSI is running. Click the **iSCSI** link.



**Figure 3-35:**

System Manager displays the “SVM Settings” for the “svmluns” SVM.

27. Make sure the “SVM” dropdown is set to **svmluns**.



**Tip:** A number of different pages in System Manager use this sort of dropdown to limit a view's context. It is worthwhile to get in the habit of looking for such a dropdown in this location on each System Manager page you visit, as these can help you avoid potential confusion when you otherwise would not see expected entries listed in a view.

28. In the “SVM Settings” pane, the **iSCSI** tab is pre-selected.
29. In the right pane, the **Service** tab is pre-selected.
30. Observe that the “iSCSI Service” field shows “iSCSI service is running”.

31. The “iSCSI Interfaces” box displays details of the network interfaces that are associated with this iSCSI service on the “svmluns” SVM.

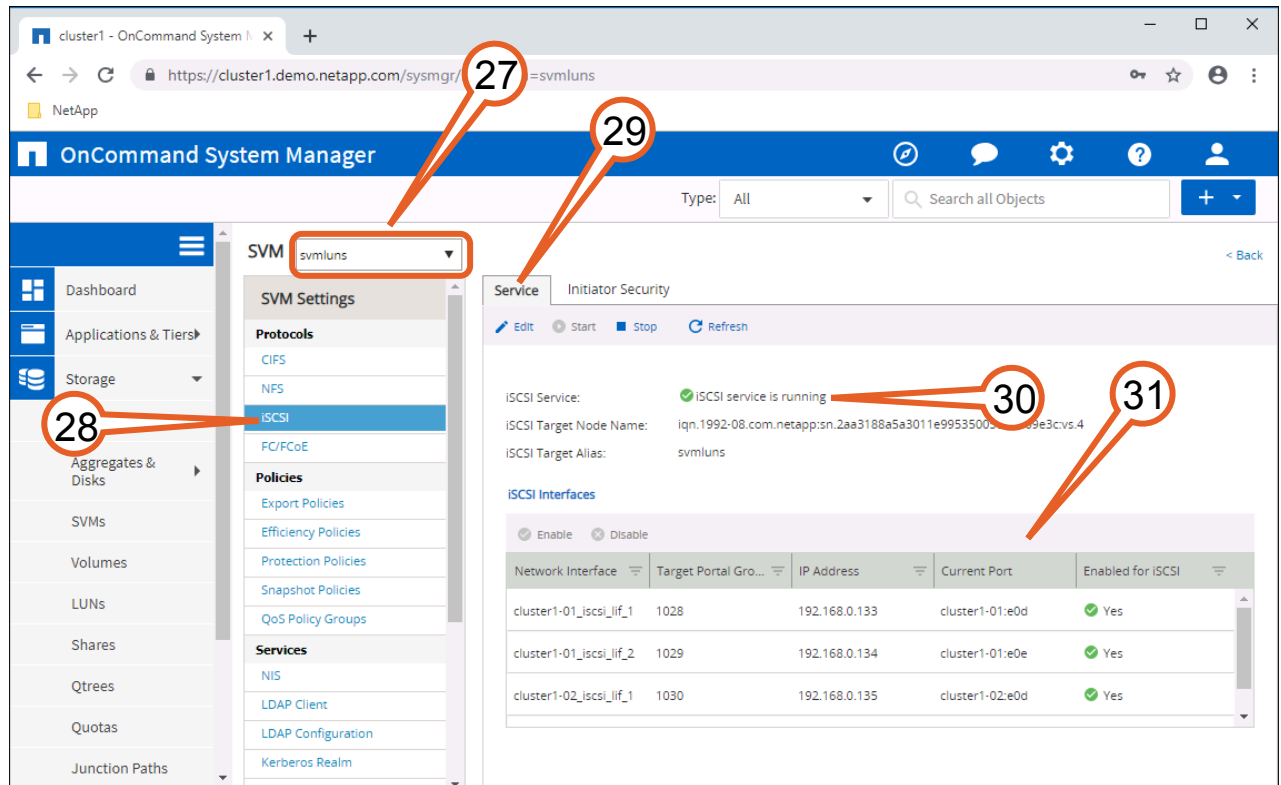


Figure 3-36:

## 3.2.2 Create, Map, and Mount a Windows LUN

In the preceding section you created a new SVM and configured it for iSCSI. In the following sub-sections you will perform the remaining steps needed to configure and use a LUN under Windows:

- Gather the iSCSI Initiator Name of the Windows client.
- Create a thin provisioned Windows volume, create a thin provisioned Windows LUN within that volume, and map the LUN so it can be accessed by the Windows client.
- Mount the LUN on a Windows client leveraging multi-pathing.

You must complete all of these subsections in order to use the LUN from the Windows client.

### 3.2.2.1 LUN Concepts

In ONTAP, LUNs are simply large files that reside within a FlexVol. A single FlexVol can support multiple LUNs, and it is quite common for a client that is using multiple LUNs to have all of those LUNs co-located on the same FlexVol. This facilitates taking application-consistent snapshots, and generally makes LUN management related to the client easier.

While it is possible to configure your SVM so that a volume containing a LUN is also accessible through CIFS or NFS, there is usually no advantage to be gained by doing so. Additionally, if you do not carefully manage export rules and file permissions, you can inadvertently expose the LUN file to unintended NAS clients. While it is perfectly acceptable to serve NAS and SAN clients off of the same SVM, best practice is to not mix NAS files and LUNs within the same volume, or to expose LUN volumes over NAS protocols.

An initiator group, or igroup, defines a list of the Fibre Channel WWPNs or iSCSI node names of the SAN clients that are permitted to access (or even see) a set of LUNs. Mapping a LUN to an igroup applies the igroup's access restrictions to the LUN.

Network access to an iSCSI (or FC) LUN occurs over a LIF. Unlike NAS LIFs, that can migrate online from one node to another in a cluster, SAN LIFs cannot be moved while online. Instead, Asymmetric Logical Unit Access (ALUA) provides redundant paths and automatic path selection as part of any ONTAP SAN solution. When there is an I/O interruption on a given LIF, the client simply (and transparently) retries by switching the I/O to another LIF.

Clients utilize multipathing software in conjunction with ALUA to support multiple paths to a LUN, and for path selection. Windows and Linux both include native multipathing software.

NetApp recommends configuring two SAN LIFs on each node that participate in a SAN configuration in order to provide path redundancy. The optimal path for accessing a LUN is usually to connect directly to the node that is hosting the LUN (i.e., the node that is directly connected to the SSDs/HDDs comprising the aggregate on which the volume hosting the LUN resides). In an HA configuration, that includes the node owning the LUN and its HA partner. LIFs on other nodes can also be used to access the LUN, but those paths may suffer a slight performance penalty as data will have to transfer between the involved nodes over the cluster network.

Selective LUN Mapping (SLM) is an ONTAP feature that reduces the number of paths from the host to the LUN. With SLM, when a new LUN map is created, the LUN is accessible only through paths on the node owning the LUN and its HA partner. This is very helpful in reducing complexity for large clusters (with many nodes) that may present many potential paths to a given LUN.

**Important:** Beginning with ONTAP 9.7, NetApp has started shipping a 2-node AFF All SAN Array option available for factory-ordered and shipped AFF systems. The new AFF All SAN Array (ASA) offers mission-critical workloads (such as databases), a symmetric active-active pathing in a 2-node configuration to offer an even higher level of continuous availability in a simple setup for SAN workloads. While a standard AFF offers a more scalable multi-protocol solution with active-passive (non-optimized) pathing, the AFF ASA SAN system provide a higher level of uninterrupted access for SAN protocols during a planned or unplanned storage failover event.

### 3.2.2.2 Gather the Windows Client iSCSI Initiator Name

You need to determine the Windows client's iSCSI initiator name so that when you create the LUN you can set up an appropriate initiator group to control access to the LUN.

On the desktop of the Windows client named Jumphost (the main Windows host you use in the lab), perform the following tasks:

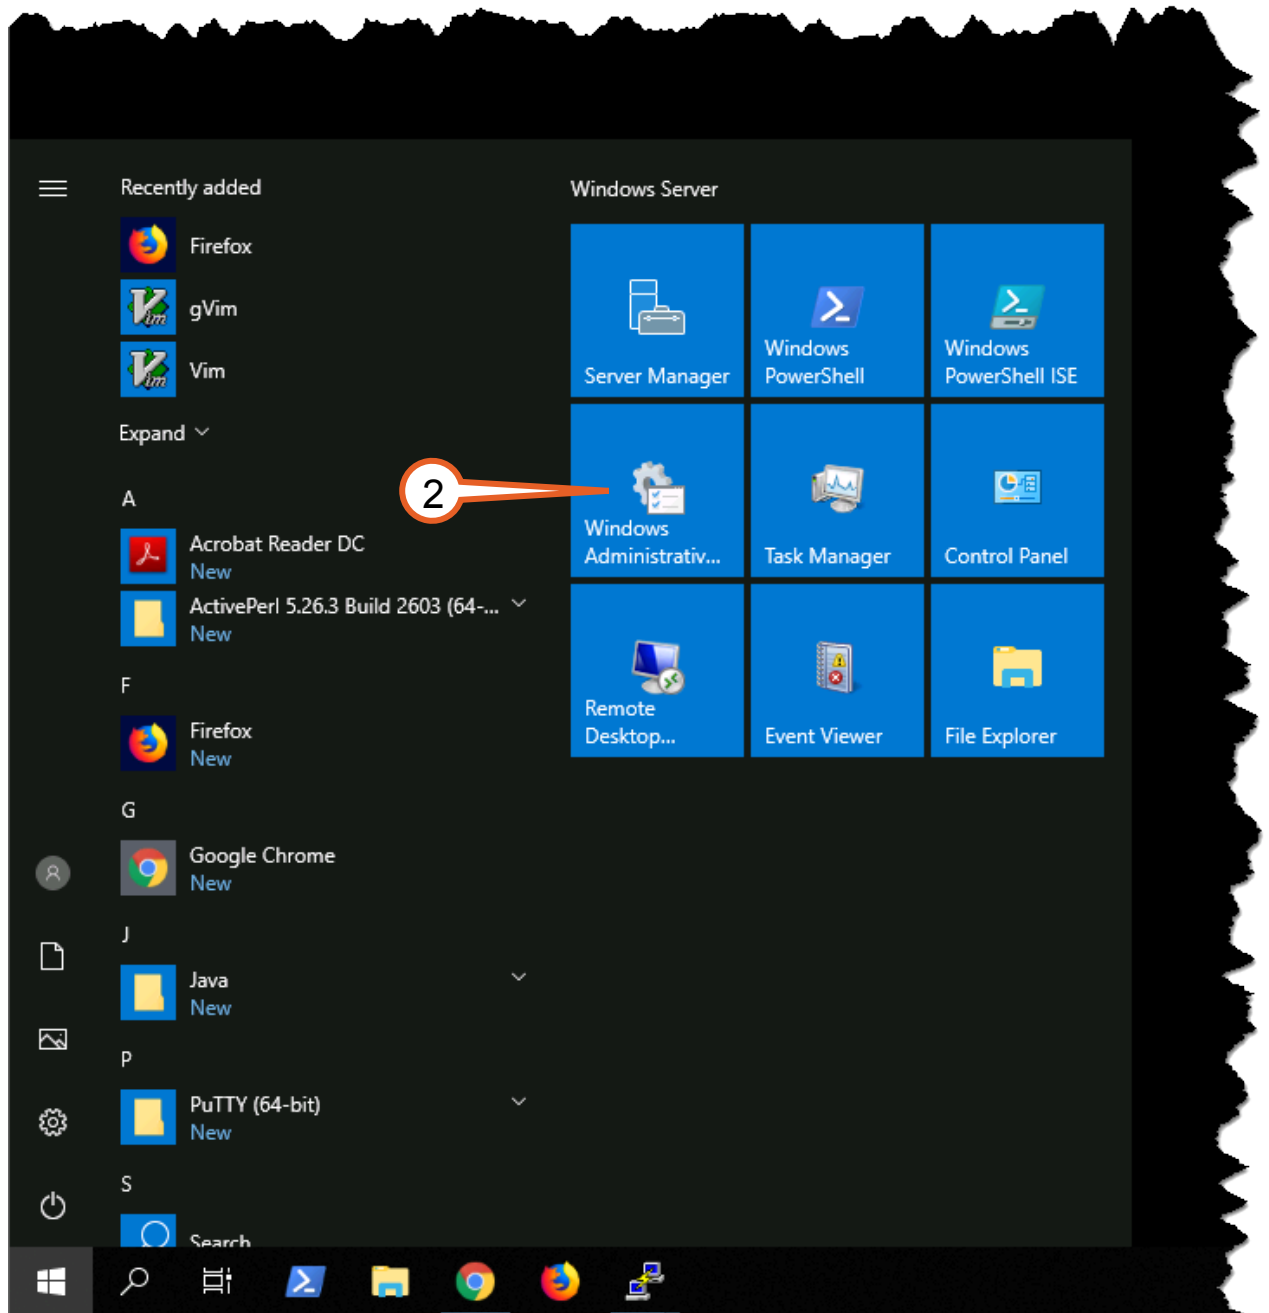
1. Click on the **Windows** icon on the far left side of the task bar.



**Figure 3-37:**

The “Start” screen opens.


2. Click on **Windows Administrative Tools**.



**Figure 3-38:**

Windows Explorer opens to the List of Administrative Tools.

3. Double-click the entry for the **iSCSI Initiator** tool.



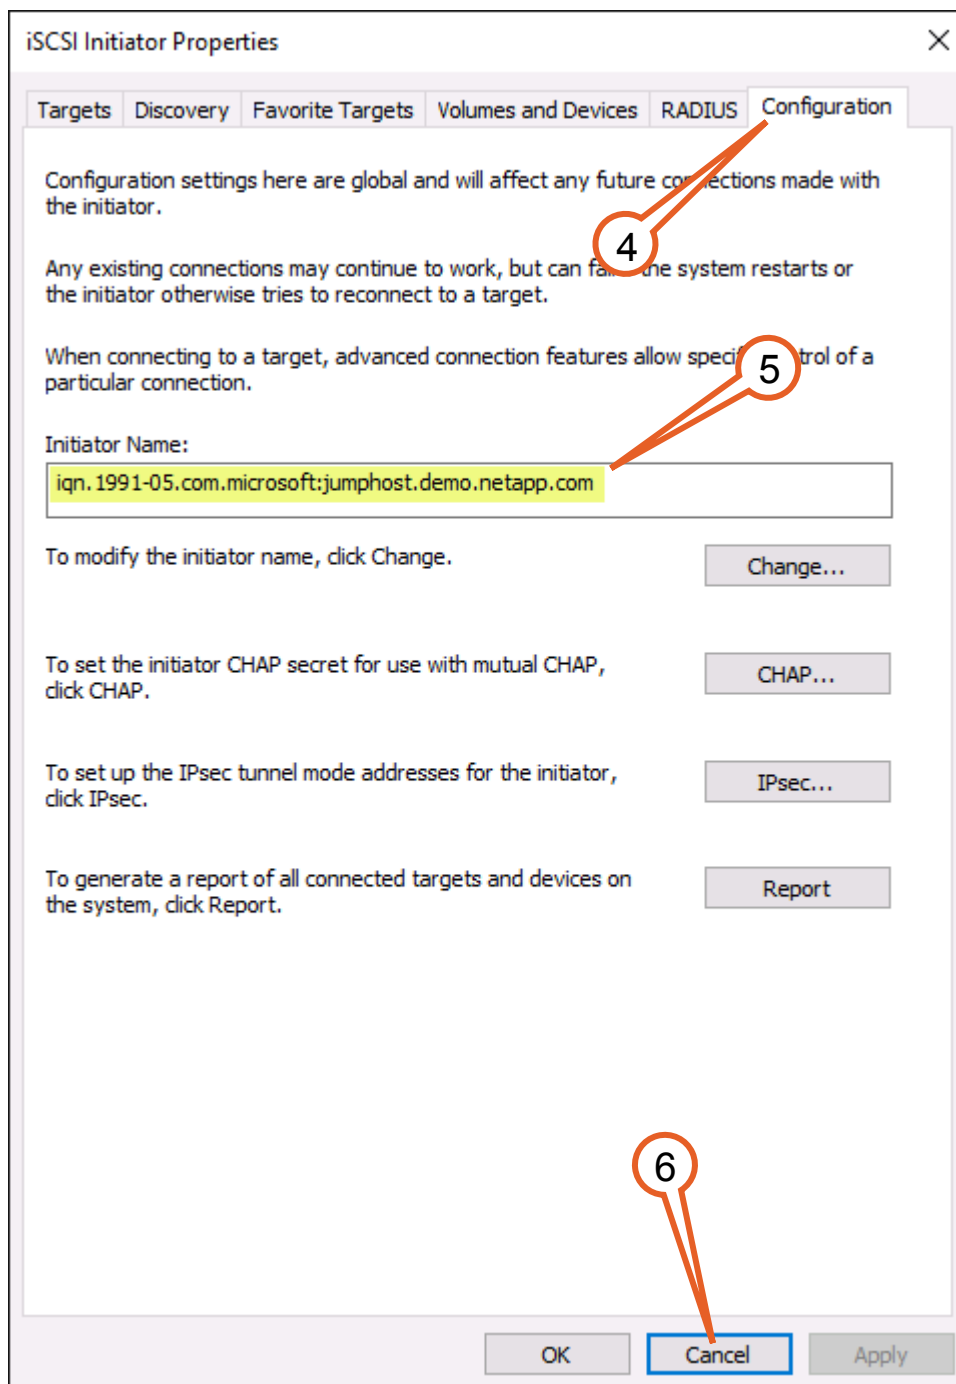
**Figure 3-39:**

The “iSCSI Initiator Properties” window opens.

4. Select the **Configuration** tab.
5. Take note of the value in the “Initiator Name” field, which contains the initiator name for Jumphost.

**Attention:** The initiator name is “**iqn.1991-05.com.microsoft:jumphost.demo.netapp.com**”. You will need this value later, so you might want to copy this value from the properties window and paste it into a text file on your lab's desktop so you have it readily available when that time comes.

6. Click **Cancel**.



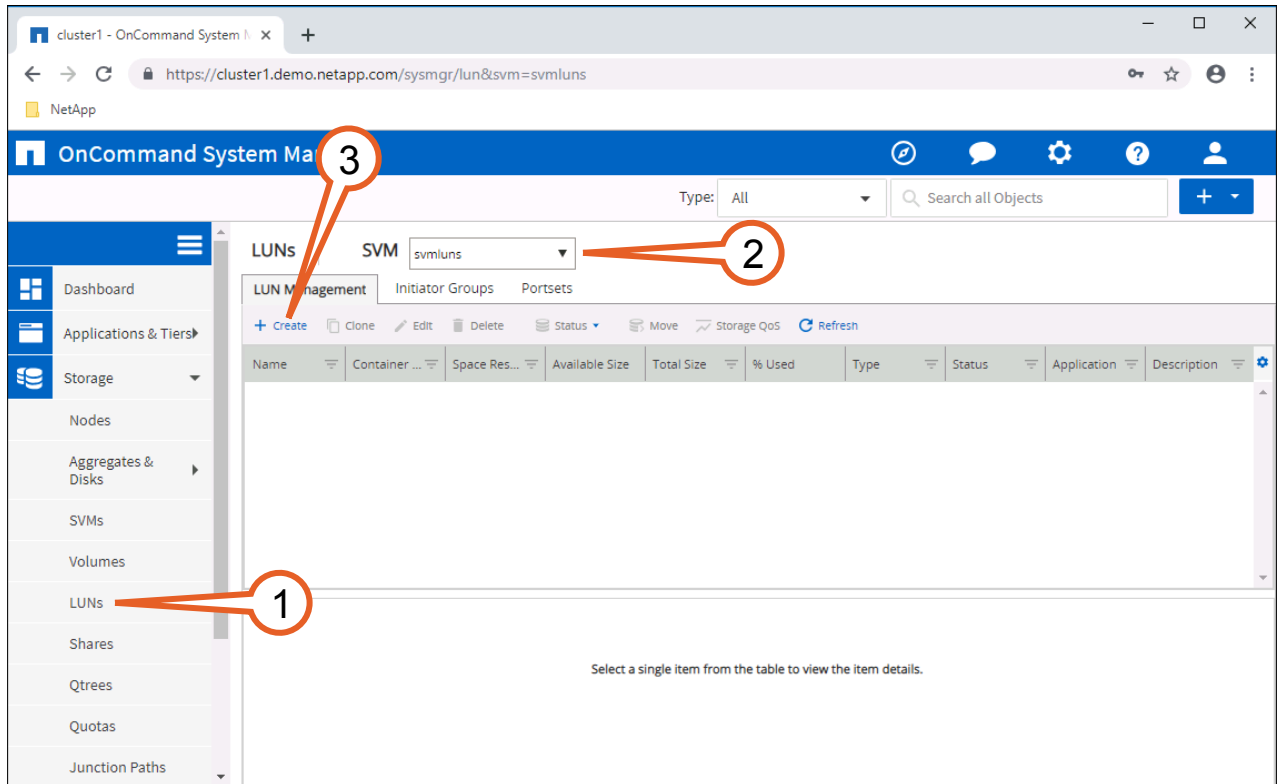
**Figure 3-40:**

The “iSCSI Properties” window closes, and focus returns to the “Windows Explorer Administrator Tools” window. Leave this window open because you will need to access other Administrator tools later in the lab.

### 3.2.2.3 Create and Map a Windows LUN

In this lab activity you create a new thin provisioned Windows LUN named “windows.lun” in the volume “winluns” on the SVM “svmluns”. You will also create an initiator group for the LUN and populate it so that only the Windows host Jumphost can access it.

1. In the left pane of System Manager, navigate to **Storage > LUNs**.
2. Make sure the “LUNS on SVM” drop-down is set to **svmluns**.
3. Click **Create**.

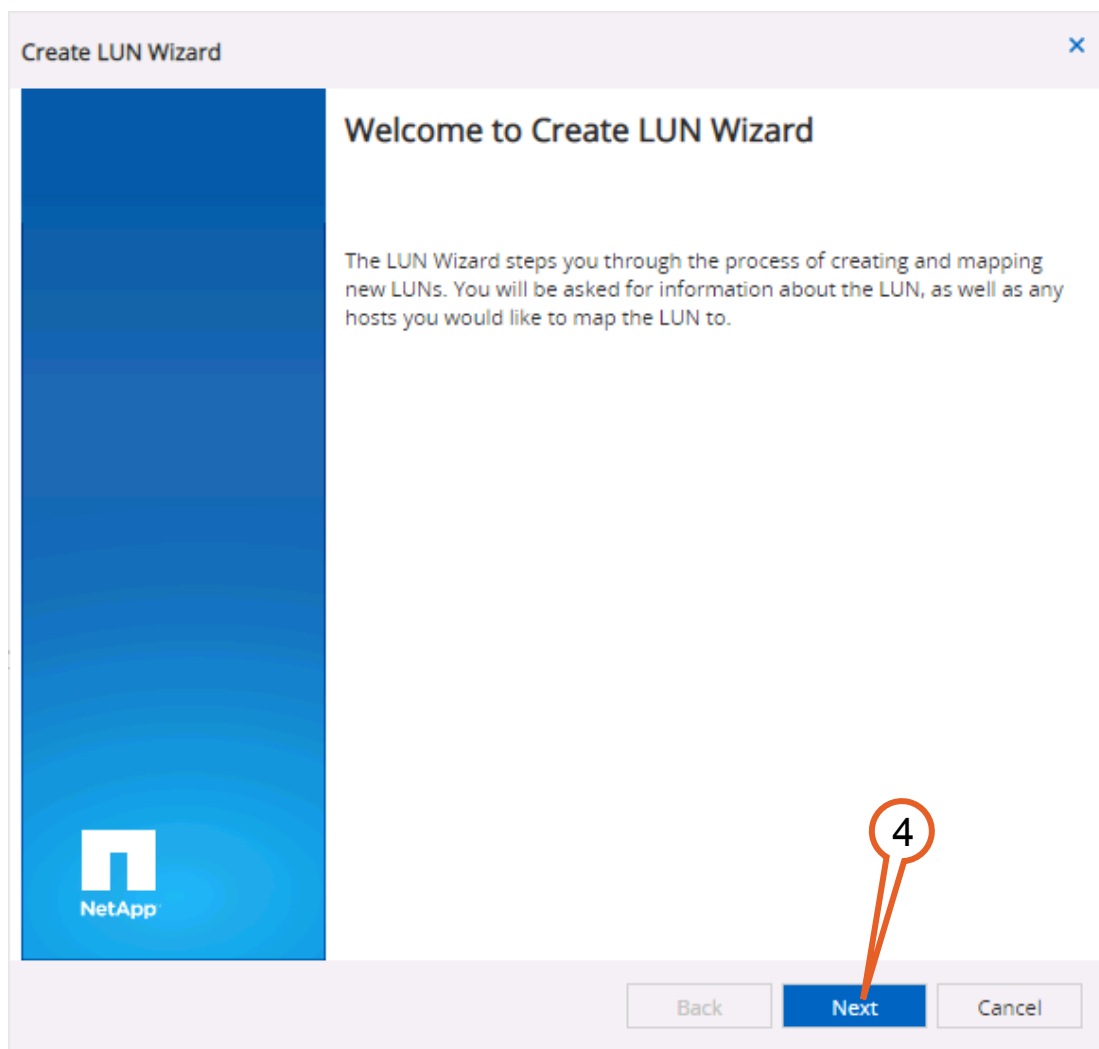


**Figure 3-41:**

The “Create LUN Wizard” dialog opens.



4. Click **Next** to advance to the next step in the wizard.



**Figure 3-42:**

The wizard advances to the “General Properties” step.

5. Set the fields in the window as follows.

- “Name”: `windows.lun`
- “Description”: `windows LUN`
- “Type”: **Windows 2008 or later.**
- “Size”: **10 GB.**
- “Space Reserve”: **Disable.**



**Note:** This will thin provision the LUN.

6. Click **Next** to continue.

**Create LUN Wizard**

**General Properties**  
You can specify the name, size, type, and an optional description for the LUN that you would like to create.

You can enter a valid name for the LUN and an optional short description

Name: windows.lun

Description: Windows LUN (optional)

You can specify the size of the LUN. Storage will be optimized according to the type selected.

Type: Windows 2008 or later (optional)

Size: 10 GB

Space Reserve: Disable (optional)

Buttons: Back, Next, Cancel

**Figure 3-43:**

The wizard advances to the LUN Container step.

7. Select the radio button to **Create a new flexible volume**, and set the fields under that heading as follows.

- “Aggregate Name”: **aggr1\_cluster1\_01**
- “Volume Name”: **winluns**
- “Tiering Policy”: **none**.



**Note:** Tiering Policy is used with FabricPool which is not covered in this lab.

8. When finished click **Next**.

Create LUN Wizard

**LUN Container**  
You can let the wizard create a volume or you can choose an existing volume as the LUN container.

The wizard automatically chooses the aggregate with most free space for creating flexible volume for the LUN. But you can choose a different aggregate of your choice. You can also select an existing volume/qtree to create your LUN.

☐ Select an existing volume or qtree for this LUN

Volume/Qtree:

☒ Create a new flexible volume in

Aggregate Name:

Volume Name:

Tiering Policy:

[Tell me more about cloud tier and tiering policies.](#)

**Figure 3-44:**

The wizard advances to the Initiator Mappings step.

9. Click **Add Initiator Group**.

Create LUN Wizard

**Initiators Mapping**  
You can connect your LUN to the initiator hosts by selecting from the initiator group and by optionally providing LUN ID for the initiator group.

Map	Initiator Group Name	Type	LUN ID (Optional)
-----	----------------------	------	-------------------

☐ Show All Initiator Groups

Add Initiator Group

Back Next Cancel

**Figure 3-45:**

The “Create Initiator Group” dialog opens.

10. Set the fields in the window as follows.

- “Name”: `winigrp`
- “Operating System”: **Windows**
- “Type”: Select the **iSCSI** radio button

11. Click **Initiators**.

The screenshot shows the 'Create Initiator Group' dialog box with the 'Initiators' tab selected. The 'Name' field contains 'winigrp'. The 'Operating System' dropdown is set to 'Windows'. Under the 'Type' section, 'iSCSI' is selected with a radio button. The 'Portset' section includes a text box and a 'Choose' button. At the bottom right are 'Create' and 'Cancel' buttons. Red circles with numbers 10 and 11 highlight the 'winigrp' field, the 'iSCSI' radio button, and the 'Initiators' tab respectively.

**Figure 3-46:**

The “Initiators” tab displays.

12. Click the first line inside the large text box that displays “Add Initiator”. The line will become editable, and you should populate it with the value of the iSCSI Initiator name for Jumphost that you noted earlier.



**Important:** The iSCSI Initiator name is  
**iqn.1991-05.com.microsoft:jumphost.demo.netapp.com.**

13. After you enter the value, click **Create**.

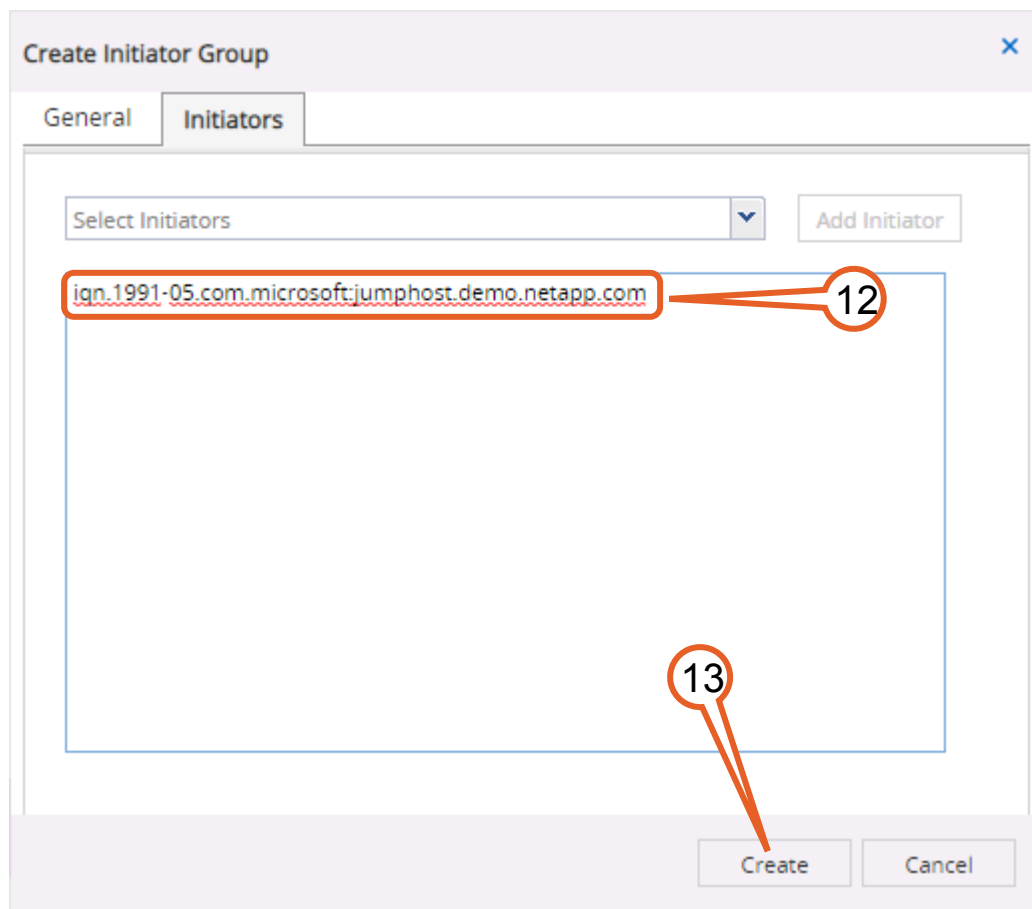


Figure 3-47:

14. If an "Initiator-Group Summary" window opens, click **OK** to acknowledge the confirmation, otherwise proceed to the next step.

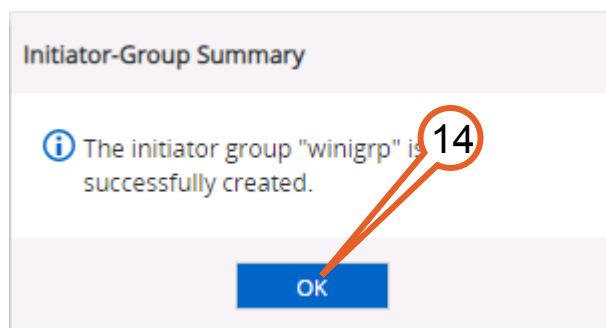


Figure 3-48:

The "Initiator-Group Summary" window closes, and focus returns to the "Initiator Mapping" step of the Create LUN wizard.

15. Click the checkbox under the map column next to the **winigrp** initiator group.

**Important:** This is a critical step because this is where you actually map the new LUN to the new igroup.

16. Click **Next**.

The screenshot shows the 'Create LUN Wizard' window. At the top, a blue header bar contains the text 'Initiators Mapping' and a sub-header 'You can connect your LUN to the initiator hosts by selecting from the initiator group and by optionally providing LUN ID for the initiator group.' Below this is a table with four columns: 'Map', 'Initiator Group Name', 'Type', and 'LUN ID (Optional)'. The first row of the table has a checked checkbox in the 'Map' column, the text 'winigrp' in the 'Initiator Group Name' column, and 'Windows' in the 'Type' column. The 'LUN ID (Optional)' column is empty. Below the table, there is a checkbox labeled 'Show All Initiator Groups' and a button labeled 'Add Initiator Group'. At the bottom of the window, there are three buttons: 'Back', 'Next', and 'Cancel'. The 'Next' button is highlighted in blue. Two orange callout boxes with numbers '15' and '16' are present. Callout 15 points to the checked checkbox in the 'Map' column. Callout 16 points to the 'Next' button.

Map	Initiator Group Name	Type	LUN ID (Optional)
<input checked="" type="checkbox"/>	winigrp	Windows	

☐ Show All Initiator Groups      Add Initiator Group

Back    **Next**    Cancel

**Figure 3-49:**

The wizard advances to the “Storage Quality of Service Properties” step. You will not create any QoS policies in this lab. If you are interested in learning about QoS, please see the *Hands-on Lab for Advanced Concepts for NetApp ONTAP*.

17. Click **Next** to continue.

Create LUN Wizard

Storage Quality of Service Properties  
Limit LUN throughput by assigning it to a Quality of Service policy group

☒ Manage Storage Quality of Service

Apply QoS policy to the LUN by assigning it to a policy group and specify the QoS maximum throughput and QoS minimum throughput values. Storage objects assigned to the same QoS policy will share the same QoS maximum throughput value.

[Tell me more about Storage Quality of Service](#)

Assign to: ☒ New Policy Group ☐ Existing Policy Group

Policy Group Name:

Minimum Throughput:  (IOPS) ⓘ

Maximum Throughput:  (MB/s) ▼

(IOPS)

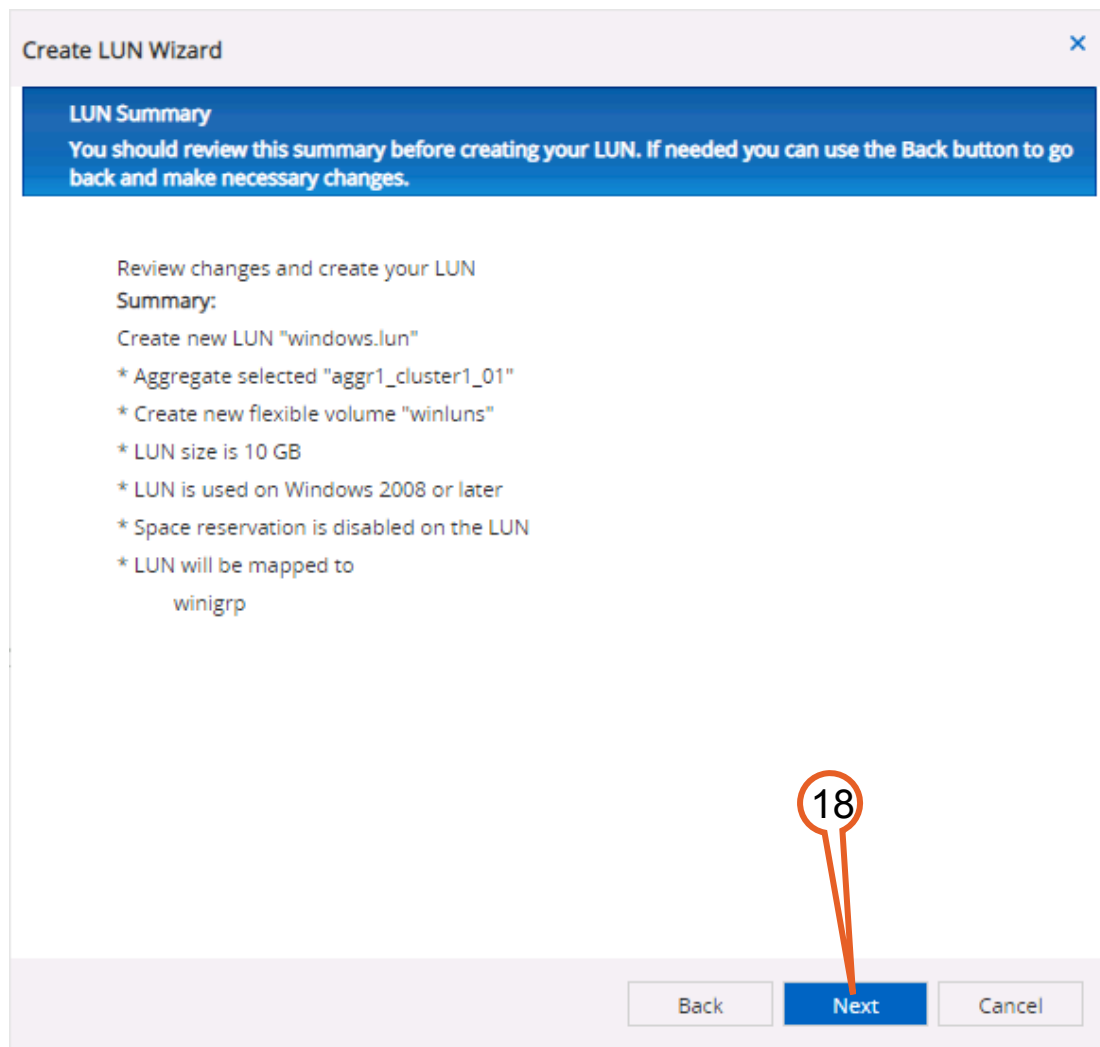
Back Next Cancel

**Figure 3-50:**

The wizard advances to the “LUN Summary” step where you can review your selections before you create the LUN.



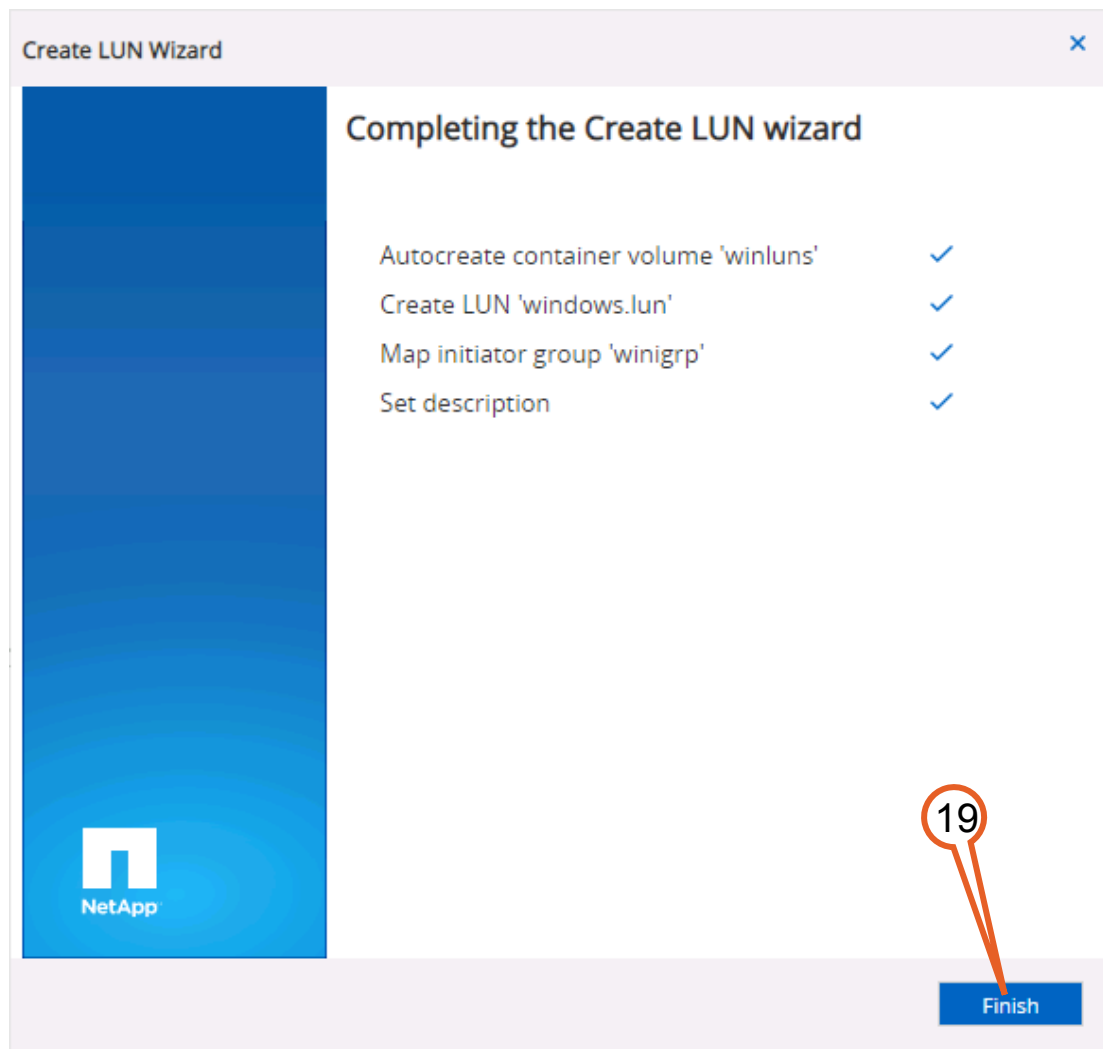
18. If everything looks correct, click **Next**.



**Figure 3-51:**

The wizard begins creating the volume that contains the LUN, creating the LUN, and mapping the LUN to the new igroup. As it finishes each step, the wizard displays a blue check mark in the window next to that step.

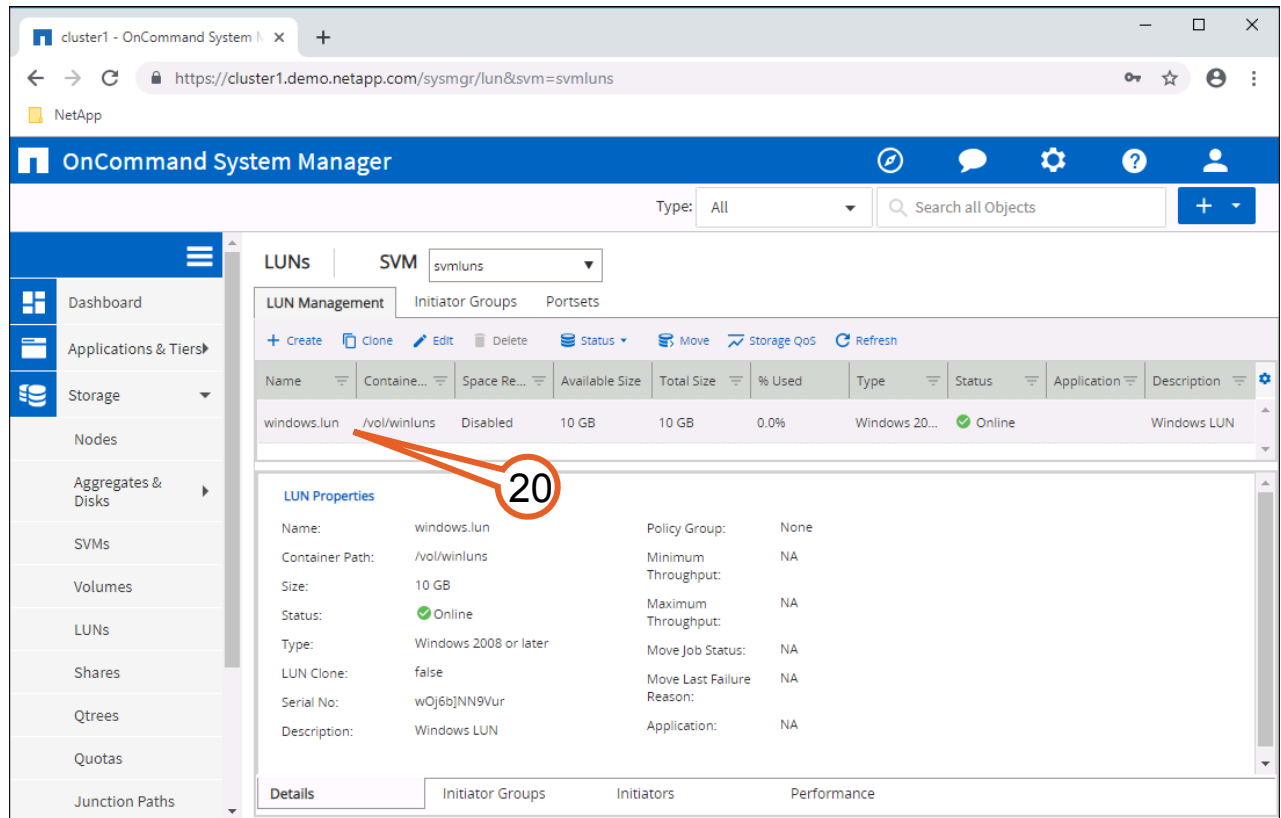
19. Click **Finish** to terminate the wizard.



**Figure 3-52:**

The “Create LUN” wizard window closes, and focus returns to the LUNs view in System Manager.

20. The new LUN “windows.lun” now shows up in the LUNs view, and if you select it you can review its details in the bottom pane.



**Figure 3-53:**

ONTAP includes support for LUN space reclamation, an optional feature that enables ONTAP to reclaim space from a thin provisioned LUN when the client deletes data from it, and which also allows ONTAP to notify the client when the LUN cannot accept writes due to lack of space on the volume. This feature is supported by VMware ESX 5.0 and later, Red Hat Enterprise Linux 6.2 and later, and Microsoft Windows 2012 and later. Jumpshot is running Windows 2019, so you will enable the space reclamation feature for your Windows LUN. You can only enable space reclamation through the ONTAP command line.

21. If you do not already have a PuTTY session open to `cluster1`, open one now by following the instructions in the “Accessing the Command Line” section at the beginning of this guide. Log in using the username `admin` and the password `Netapp1!`.
22. In the cluster1 CLI, view whether space reclamation is enabled for the LUN.

```
cluster1::> lun show -vserver svmluns -path /vol/winluns/windows.lun
               -fields space-allocation
vserver path          space-allocation
-----
svmluns /vol/winluns/windows.lun disabled
cluster1::>
```

23. Enable space reclamation for the LUN windows.lun.

```
cluster1::> lun modify -vserver svmluns -path /vol/winluns/windows.lun
               -space-allocation enabled
cluster1::>
```

24. View the LUN's space reclamation setting again.

```
cluster1::> lun show -vserver svmluns -path /vol/winluns/windows.lun
               -fields space-allocation
vserver path          space-allocation
```

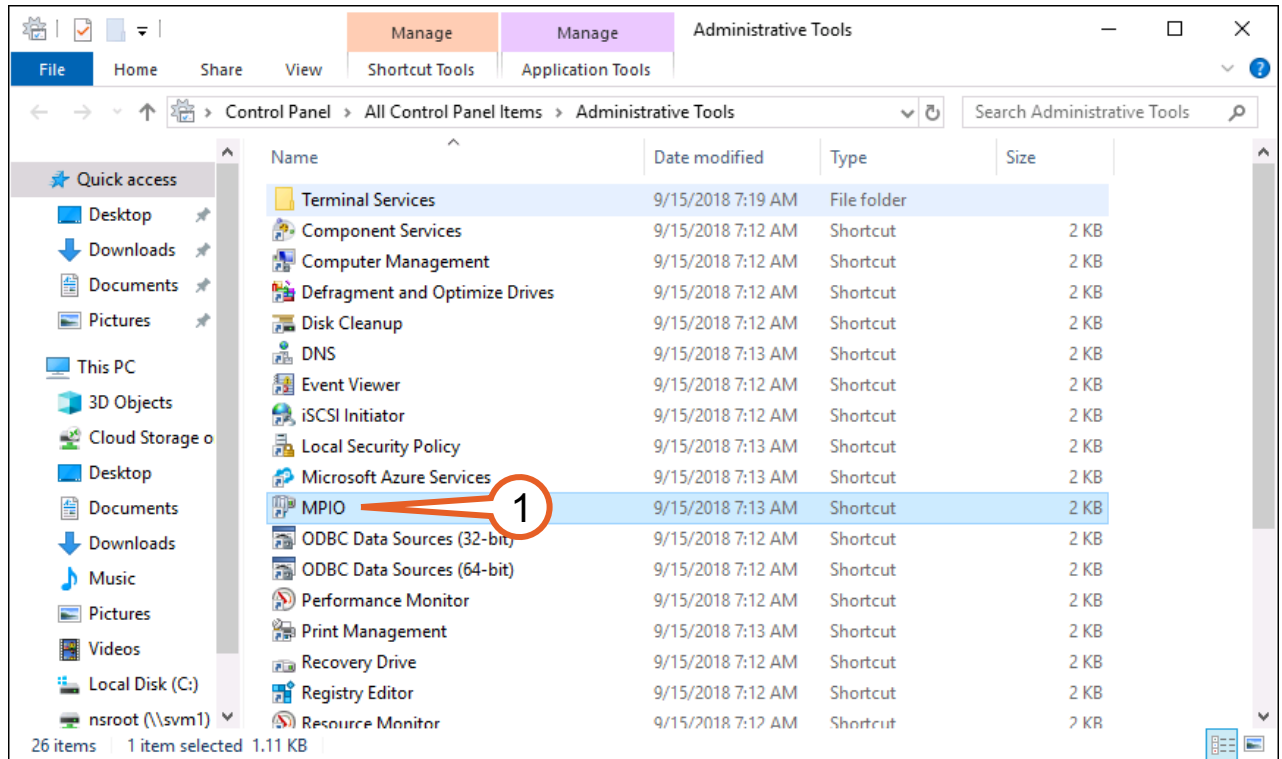
```
svmluns /vol/winluns/windows.lun enabled
cluster1::>
```

### 3.2.2.4 Mount the LUN on a Windows Client

In this lab activity you mount the LUN on the Windows client.

Begin by validating that the Multi-Path I/O (MPIO) software is working properly on this windows host.

1. On the desktop of Jumphost, in the “Administrative Tools” window (which you should still have open from a previous exercise), double-click the **MPIO** tool.



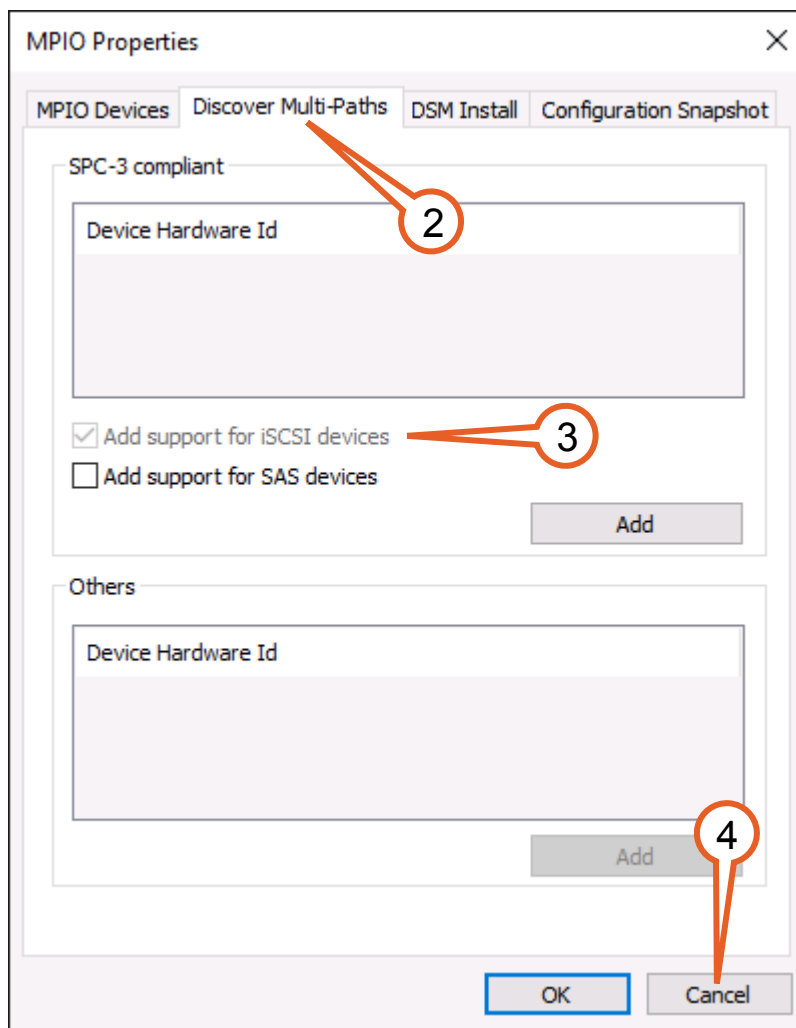
**Figure 3-54:**

2. Select the **Discover Multi-Paths** tab.
3. Examine the **Add Support for iSCSI devices** checkbox. If this checkbox is NOT greyed out then MPIO is improperly configured.



**Tip:** This checkbox should be greyed out for this lab. If it is not, place a check in that checkbox, click **Add**, then if prompted to reboot click **Yes**. Once the system finishes rebooting, return to this window to verify that the checkbox is now greyed out, indicating that MPIO is properly configured.


4. Click **Cancel**.



**Figure 3-55:**

The “MPIO Properties” dialog closes and focus returns to the “Administrative Tools” window for Jumphost. Now you can start connecting Jumphost to the LUN.

5. In “Administrative Tools”, double-click the **iSCSI Initiator** tool.

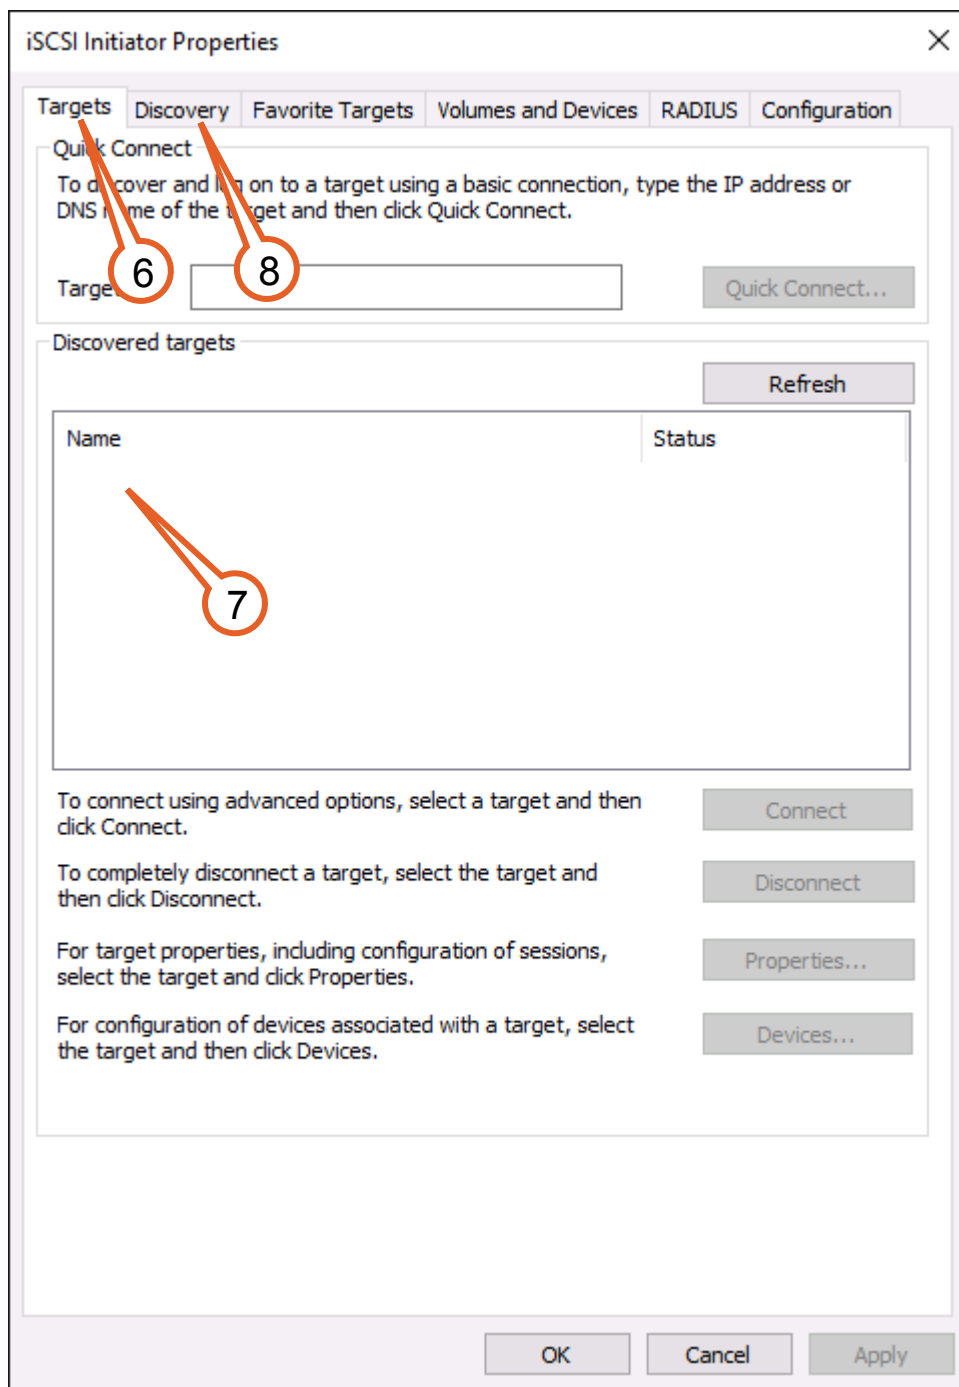


**Figure 3-56:**

The “iSCSI Initiator Properties” dialog opens.

6. Select the **Targets** tab.
7. Notice that there are no targets listed in the “Discovered Targets” list box. This indicates that there are currently no iSCSI targets mapped to this host.

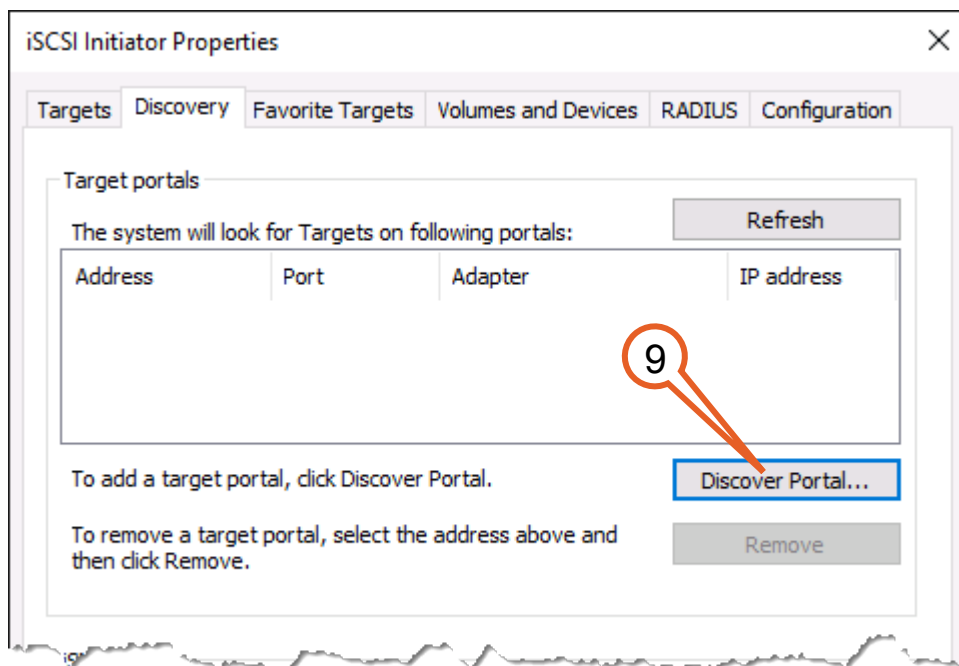
8. Click the **Discovery** tab.



**Figure 3-57:**

The **Discovery** tab is where you begin the process of discovering LUNs, and to do that you must define a target portal to scan. You are going to manually add a target portal.


9. Click the **Discover Portal...** button.



**Figure 3-58:**

The “Discover Target Portal” dialog opens. Here you specify the first of the IP addresses that the ONTAP Create LUN wizard assigned to your iSCSI LIFs when you created the svmluns SVM. Recall that the wizard assigned your LIF’s IP addresses in the range 192.168.0.133-192.168.0.136.

10. Set the “IP Address or DNS name” textbox to 192.168.0.133, the first address in the range for your LIFs.
11. Click **OK**.



**Figure 3-59:**

The “Discover Target Portal” dialog closes, and focus returns to the “iSCSI Initiator Properties” dialog.



12. The “Target Portals” list now contains an entry for the IP address you entered in the previous step.

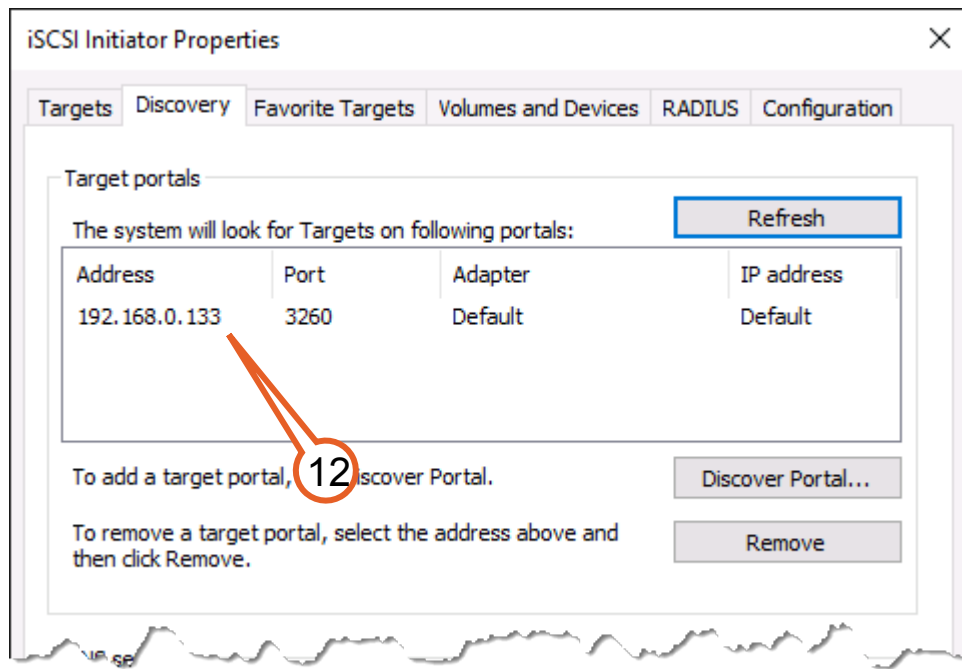


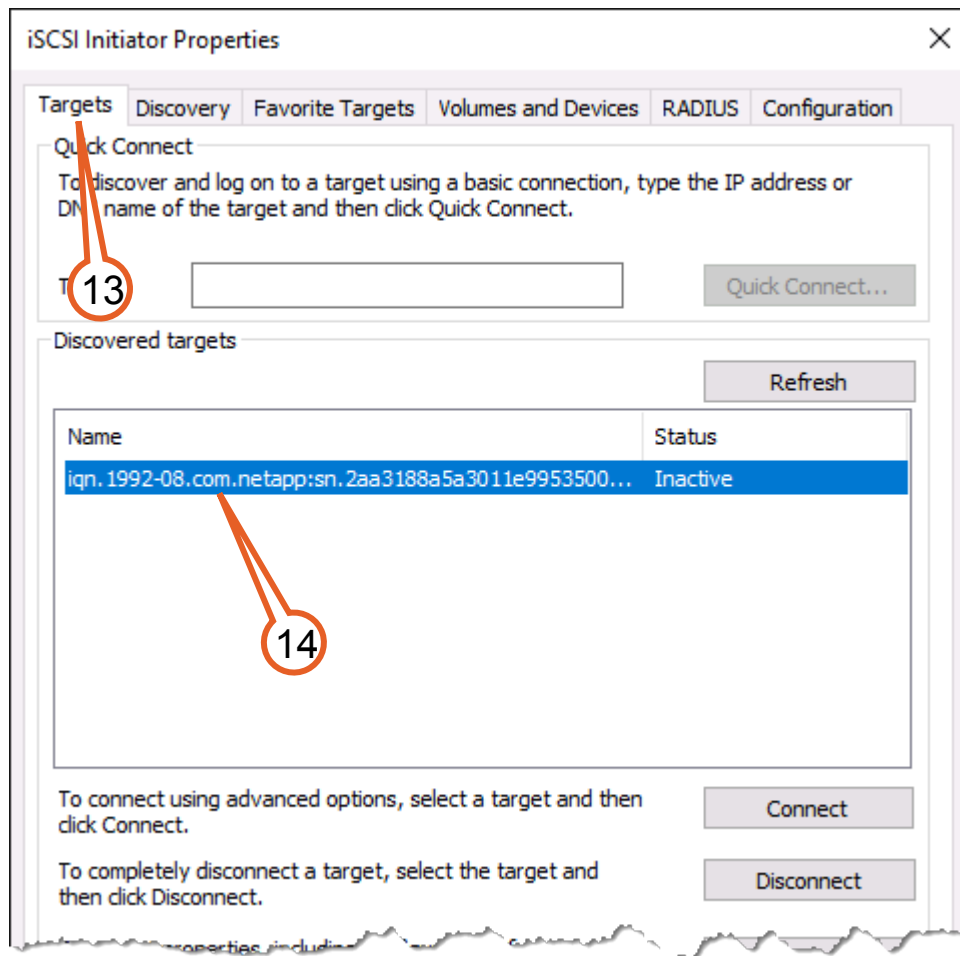
Figure 3-60:

13. Click on the **Targets** tab.
14. In the “Discovered targets”, list select the only listed target, **iqn.1992-08.com.netapp:....**. Observe that the target's status is Inactive, because although you have discovered it you have not yet connected to it. Also note that the “Name” of the discovered target in your lab will have a different value than what you see in this guide; that name string is uniquely generated for each instance of the lab.



**Note:** Make a mental note of that string value as you will see it a lot as you continue to configure iSCSI in later steps of this procedure.

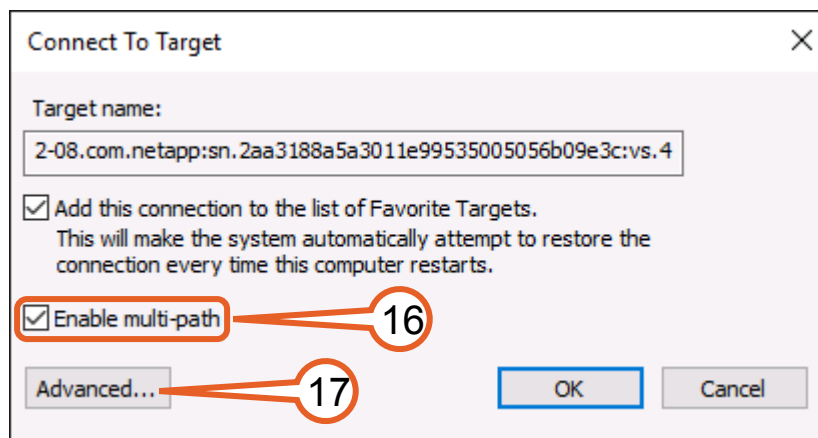
15. Click **Connect**.



**Figure 3-61:**

The “Connect to Target” dialog box opens.

16. Click the **Enable multi-path** checkbox.
17. Click **Advanced...**



**Figure 3-62:**

The “Advanced Settings” dialog opens.

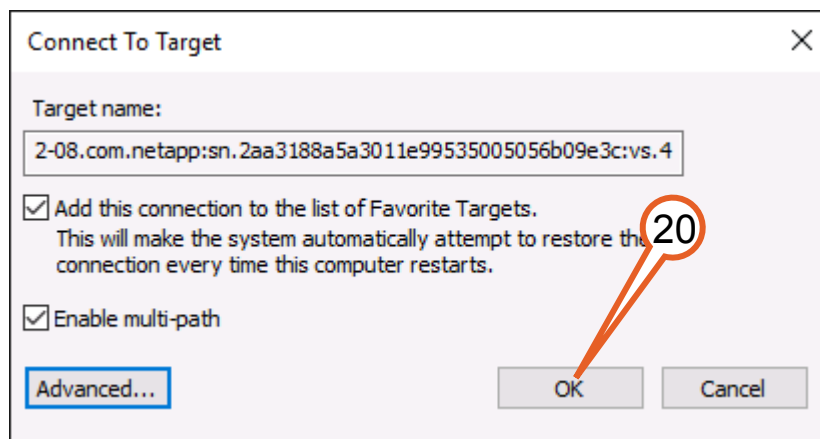
18. In the “Target portal IP” dropdown menu select the entry containing the IP address you specified when you discovered the target portal, which should be **192.168.0.133**. The listed values are IP Address and Port number combinations, and the specific value you want to select is **192.168.0.133 / 3260**.
19. When finished, click **OK**.

The screenshot shows the 'Advanced Settings' dialog box with the 'IPsec' tab selected. The 'Connect using' section has three dropdown menus: 'Local adapter' (Default), 'Initiator IP' (Default), and 'Target portal IP' (192.168.0.133 / 3260). The 'Target portal IP' dropdown is highlighted with a red box and a callout labeled '18'. Below this, the 'CRC / Checksum' section has two checkboxes: 'Data digest' and 'Header digest'. The 'Enable CHAP log on' checkbox is also present. The 'CHAP Log on information' section includes a text box for 'Name' (iqn.1991-05.com:microsoft:jumphost.demo.netapp.com) and a text box for 'Target secret'. The 'Perform mutual authentication' checkbox is also present. The 'Use RADIUS to generate user authentication credentials' and 'Use RADIUS to authenticate target credentials' checkboxes are also present. The 'OK' button is highlighted with a red box and a callout labeled '19'.

**Figure 3-63:**

The “Advanced Setting” dialog closes, and focus returns to the “Connect to Target” dialog.


20. Click **OK**.



**Figure 3-64:**

The “Connect to Target” window closes, and focus returns to the “iSCSI Initiator Properties” window.

21. Notice that the status of the listed discovered target has changed from “Inactive” to “Connected”.

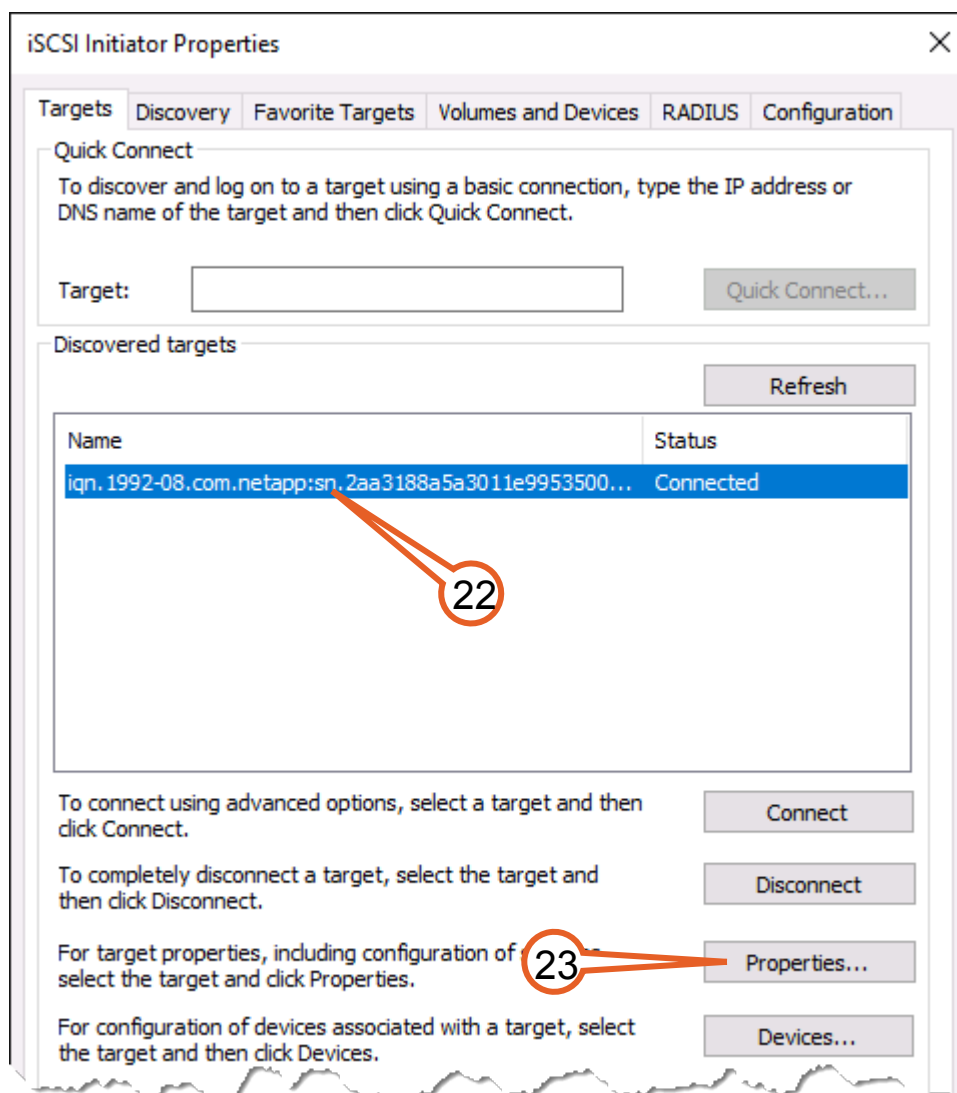


**Figure 3-65:**

Up to this point you have added a single path to your iSCSI LUN. The 192.168.0.133 address you selected corresponds to the cluster1-01\_iscsi\_lif\_1 SAN LIF you created for the “svmluns” SVM in the preceding exercise. You created three additional SAN LIFs during that same exercise, and now you are going to add each of their IP addresses as alternate paths to this same LUN. To begin this procedure you must edit the properties of your existing iSCSI connection.

22. Still on the “Targets” tab, select the discovered target entry for your existing connection.

23. Click **Properties**.



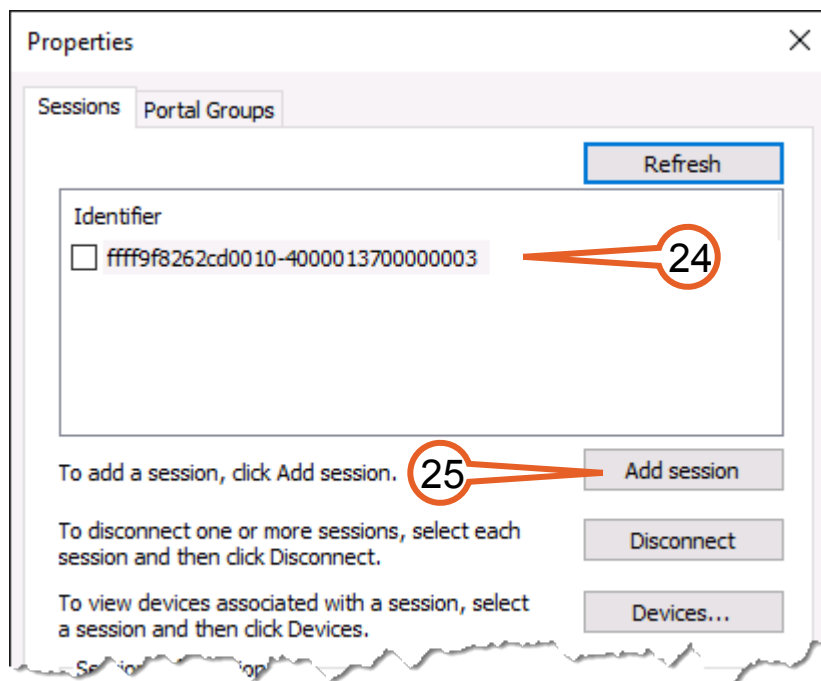
**Figure 3-66:**

The “Properties” dialog opens. From this window you start to connect alternate paths for your newly connected LUN. You will repeat this procedure 3 times, once for each of the remaining LIFs that are present on the svmluns SVM.

LIF IP Address
192.168.0.134
192.168.0.135
192.168.0.136

24. The Identifier list contains an entry for every path you have specified so far, so it can serve as a visual indicator of your progress for specifying all your paths. The first time you enter this window you will see one entry for the LIF you used to first connect to this LUN. When you are done you will see four entries in this window.

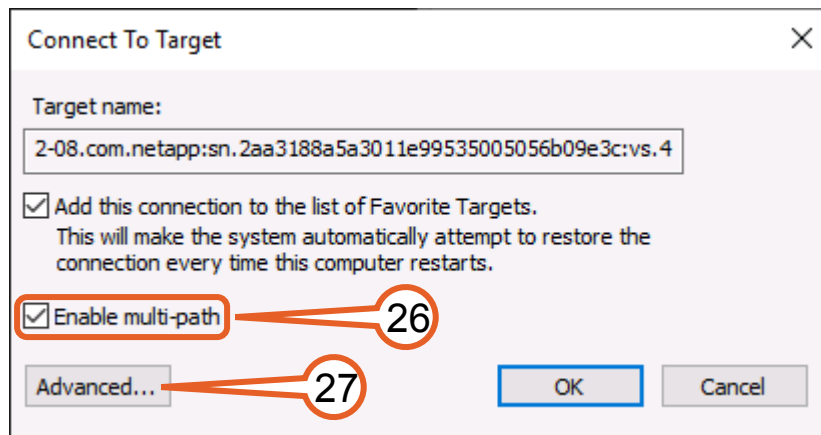
25. Click **Add Session**.



**Figure 3-67:**

The “Connect to Target” dialog opens.

26. Check the **Enable multi-path** checkbox.  
27. Click **Advanced....**



**Figure 3-68:**

The “Advanced Setting” dialog opens.

28. Select the “Target port IP” entry that contains the IP address of the specific LIF you are adding as an alternate path. The following screenshot shows the 192.168.0.134 address, but the value you specify depends of which specific path you are configuring.

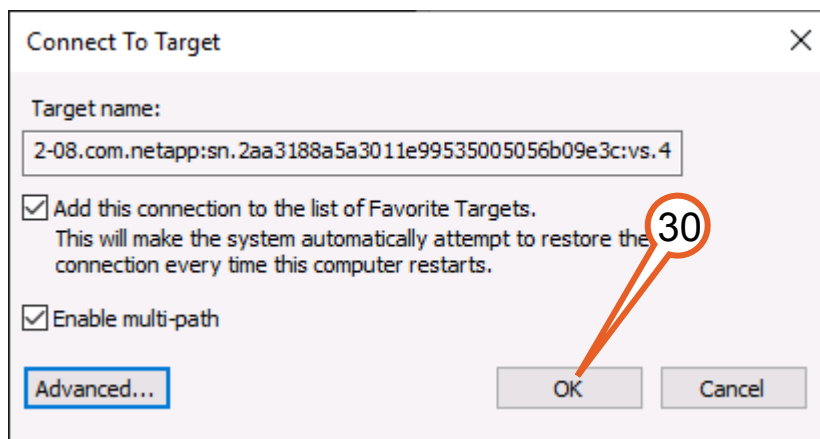
29. When finished, click **OK**.

The screenshot shows the 'Advanced Settings' dialog box with the 'IPsec' tab selected. The 'Connect using' section has three dropdown menus: 'Local adapter' (Default), 'Initiator IP' (Default), and 'Target portal IP' (192.168.0.134 / 3260). The 'Target portal IP' dropdown is highlighted with a red box and a callout '28'. Below this, there are checkboxes for 'CRC / Checksum' (Data integrity), 'Header digest', and 'Enable CHAP log on'. The 'CHAP Log on information' section contains a text box for 'Name' (iqn.1991-05.com.microsoft:jumphost.demo.netapp.com) and a text box for 'Target secret'. There are also checkboxes for 'Perform mutual authentication', 'Use RADIUS to generate user authentication credentials', and 'Use RADIUS to authenticate target credentials'. The 'OK' button is highlighted with a red box and a callout '29'.

**Figure 3-69:**

The “Advanced Settings” dialog closes, and focus returns to the “Connect to Target” dialog.

30. Click **OK**.



**Figure 3-70:**

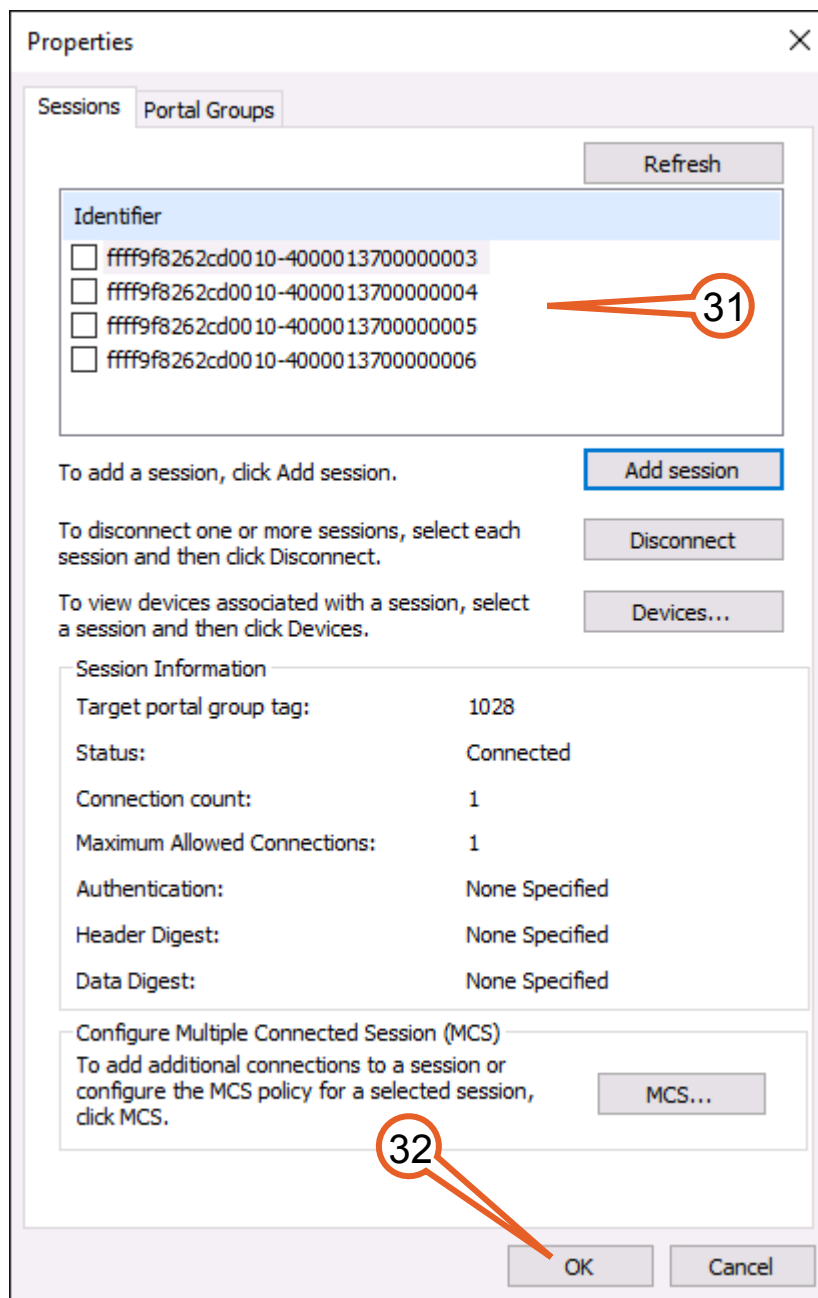
The “Connect to Target” dialog closes, and focus returns to the “Properties” dialog where there are now two entries shown in the identifier list.

Repeat steps 24 - 30 for each of the last two remaining LIF IP addresses. When you have finished adding all the additional paths, the “Identifier” list in the Properties window should contain 4 entries.

31. There are 4 entries in the “Identifier” list when you are finished, indicating that there are 4 sessions, one for each path. Note that it is okay if the identifier values in your lab differ from those in the screenshot.




32. Click **OK**.



**Figure 3-71:**

The “Properties” dialog closes, and focus returns to the “iSCSI Properties” dialog.

33. Click **OK**.




**Figure 3-72:**

The “iSCSI Properties” window closes, and focus returns to the desktop of Jumphost. If the “Administrative Tools” window is not still open on your desktop, open it again now.

The Jumphost should now be connected to the LUN using multi-pathing, so it is time to format your LUN and build a filesystem on it.

34. In “Administrative Tools”, double-click the **Computer Management** tool.



**Figure 3-73:**

The “Computer Management” window opens.

35. In the left pane of the “Computer Management” window, navigate to **Computer Management (Local) > Storage > Disk Management**.
36. In the middle pane, locate the entry for Disk 1 and observe that it is currently offline.



**Tip:** If you see more disks listed than Disk 0 (the boot disk) and Disk 1 (the LUN), MPIO has not correctly recognized that the multiple paths you set up are all for the same LUN. If this occurs, you need to quit **Computer Manager** and go back to the **iSCSI Initiator** tool to review your path configuration steps to find and correct any configuration errors. After that you can return to the Computer Management tool and try again.

37. Right-click on the listing for **Disk 1** and select **Online** from the context menu.

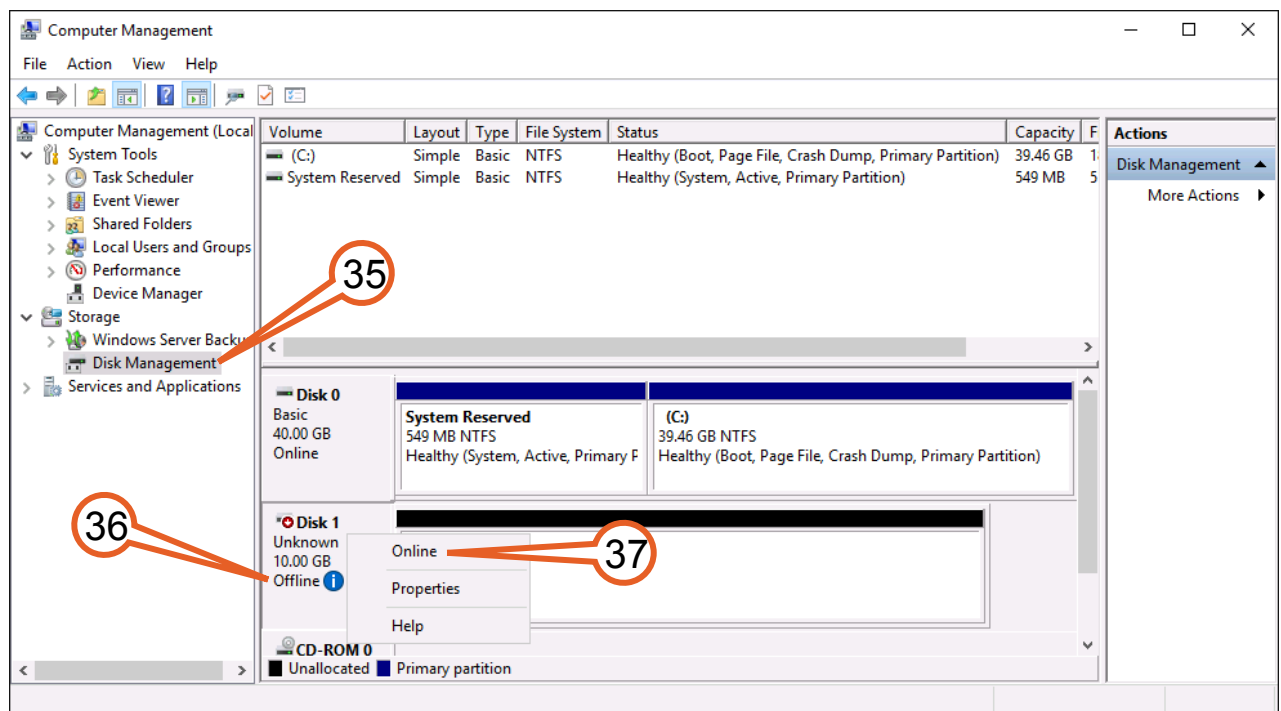


Figure 3-74:

38. The entry for Disk 1 now shows as “Unknown” and “Not Initialized”.
39. Right-click inside the entry for **Disk 1** and select **Initialize Disk** from the context menu.

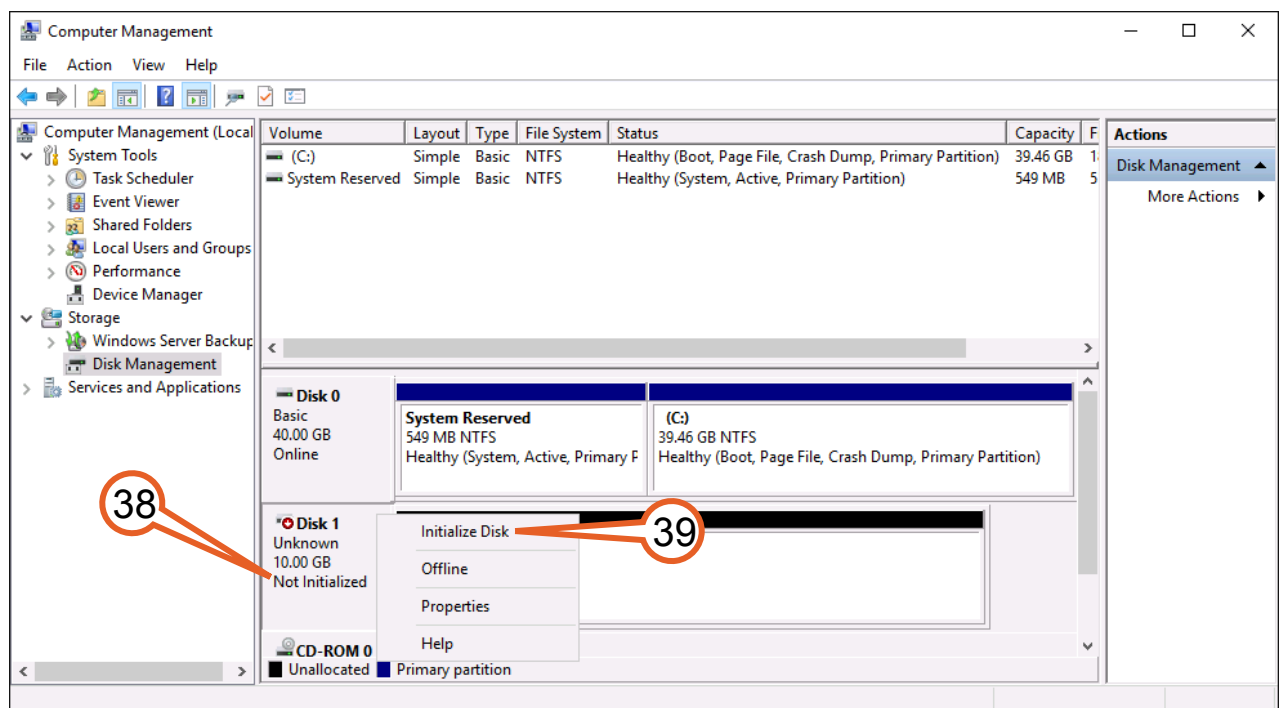



Figure 3-75:

The “Initialize Disk” dialog opens.

40. When you launch Disk Management, an “Initialize Disk” dialog will open informing you that you must initialize a new disk before Logical Disk Manager can access it.

Click **OK** to initialize the disk.




**Figure 3-76:**

The “Initialize Disk” window closes, and focus returns to the “Disk Management” view in the Computer Management window.

41. The disk now reports a type of “Basic”, and a status of “online”.
42. In the partition area, the disk reports as “Unallocated”.

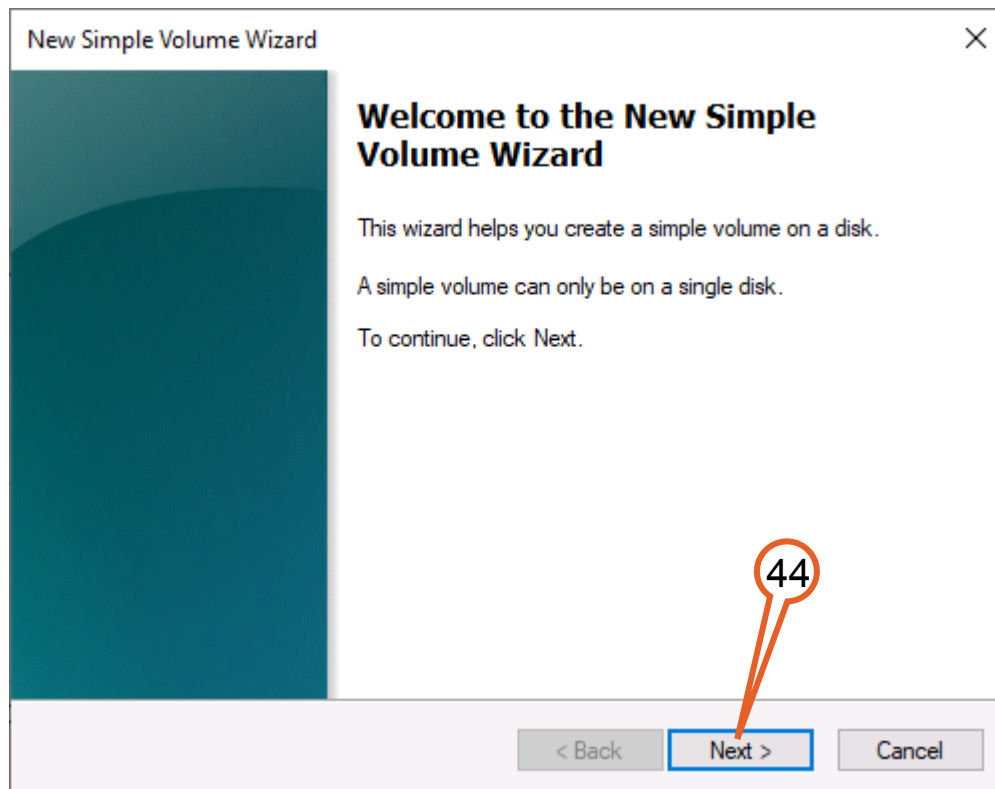
43. Right-click inside the disk's partition area (where it says **Unallocated**) and select **New Simple Volume...** from the context menu.



**Figure 3-77:**

The “New Simple Volume Wizard” window opens.

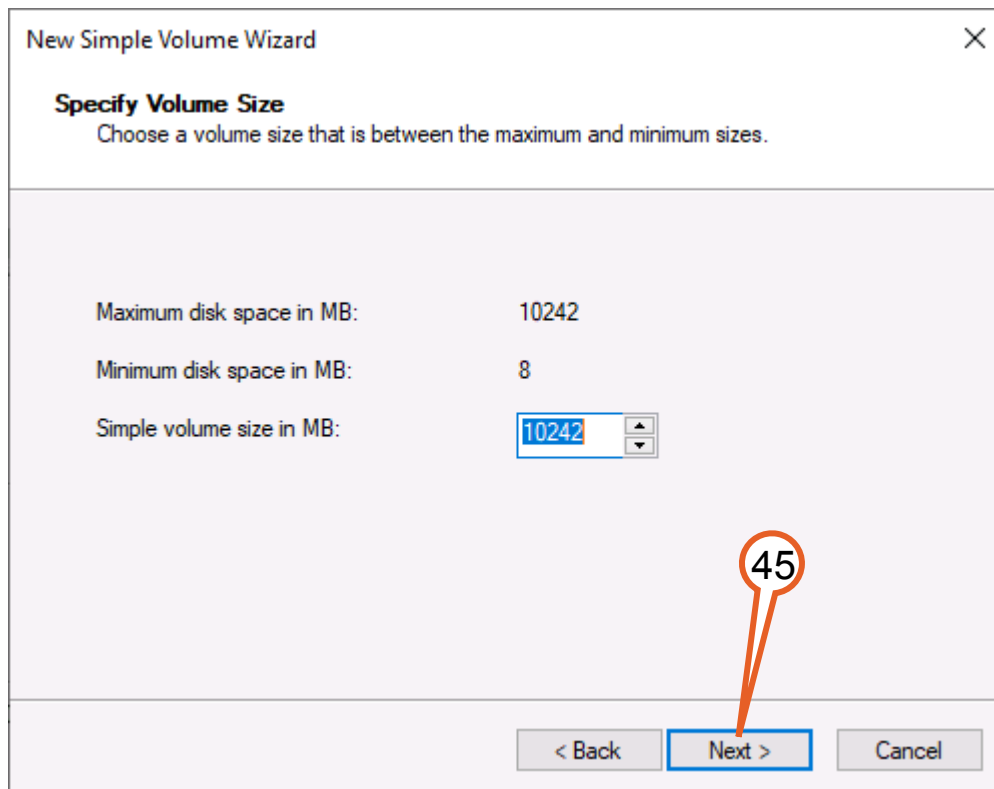
44. Click the **Next** button to advance the wizard.



**Figure 3-78:**

The wizard advances to the “Specify Volume Size” step.

45. The wizard defaults to allocating all of the space in the volume, so click the **Next** button.



The image shows a Windows-style dialog box titled "New Simple Volume Wizard" with a close button (X) in the top right corner. Below the title bar, the section "Specify Volume Size" is displayed, followed by the instruction "Choose a volume size that is between the maximum and minimum sizes." The main area of the dialog contains three rows of information: "Maximum disk space in MB:" with the value "10242", "Minimum disk space in MB:" with the value "8", and "Simple volume size in MB:" with a text box containing "10242" and a spinner control to its right. At the bottom of the dialog, there are three buttons: "< Back", "Next >", and "Cancel". The "Next >" button is highlighted with a blue border. An orange callout bubble with the number "45" inside points to the "Next >" button.

**Figure 3-79:**

The wizard advances to the “Assign Drive Letter or Path” step.

46. The wizard automatically selects the next available drive letter, which should be **E**.



47. Click **Next**.

The screenshot shows the 'New Simple Volume Wizard' window with the title bar 'New Simple Volume Wizard' and a close button. The main heading is 'Assign Drive Letter or Path' with a subtitle 'For easier access, you can assign a drive letter or drive path to your partition.' There are three radio button options: 'Assign the following drive letter:' (selected), 'Mount in the following empty NTFS folder:', and 'Do not assign a drive letter or drive path'. The first option has a dropdown menu showing 'E' with a callout bubble labeled '46' pointing to it. The second option has a text input field and a 'Browse...' button. The third option is unselected. At the bottom, there are three buttons: '< Back', 'Next >' (highlighted with a blue border and a callout bubble labeled '47' pointing to it), and 'Cancel'.

**Figure 3-80:**

The wizard advances to the “Format Partition” step.

48. Set the “Volume Label” field to `WINLUN`.

49. Click **Next**.

New Simple Volume Wizard

**Format Partition**  
To store data on this partition, you must format it first.

Choose whether you want to format this volume, and if so, what settings you want to use.

☐ Do not format this volume

☒ Format this volume with the following settings:

File system: NTFS

Allocation unit size: Default

Volume label: WINLUN

☒ Perform a quick format

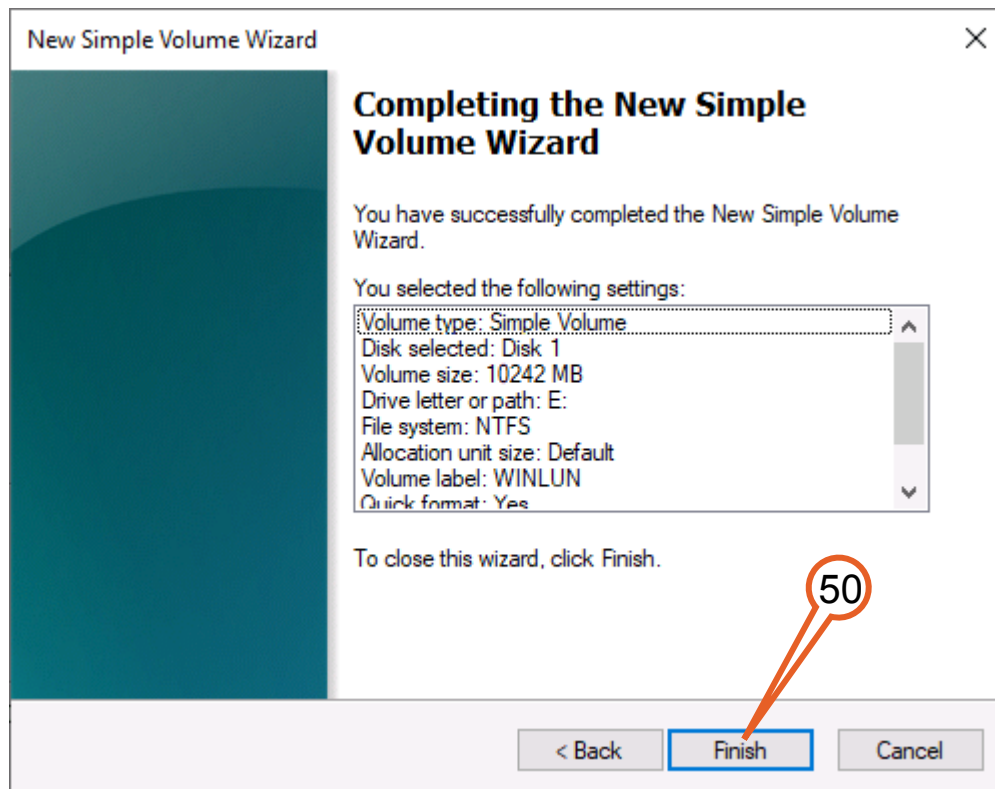
☐ Enable file and folder compression

< Back Next > Cancel

**Figure 3-81:**

The wizard advances to the “Completing the New Simple Volume Wizard” step.

50. Click **Finish**.




**Figure 3-82:**

The “New Simple Volume Wizard” window closes, and focus returns to the “Disk Management” view of the Computer Management window.

51. Windows may open a pop-up dialog (which may appear behind other windows that are already open) indicating that you need to format the new disk in drive E:. This is an erroneous warning due to a race condition in Windows. You already formatted the volume during the New Simple Volume wizard, so click **Cancel** to dismiss the dialog. If you do not see this dialog now, you will likely see it later in step 64.

**Figure 3-83:**



52. The new WINLUN volume displays as “Healthy” in the disk list at the bottom of the Computer Management window, indicating that the new LUN is mounted and ready to use.

53. Before you complete this section of the lab, take a look at the MPIO configuration for this LUN, right-click inside the partition box for the **WINLUN** volume, and select **Properties** from the context menu.





Figure 3-84:

The "WINLUN (E:) Properties" window opens.

54. Click the **Hardware** tab.
55. In the "All disk drives" list, select the **NETAPP LUN C-Mode Multi-Path Disk** entry.

56. Click **Properties**.



**Figure 3-85:**

The "NETAPP LUN C-Mode Multi-Path Disk Device Properties" window opens.

57. Click the **MPIO** tab.
58. Notice that you are using the Microsoft DSM for multi-path access.
59. The MPIO policy is set to "Round Robin with Subset". A number of different multi-pathing policies are available, but the configuration shown here sends LUN I/O down alternating active/optimized paths. You can click the **More information about MPIO policies** link at the bottom of the dialog window for details about all the available policies.
60. The top two paths show both a "Path State" and "TPG State" as "Active/Optimized". These paths are connected to the node cluster1-01, and the Least Queue Depth policy makes active use of both paths to this node. Conversely, the bottom two paths show a "Path State" of "Unavailable", and a "TPG State" of "Active/Unoptimized". These paths are connected to the node cluster1-02, and only enter a Path State of "Active/Optimized" if the node cluster1-01 becomes unavailable, or if the volume hosting the LUN migrates over to the node cluster1-02.

61. When you finish reviewing the information in this dialog, click **OK** to exit. If you changed any of the values in this dialog click **Cancel** to discard those changes.


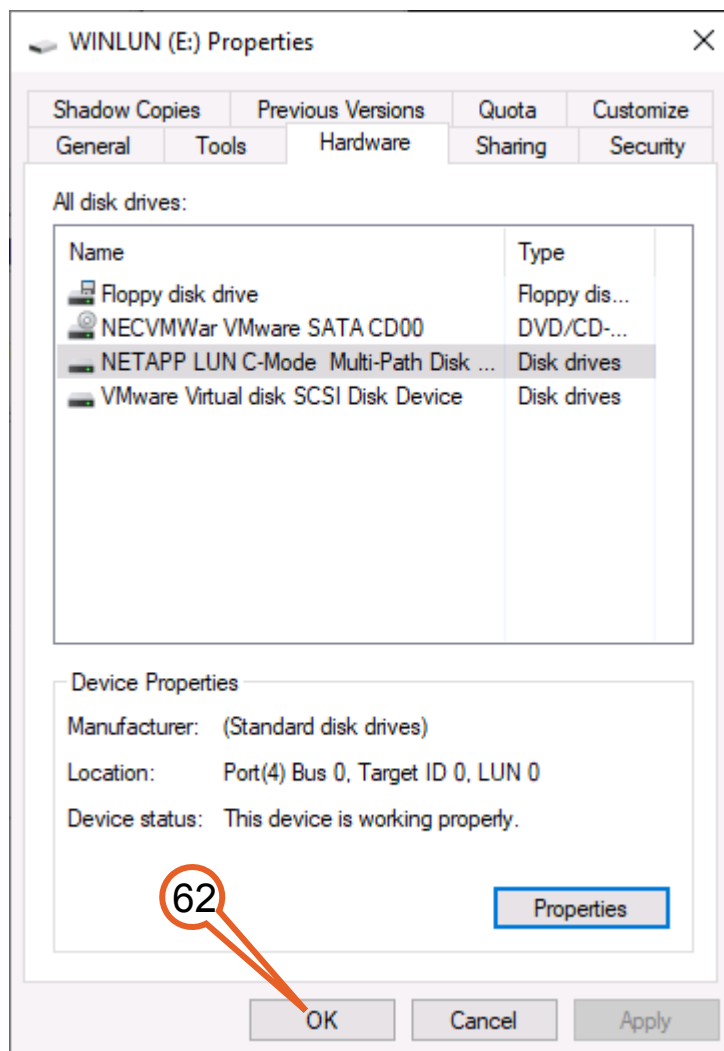


Figure 3-86:

The “NETAPP LUN C-Mode Multi-Path Disk Device Properties” window closes, and focus returns to the “WINLUN (E:) Properties” window.

62. Click **OK**.



**Figure 3-87:**

The "WINLUN (E:) Properties" window closes.

63. Close the **Computer Management** window.

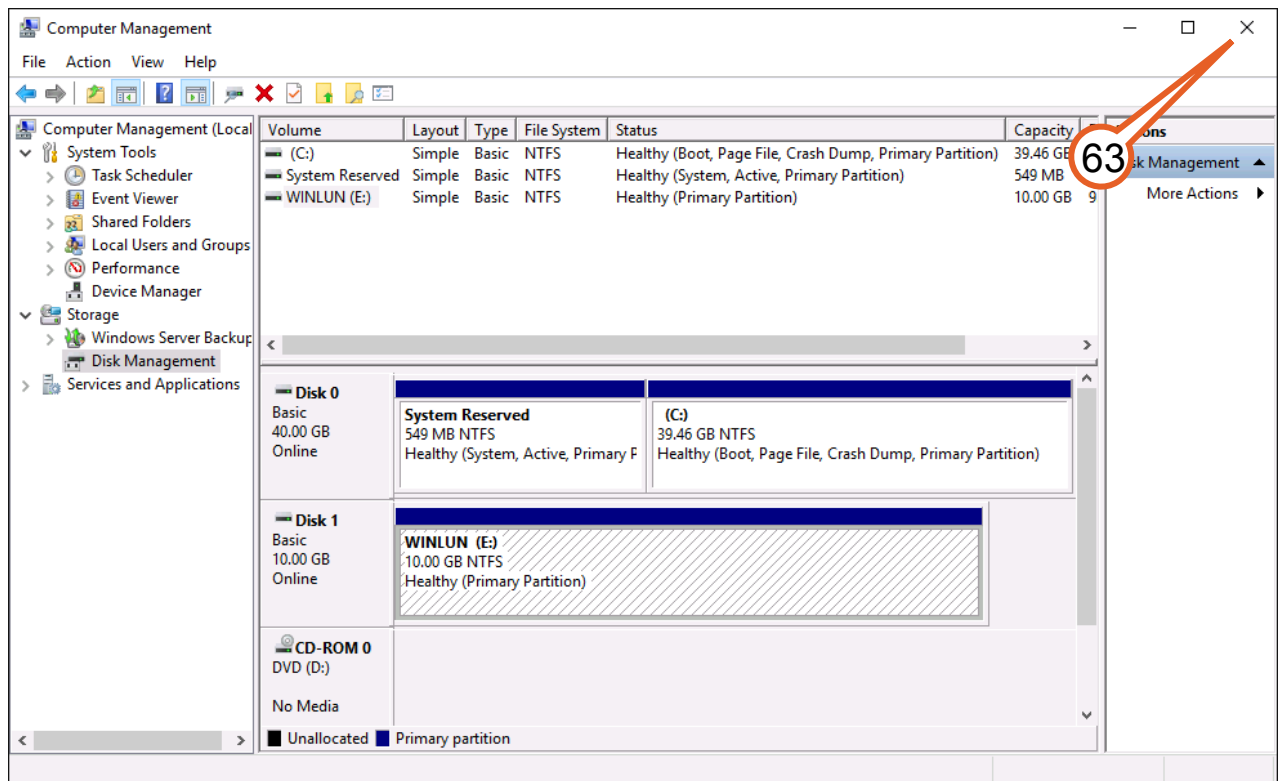


Figure 3-88:

64. If you did not see the pop-up dialog in step 51 from Microsoft Windows stating that you must format the disk in drive E: before you can use it, you may see it now. This window might have been obscured by one of the other windows on the desktop. As you may recall, you did format the LUN during the “New Simple Volume Wizard”, meaning this dialog represents an erroneous disk format message. Click **Cancel** to ignore the format request.



**Note:** Do not close the Administrative tools window yet, as you need to use it again shortly.


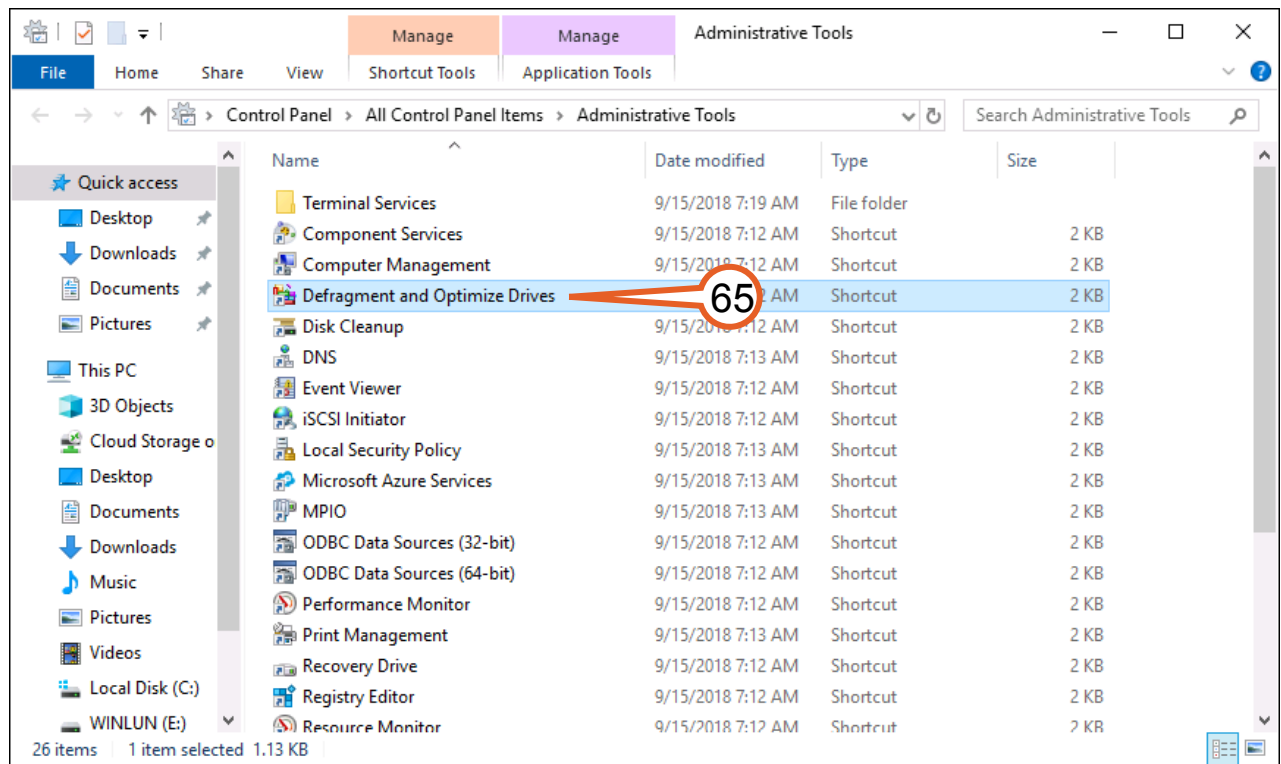


Figure 3-89:

Finally, verify that Windows has detected that the new LUN supports space reclamation. Remember that only Windows 2012 and newer OSs support this feature, and you must minimally have NetApp Windows Unified Host Utilities v6.0.2 (or later) installed. Jumphost meets this criteria.



65. In the “Administrative Tools” window, double-click **Defragment and Optimize drives**.

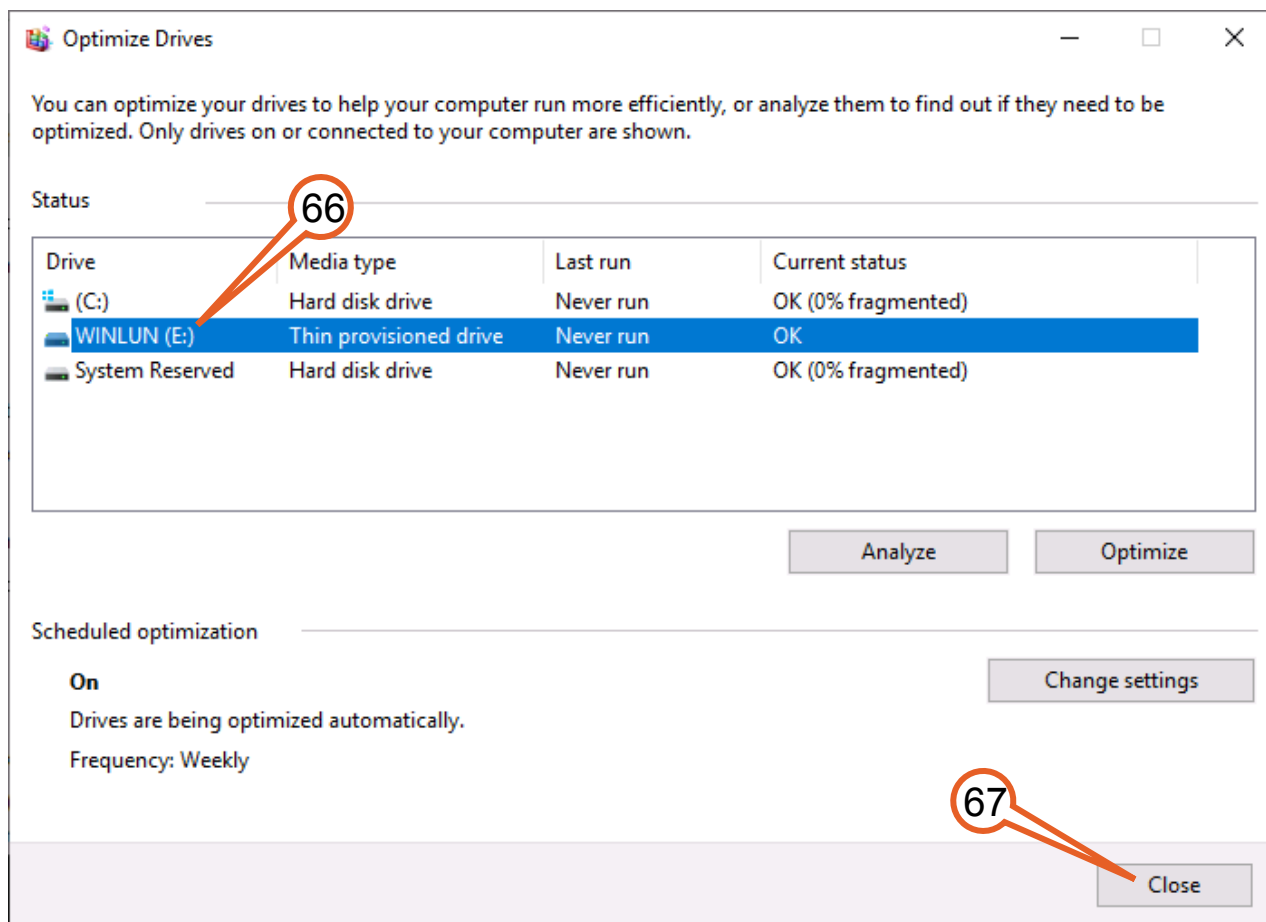


**Figure 3-90:**

The “Optimize Drives” window opens .

66. Find the **WINLUN (E:)** entry in the drive list and look at its “Media type” value. If that value is “Thin provisioned drive”, then Windows has recognized that this drive supports space reclamation. If that value is “Hard disk drive”, then it does not.

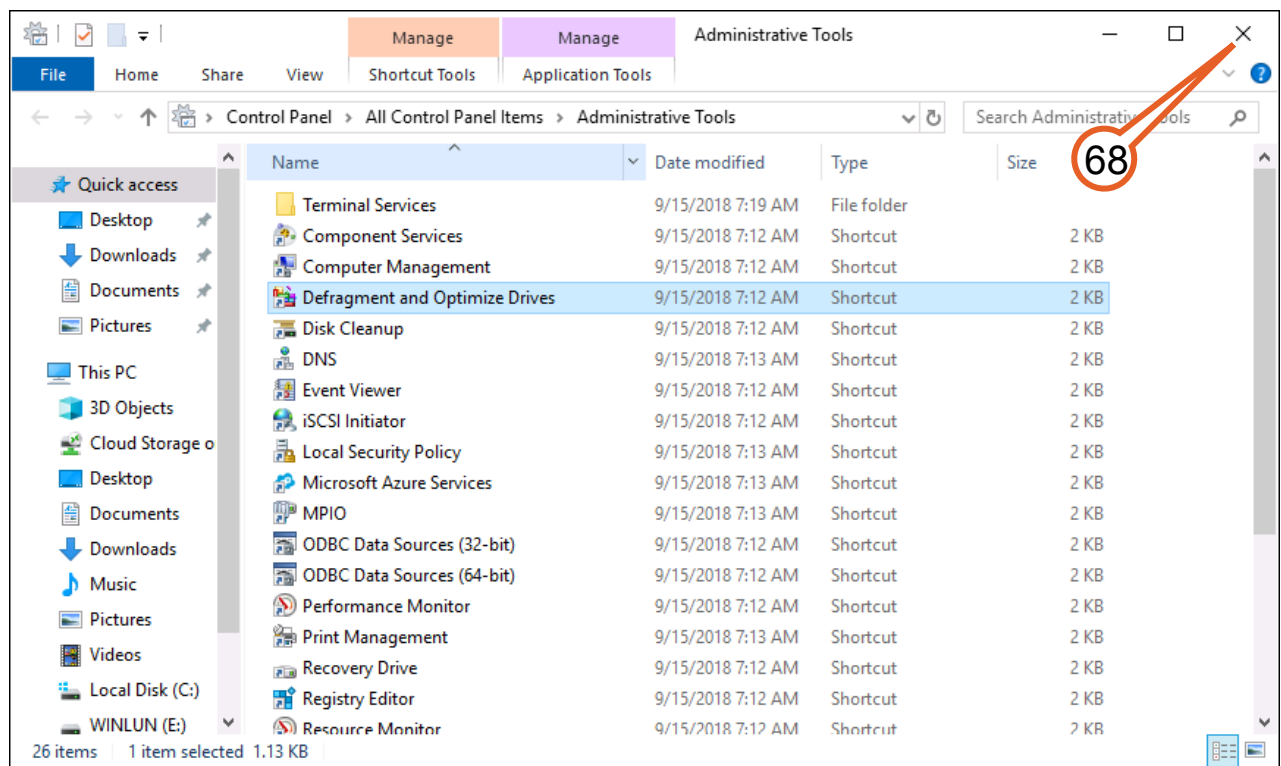
67. Click **Close**.



**Figure 3-91:**

The "Optimize Drives" window closes.

68. Close the **Administrative Tools** window.



**Figure 3-92:**

Feel free to open Windows Explorer on Jumphost, and verify that you can create a file on the E: drive.

This completes this exercise.

### 3.2.3 Create, Map, and Mount a Linux LUN

In an earlier section you created a new SVM, and configured it for iSCSI. In the following sub-sections you will perform the remaining steps needed to configure and use a LUN under Linux:

- Gather the iSCSI Initiator Name of the Linux client.
- Create a thin provisioned Linux volume, create a thin provisioned Linux LUN named “linux.lun” within that volume, and map the LUN to the Linux client.
- Mount the LUN on the Linux client.

The examples in these exercises assume that you have completed the “Create, Map, and Mount a Windows LUN” exercise, but completing that exercise is not a prerequisite for this exercise.

#### 3.2.3.1 LUN Concepts

In ONTAP, LUNs are simply large files that reside within a FlexVol. A single FlexVol can support multiple LUNs, and it is quite common for a client that is using multiple LUNs to have all of those LUNs co-located on the same FlexVol. This facilitates taking application-consistent snapshots, and generally makes LUN management related to the client easier.

While it is possible to configure your SVM so that a volume containing a LUN is also accessible through CIFS or NFS, there is usually no advantage to be gained by doing so. Additionally, if you do not carefully manage export rules and file permissions, you can inadvertently expose the LUN file to unintended NAS clients. While it is perfectly acceptable to serve NAS and SAN clients off of the same SVM, best practice is to not mix NAS files and LUNs within the same volume, or to expose LUN volumes over NAS protocols.

An initiator group, or igroup, defines a list of the Fibre Channel WWPNs or iSCSI node names of the SAN clients that are permitted to access (or even see) a set of LUNs. Mapping a LUN to an igroup applies the igroup's access restrictions to the LUN.

Network access to an iSCSI (or FC) LUN occurs over a LIF. Unlike NAS LIFs, that can migrate online from one node to another in a cluster, SAN LIFs cannot be moved while online. Instead, Asymmetric Logical Unit Access (ALUA) provides redundant paths and automatic path selection as part of any ONTAP SAN solution. When there is an I/O interruption on a given LIF, the client simply (and transparently) retries by switching the I/O to another LIF.

Clients utilize multipathing software in conjunction with ALUA to support multiple paths to a LUN, and for path selection. Windows and Linux both include native multipathing software.

NetApp recommends configuring two SAN LIFs on each node that participate in a SAN configuration in order to provide path redundancy. The optimal path for accessing a LUN is usually to connect directly to the node that is hosting the LUN (i.e., the node that is directly connected to the SSDs/HDDs comprising the aggregate on which the volume hosting the LUN resides). In an HA configuration, that includes the node owning the LUN and its HA partner. LIFs on other nodes can also be used to access the LUN, but those paths may suffer a slight performance penalty as data will have to transfer between the involved nodes over the cluster network.

Selective LUN Mapping (SLM) is an ONTAP feature that reduces the number of paths from the host to the LUN. With SLM, when a new LUN map is created, the LUN is accessible only through paths on the node owning the LUN and its HA partner. This is very helpful in reducing complexity for large clusters (with many nodes) that may present many potential paths to a given LUN.



**Important:** Beginning with ONTAP 9.7, NetApp has started shipping a 2-node AFF All SAN Array option available for factory-ordered and shipped AFF systems. The new AFF All SAN Array (ASA) offers mission-critical workloads (such as databases), a symmetric active-active pathing in a 2-node configuration to offer an even higher level of continuous availability in a simple setup for SAN workloads. While a standard AFF offers a more scalable multi-protocol solution with active-passive (non-optimized) pathing, the AFF ASA SAN system provide a higher level of uninterrupted access for SAN protocols during a planned or unplanned storage failover event.

### 3.2.3.2 Gather the Linux Client iSCSI Initiator Name

You need to determine the Linux client's iSCSI initiator name so that you can set up an appropriate initiator group to control access to the LUN.

You should already have a PuTTY connection open to the Linux host rhel1. If you do not, then open one now using the instructions found in the [“Accessing the Command Line”](#) section at the beginning of this lab guide. The username will be `root`, and the password will be `Netapp11`.

1. Change to the directory that hosts the iscsi configuration files.

```
[root@rhel1 ~]# cd /etc/iscsi
[root@rhel1 iscsi]# ls
initiatorname.iscsi  iscsid.conf
[root@rhel1 iscsi]#
```

2. Display the name of the iscsi initiator.

```
[root@rhel1 iscsi] cat initiatorname.iscsi
InitiatorName=iqn.1994-05.com.redhat:rhel1.demo.netapp.com
[root@rhel1 iscsi]#
```



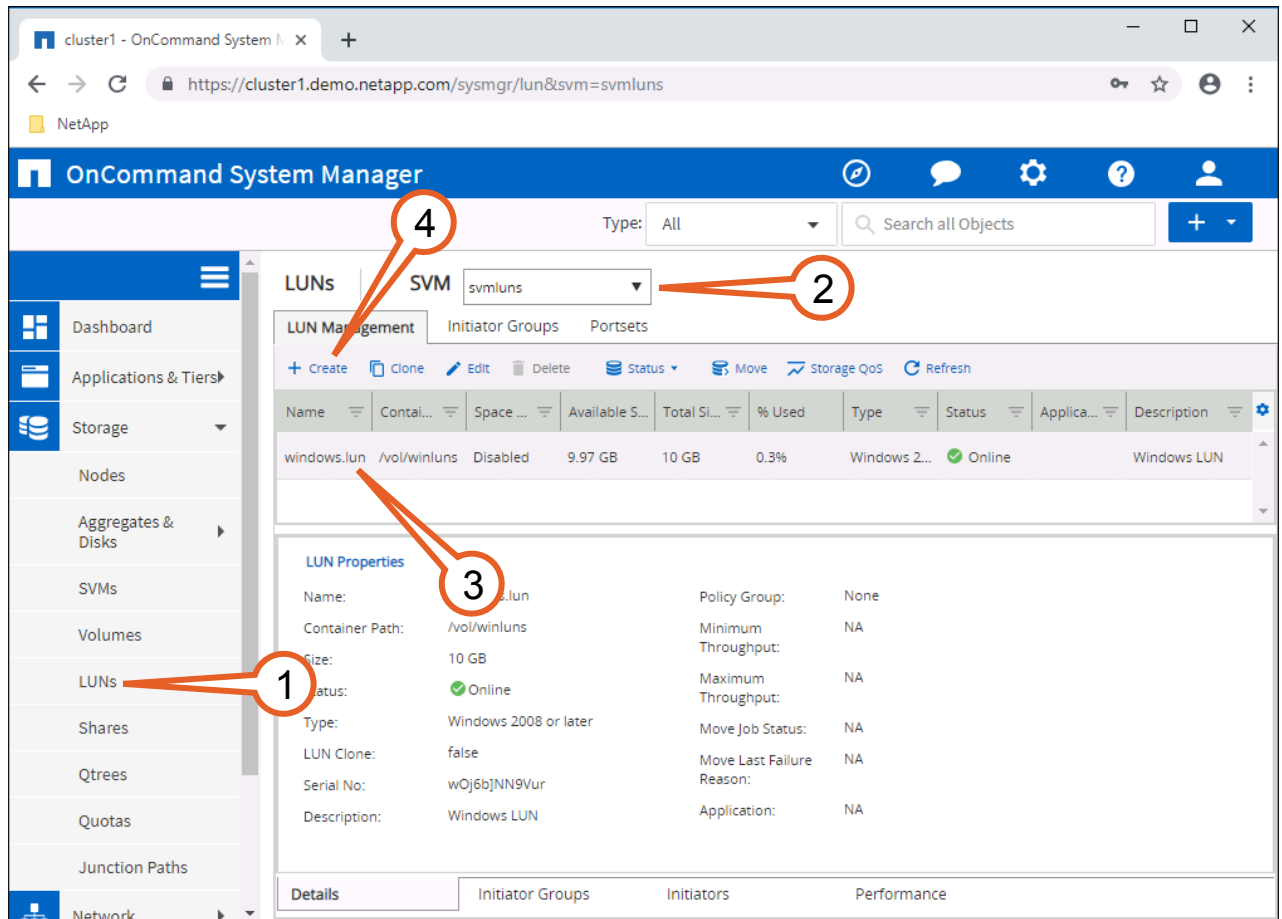
**Important:** The initiator name for rhel1 is `iqn.1994-05.com.redhat:rhel1.demo.netapp.com`.

### 3.2.3.3 Create and Map a Linux LUN

In this activity, you create a new thin provisioned Linux LUN on the SVM “svmluns” under the volume “linluns”, and also create an initiator igroup for the LUN so that only the Linux host rhel1 can access it. An initiator group,

or igroup, defines a list of the Fibre Channel WWPNNs or iSCSI node names for the hosts that are permitted to see the associated LUNs.

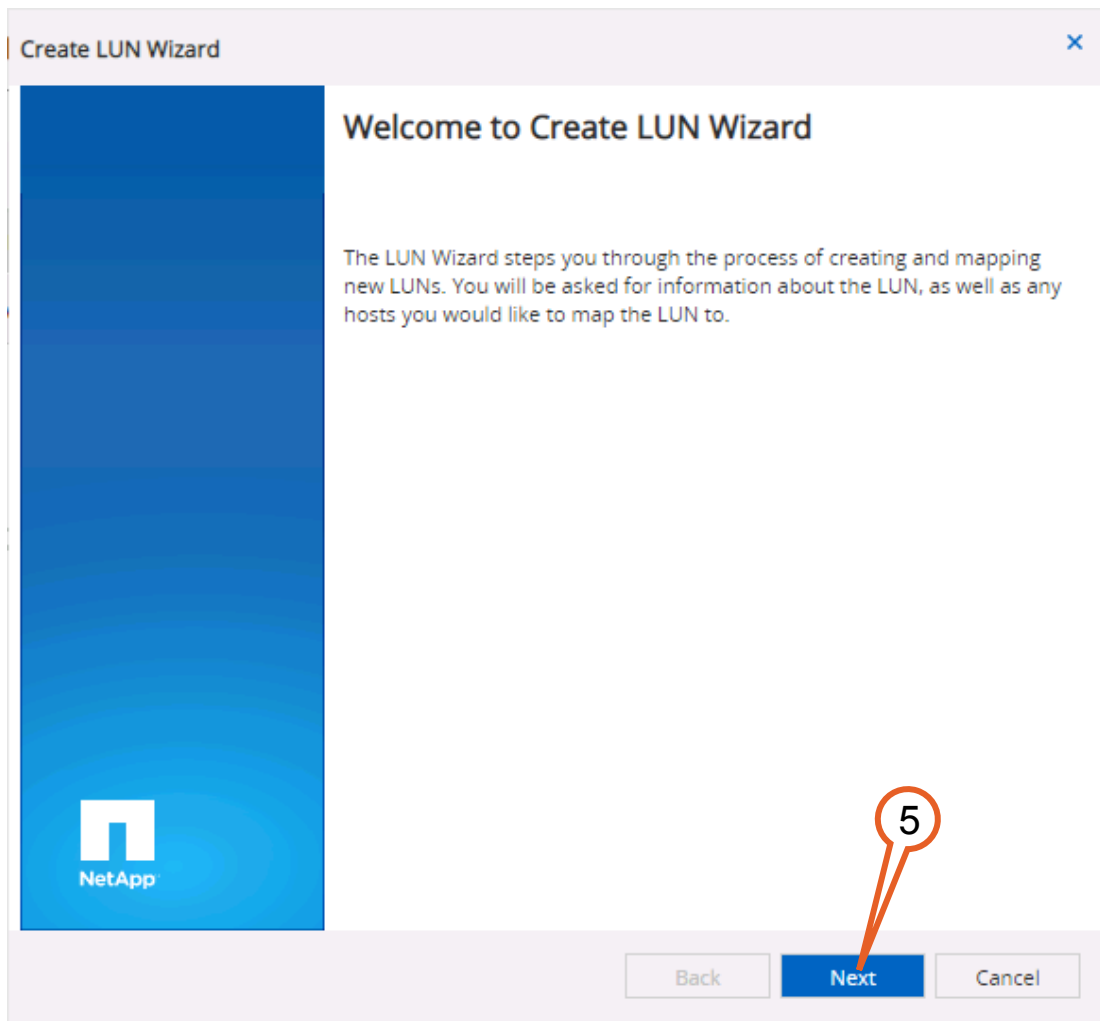
1. In the left pane of System Manager, navigate to **Storage > LUNs**.
2. Make sure the “LUNS on SVM” drop-down is set to **svmluns**.
3. You may or may not see an entry for “windows.lun” in LUN Management list, depending on whether or not you completed the lab sections for creating a Windows LUN.
4. Click **Create**.



**Figure 3-93:**

The “Create LUN Wizard” dialog opens.

5. Click **Next** to advance to the next step in the wizard.



**Figure 3-94:**

The wizard advances to the General Properties step.

6. Set the fields in the window as follows.
  - “Name”: **linux.lun**
  - “Description”: **Linux LUN**
  - “Type”: **Linux**.
  - “Size”: **10 GB**.
  - “Space Reserve”: **Disable**.

7. Click **Next** to continue.

**Create LUN Wizard** [X]

**General Properties**  
You can specify the name, size, type, and an optional description for the LUN that you would like to create.

You can enter a valid name for the LUN and an optional short description

Name:

Description:  (optional)

You can specify the size of the LUN. Storage will be optimized according to the type selected.

Type:  (optional) **6**

[Tell me more about LUN types](#)

Size:   (optional)

Space Reserve:  (optional)

[Tell me more about space reservation](#)

**7**

Back **Next** Cancel

**Figure 3-95:**

The wizard advances to the LUN Container step.

8. Select the radio button to **Create a new flexible volume**, and set the fields under that heading as follows.

- “Aggregate Name”: `aggr1_cluster1_01`
- “Volume Name”: `linluns`
- “Tiering Policy”: **none**.



**Note:** Tiering Policy is used with FabricPool, coverage of which lies outside the scope of this lab.

9. When finished click **Next**.

Create LUN Wizard

**LUN Container**  
You can let the wizard create a volume or you can choose an existing volume as the LUN container.

The wizard automatically chooses the aggregate with most free space for creating flexible volume for the LUN. But you can choose a different aggregate of your choice. You can also select an existing volume/qtree to create your LUN.

☐ Select an existing volume or qtree for this LUN

Volume/Qtree:

☒ Create a new flexible volume in

Aggregate Name:

Volume Name:

Tiering Policy:

[Tell me more about cloud tier and tiering policies.](#)

**Figure 3-96:**

The wizard advances to the Initiator Mapping step.



10. Click **Add Initiator Group**.

Create LUN Wizard

**Initiators Mapping**  
You can connect your LUN to the initiator hosts by selecting from the initiator group and by optionally providing LUN ID for the initiator group.

Map	Initiator Group Name	Type	LUN ID (Optional)
-----	----------------------	------	-------------------

☐ Show All Initiator Groups

Add Initiator Group

Back Next Cancel

**Figure 3-97:**

The “Create Initiator Group” window opens.

11. Set the fields in the window as follows.

- “Name”: `linigrp`
- “Operating System”: **Linux**.
- “Type”: Click the **iSCSI** radio button.

12. Click the **Initiators** tab.

The screenshot shows the 'Create Initiator Group' dialog box with the 'Initiators' tab selected. The 'Name' field contains 'linigrp'. The 'Operating System' dropdown menu is set to 'Linux'. Under the 'Type' section, the 'iSCSI' radio button is selected. The 'Portset' section includes a text box and a 'Choose' button. Red callout boxes with numbers 11 and 12 point to the 'Name' field, the 'Operating System' dropdown, and the 'iSCSI' radio button respectively.

**Figure 3-98:**

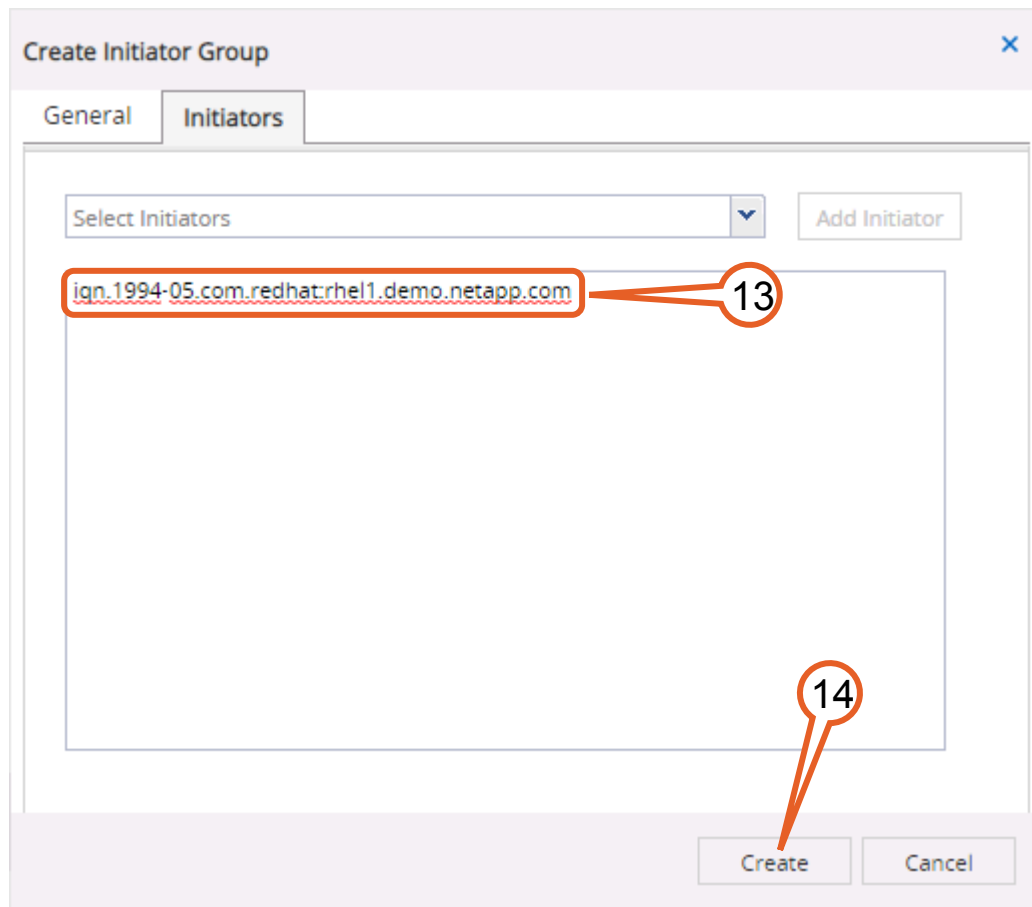
The Initiators tab displays.

13. Click the first line inside the large text box in the window (the box below the dropdown displaying “Select Initiators”). Populate the entry with the value of the iSCSI Initiator name for rhel1 that you saved earlier. In case you misplaced that value, it was:



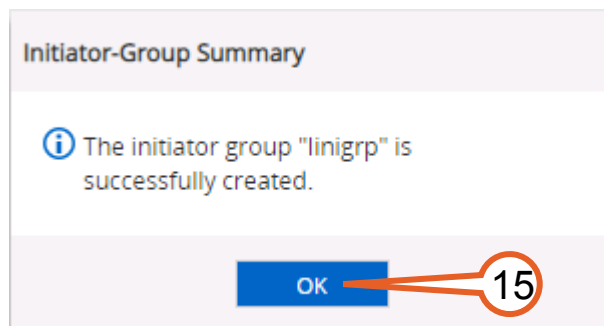
**Important:** The iSCSI Initiator name is **iqn.1994-05.com.redhat:rhel1.demo.netapp.com**.

14. When you finish entering the value, click **Create**.



**Figure 3-99:**

15. If an “Initiator-Group Summary” window opens, click **OK** to acknowledge the confirmation, otherwise proceed to the next step.



**Figure 3-100:**

The “Initiator-Group Summary” window closes, and focus returns to the “Initiators Mapping” step of the Create LUN wizard.

16. Click the checkbox under the “Map” column next to the **linigrp** initiator group.

**Caution:** This is a critical step because this is where you actually map the new LUN to the new igrp.

17. Click **Next** to continue.

Create LUN Wizard

**Initiators Mapping**  
You can connect your LUN to the initiator hosts by selecting from the initiator group and by optionally providing LUN ID for the initiator group.

Map ▾	Initiator Group Name	Type	LUN ID (Optional)
<input checked="" type="checkbox"/>	linigrp	Linux	

☐ Show All Initiator Groups      Add Initiator Group

Back    **Next**    Cancel

**Figure 3-101:**

The wizard advances to the Storage Quality of Service Properties step. You will not create any QoS policies in this lab. If you are interested in learning about QoS, please see the *Hands-on Lab for Advanced Concepts for NetApp ONTAP* lab.

18. Click **Next** to continue.

Create LUN Wizard

Storage Quality of Service Properties  
Limit LUN throughput by assigning it to a Quality of Service policy group

☐ Manage Storage Quality of Service

Apply QoS policy to the LUN by assigning it to a policy group and specify the QoS maximum throughput and QoS minimum throughput values. Storage objects assigned to the same QoS policy will share the same QoS maximum throughput value.

[Tell me more about Storage Quality of Service](#)

Assign to: ☒ New Policy Group ☐ Existing Policy Group

Policy Group Name:

Minimum Throughput:  (IOPS) [i](#)

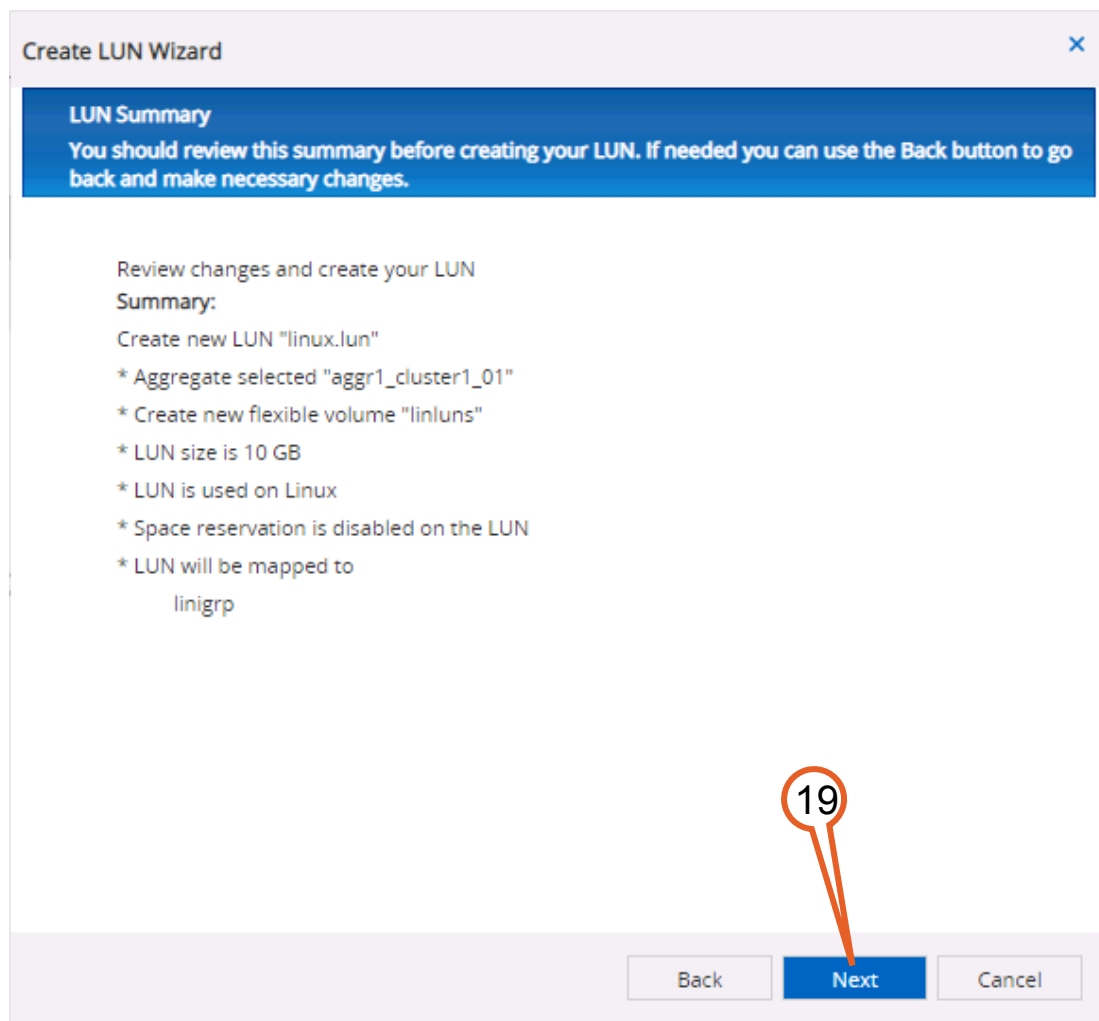
Maximum Throughput:  (MB/s)  (IOPS)

Back Next Cancel

**Figure 3-102:**

The wizard advances to the LUN Summary step, where you can review your selections before proceeding to create the LUN.

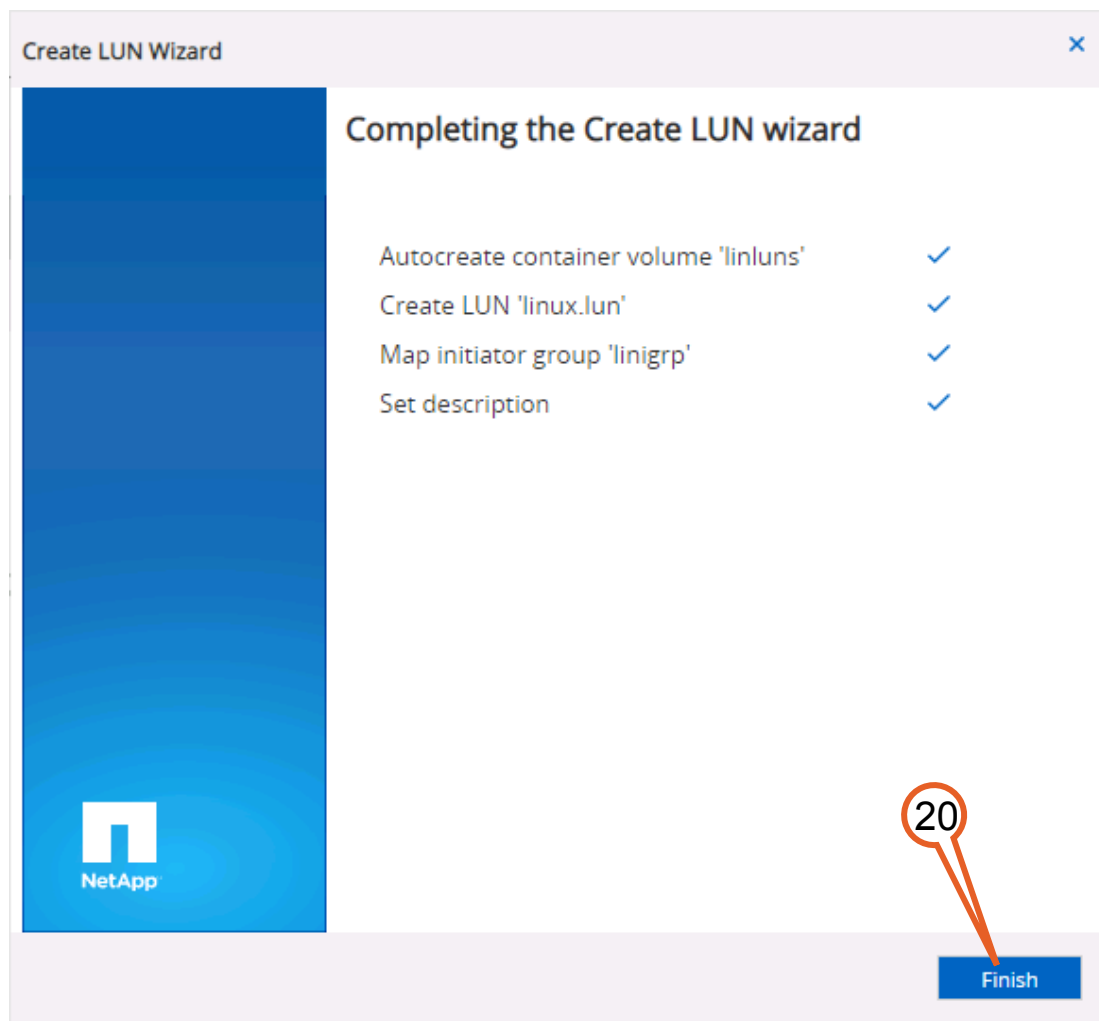
19. If everything looks correct, click **Next**.



**Figure 3-103:**

The wizard begins creating the volume that will contain the LUN, creating the LUN, and mapping the LUN to the new igroup. As it finishes each step the wizard displays a green check mark in the window next to that step.

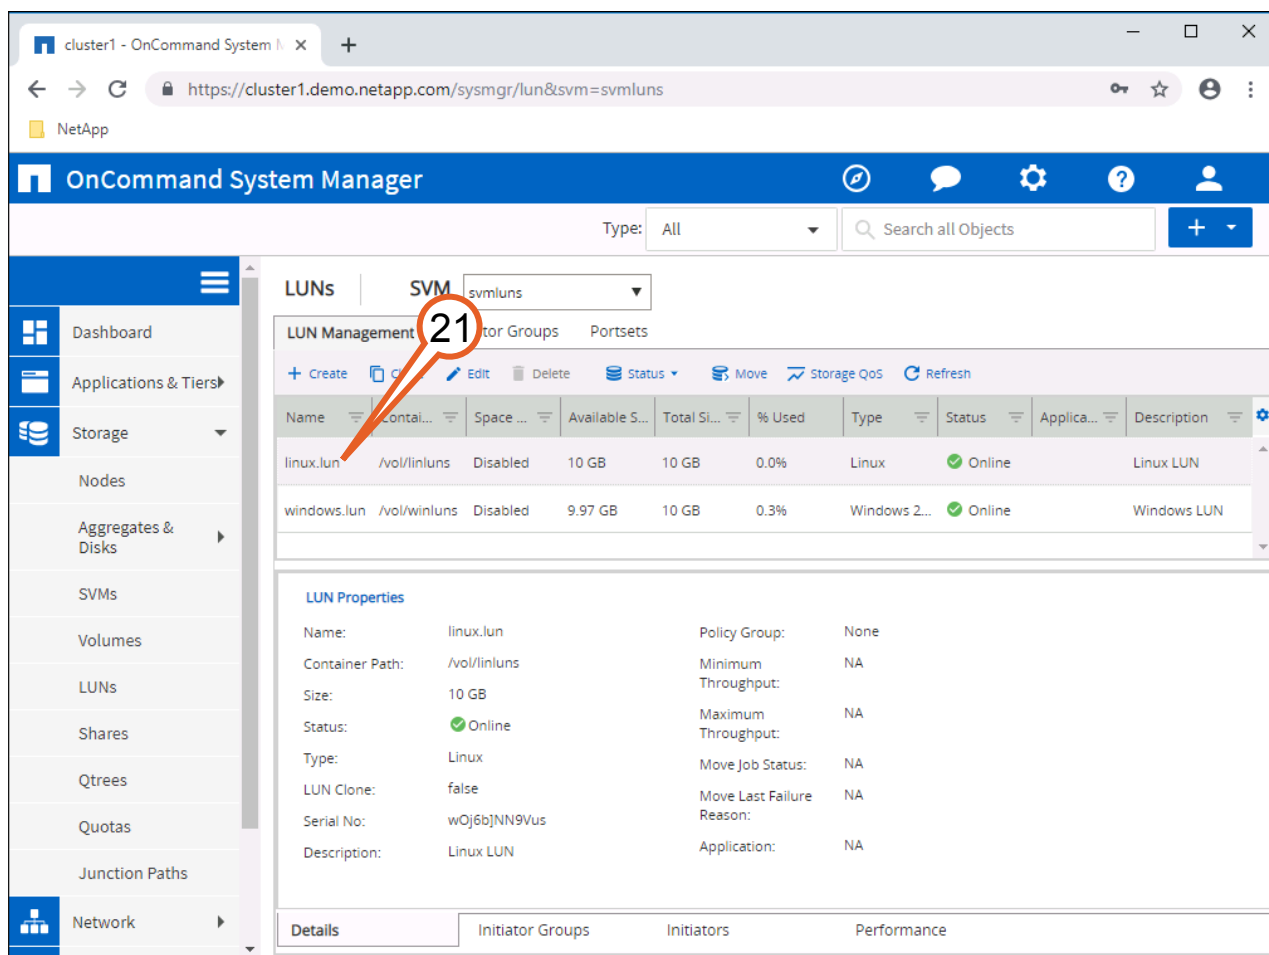
20. Click **Finish** to terminate the wizard.



**Figure 3-104:**

The “Create LUN wizard” window closes, and focus returns to the LUNs view in System Manager.

21. The new LUN “linux.lun” now appears in the LUNs view, and if you select it you can review its details in the bottom pane.



**Figure 3-105:**

The new Linux LUN now exists, and is mapped to your rhel1 client.

ONTAP includes support for LUN space reclamation, an optional feature that enables ONTAP to reclaim space from a thin provisioned LUN when the client deletes data from it, and which also allows ONTAP to notify the client when the LUN cannot accept writes due to lack of space on the volume. This feature is supported by VMware ESX 5.0 and later, Red Hat Enterprise Linux 6.2 and later, and Microsoft Windows 2012 and later. The RHEL clients used in this lab are running version 7.5, so you will enable the space reclamation feature for your Linux LUN. You can only enable space reclamation through the ONTAP command line.

22. If you do not already have a PuTTY session open to `cluster1`, open one now by following the instructions in the “[Accessing the Command Line](#)” section at the beginning of this guide. Log in with username `admin`, and the password `Netapp1!`.
23. In the cluster1 CLI, view whether space reclamation is enabled for the LUN.

```
cluster1::> lun show -vserver svmluns -path /vol/linluns/linux.lun
-fields space-allocation
vserver path space-allocation
-----
svmluns /vol/linluns/linux.lun disabled
cluster1::>
```

24. Enable space reclamation for the LUN linux.lun.

```
cluster1::> lun modify -vserver svmluns -path /vol/linluns/linux.lun
```



```
-space-allocation enabled
cluster1::>
```

25. View the LUN's space reclamation setting again.

```
cluster1::> lun show -vserver svmluns -path /vol/linluns/linux.lun
-fields space-allocation
vserver path space-allocation
-----
svmluns /vol/linluns/linux.lun enabled
cluster1::>
```

### 3.2.3.4 Mount the LUN on a Linux Client

In this activity you use the Linux command line to configure the host rhel1 to connect to the Linux LUN /vol/linluns/linux.lun you created in the preceding section.

This section assumes that you know how to use the Linux command line. If you are not familiar with these concepts, we recommend that you skip this section of the lab.

1. If you do not currently have a PuTTY session open to rhel1, open one now and log in as user `root` with the password "**Netapp1!**".
2. The NetApp Linux Unified Host Utilities kit has been pre-installed on both Red Hat Linux hosts in this lab, and the iSCSI initiator name has already been configured for each host. Confirm that is the case:

```
[root@rhel1 ~]# cd
[root@rhel1 ~]# rpm -qa | grep netapp
netapp_linux_unified_host_utilities-7-1.x86_64
[root@rhel1 ~]# cat /etc/iscsi/initiatorname.iscsi
InitiatorName=ign.1994-05.com.redhat:rhel1.demo.netapp.com
[root@rhel1 ~]#
```

3. In the `/etc/iscsi/iscsid.conf` file the `node.session.timeo.replacement_timeout` value is set to **5** to better support timely path failover, and the `node.startup` value is set to **automatic** so that the system will automatically log in to the iSCSI node at startup.

```
[root@rhel1 ~]# grep replacement_time /etc/iscsi/iscsid.conf
#node.session.timeo.replacement_timeout = 120
node.session.timeo.replacement_timeout = 5
[root@rhel1 ~]# grep node.startup /etc/iscsi/iscsid.conf
# node.startup = automatic
node.startup = automatic
[root@rhel1 ~]#
```

4. You will find that the Red Hat Linux hosts in the lab have pre-installed DM-Multipath packages and a `/etc/multipath.conf` file pre-configured to support multi-pathing, so that the RHEL host can access the LUN using all of the SAN LIFs you created for the svmluns SVM.

```
[root@rhel1 ~]# rpm -q device-mapper
device-mapper-1.02.146-4.el7.x86_64
[root@rhel1 ~]# rpm -q device-mapper-multipath
device-mapper-multipath-0.4.9-119.el7.x86_64
[root@rhel1 ~]# cat /etc/multipath.conf
# For a complete list of the default configuration values, see
# /usr/share/doc/device-mapper-multipath-0.4.9/multipath.conf.defaults
# For a list of configuration options with descriptions, see
# /usr/share/doc/device-mapper-multipath-0.4.9/multipath.conf.annotated
#
# REMEMBER: After updating multipath.conf, you must run
#
# service multipathd reload
#
# for the changes to take effect in multipathd

# NetApp recommended defaults
defaults {
    flush_on_last_del yes
    max_fds max
    queue_without_daemon no
    user_friendly_names no
    dev_loss_tmo infinity
    fast_io_fail_tmo 5
}
```

```

}

blacklist {
  devnode "^sda"
  devnode "^hd[a-z]"
  devnode "^(ram|raw|loop|fd|md|dm-|sr|scd|st)[0-9]*"
  devnode "^ccis.*"
}

devices {
  # NetApp iSCSI LUNs
  device {
    vendor          "NETAPP"
    product         "LUN"
    path_grouping_policy group_by_prio
    features        "3 queue_if_no_path pg_init_retries 50"
    prio           "alua"
    path_checker    tur
    failback        immediate
    path_selector   "round-robin 0"
    hardware_handler "1 alua"
    rr_weight       uniform
    rr_min_io       128
    getuid_callout  "/lib/udev/scsi_id -g -u -d /dev/%n"
  }
}
[root@rhell ~]#

```

5. You now need to start the iSCSI software service on rhel1, and configure it to start automatically at boot time. Note that a force-start is only necessary the very first time you start the iscsid service on host.

```

[root@rhell ~]# systemctl status iscsid.service
# iscsid.service - Open-iSCSI
   Loaded: loaded (/usr/lib/systemd/system/iscsid.service; disabled; vendor preset: disabled)
   Active: inactive (dead)
     Docs: man:iscsid(8)
           man:iscsiadm(8)
[root@rhell ~]# systemctl start iscsid
[root@rhell ~]# systemctl enable iscsid
Created symlink from /etc/systemd/system/multi-user.target.wants/iscsid.service to /usr/lib/systemd/system/iscsid.service.
[root@rhell ~]# systemctl status iscsid.service
#m iscsid.service - Open-iSCSI
   Loaded: loaded (/usr/lib/systemd/system/iscsid.service; enabled; vendor preset: disabled)
   Active: active (running) since Apr 2019-04-09 17:46:56 UTC; 12s ago
     Docs: man:iscsid(8)
           man:iscsiadm(8)
  Main PID: 30908 (iscsid)
    CGroup: /system.slice/iscsid.service
            ##30907 /usr/sbin/iscsid
            ##30908 /usr/sbin/iscsid

Apr 09 17:46:56 rhell systemd[1]: Starting Open-iSCSI...
Apr 09 17:46:56 rhell iscsid[30906]:
Apr 09 17:46:56 rhell systemd[1]: Failed to read PID from file /var/run/isc...nt
Apr 09 17:46:56 rhell systemd[1]: Started Open-iSCSI.
Apr 09 17:46:56 rhell iscsid[30907]: iSCSI daemon with pid=30908 started!
Hint: Some lines were ellipsized, use -l to show in full.
[root@rhell ~]#

```

6. Next discover the available targets using the `iscsiadm` command. Note that the exact values used for the node paths may differ in your lab from what is shown in this example, and that after running this command there will still not be active iSCSI sessions because you have not yet created the necessary device files.

```

[root@rhell ~]# iscsiadm --mode discovery --op update --type sendtargets
--portal 192.168.0.133
192.168.0.133:3260,1028 iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5
192.168.0.136:3260,1031 iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5
192.168.0.135:3260,1030 iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5
192.168.0.134:3260,1029 iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5
[root@rhell ~]# iscsiadm --mode session
iscsiadm: No active sessions.
[root@rhell ~]#

```

7. Create the devices necessary to support the discovered nodes, after which the sessions become active.

```
[root@rhell ~]# iscsiadm --mode node -l all
Logging in to [iface: default, target:
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5, portal: 192.168.0.133,3260]
(multiple)
Logging in to [iface: default, target:
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5, portal: 192.168.0.136,3260]
(multiple)
Logging in to [iface: default, target:
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5, portal: 192.168.0.135,3260]
(multiple)
Logging in to [iface: default, target:
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5, portal: 192.168.0.134,3260]
(multiple)
Login to [iface: default, target:
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5, portal: 192.168.0.133,3260]
successful.
Login to [iface: default, target:
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5, portal: 192.168.0.136,3260]
successful.
Login to [iface: default, target:
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5, portal: 192.168.0.135,3260]
successful.
Login to [iface: default, target:
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5, portal: 192.168.0.134,3260]
successful.
[root@rhell ~]# iscsiadm --mode session
tcp: [1] 192.168.0.133:3260,1028
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5 (non-flash)
tcp: [2] 192.168.0.136:3260,1031
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5 (non-flash)
tcp: [3] 192.168.0.135:3260,1030
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5 (non-flash)
tcp: [4] 192.168.0.134:3260,1029
iqn.1992-08.com.netapp:sn.85f24836beaf11e8ba95005056b048e1:vs.5 (non-flash)
[root@rhell ~]#
```

8. At this point the Linux client sees the LUN over all four paths, but it does not yet understand that all four paths represent the same LUN.

```
[root@rhell ~]# sanlun lun show
controller(7mode/E-Series)/
lun
vserver(cDOT/FlashRay)      lun-pathname      device      host
protocol  size  product      filename      adapter
-----
svmluns          /vol/linluns/linux.lun      /dev/sdd      host5      iSCSI
10g      cDOT
svmluns          /vol/linluns/linux.lun      /dev/sde      host6      iSCSI
10g      cDOT
svmluns          /vol/linluns/linux.lun      /dev/sdc      host4      iSCSI
10g      cDOT
svmluns          /vol/linluns/linux.lun      /dev/sdb      host3      iSCSI
10g      cDOT
[root@rhell ~]#
```

9. Verify that the multipathd service is enabled and running.

```
[root@rhell ~]# systemctl status multipathd
# multipathd.service - Device-Mapper Multipath Device Controller
Loaded: loaded (/usr/lib/systemd/system/multipathd.service; enabled; vendor preset:
enabled)
Active: active (running) since Tue 2020-12-15 06:008:24 UTC; 2 days ago
Main PID: 535 (multipathd)
CGroup: /system.slice/multipathd.service
#535 /sbin/multipathd

Dec 18 00:26:23 rhell multipathd[535]: sde: add path (uevent)
Dec 18 00:26:23 rhell multipathd[535]: 3600a0980774f6a36625d4d4f7775564d: load table [0
20971520 multipath 4 qu...:64 1]
Dec 18 00:26:23 rhell multipathd[535]: 3600a0980774f6a36625d4d4f7775564d: event checker
started
Dec 18 00:26:23 rhell multipathd[535]: sde [8:64]: path added to devmap
3600a0980774f6a36625d4d4f7775564dDec 18 00:26:23 rhell multipathd[535]: sdb: add path
(uevent)
Dec 18 00:26:23 rhell multipathd[535]: sdc: add path (uevent)
Dec 18 00:26:23 rhell multipathd[535]: sdd: add path (uevent)
```

```
Dec 18 00:26:23 rhell multipathd[535]: 3600a0980774f6a36625d4d4f7775564d: load table [0
20971520 multipath 4 qu...:48 1]
Warning: Journal has been rotated since unit was started. Log output is incomplete or
unavailable.
Hint: Some lines were ellipsized, use -l to show in full.
[root@rhell ~]#
```



**Tip:** If multipathd was not running, then you could enable and start it by issuing the `systemctl enable multipathd` and `systemctl start multipathd` commands.

- The `multipath` command displays the configuration of DM-Multipath, and the `multipath -ll` command displays a list of the multipath devices. DM-Multipath maintains a device file under `/dev/mapper` that you use to access the multipathed LUN (in order to create a filesystem on it and to mount it). The first line of output from the `multipath -ll` command lists the name of that device file (in this example “3600a0980774f6a36625d4d4f7775564d”). The autogenerated name for this device file will likely differ in your copy of the lab. Also pay attention to the output of the `sanlun lun show -p` command which shows information about the ONTAP path of the LUN, the LUN's size, its device file name under `/dev/mapper`, the multipath policy, and also information about the various device paths themselves.

```
[root@rhell ~]# multipath -ll
Dec 18 00:26:45 | /etc/multipath.conf line 43, invalid keyword: getuid_callout
3600a0980774f6a36625d4d4f7775564d dm-3 NETAPP ,LUN C-Mode
size=10G features='4 queue_if_no_path pg_init_retries 50 retain_attached_hw_handle'
hw_handler='1 alua' wp=rw
|+- policy='round-robin 0' prio=50 status=active
| | 6:0:0:0 sde 8:64 active ready running
| | 3:0:0:0 sdb 8:16 active ready running
|+- policy='round-robin 0' prio=10 status=enabled
| | 4:0:0:0 sdc 8:32 active ready running
| | 5:0:0:0 sdd 8:48 active ready running
[root@rhell ~]# ls -l /dev/mapper
total 0
lrwxrwxrwx 1 root root 7 Dec 18 00:26 3600a0980774f6a36625d4d4f7775564d -> ../dm-3
crw----- 1 root root 10, 236 Dec 15 06:08 control
lrwxrwxrwx 1 root root 7 Dec 15 06:08 docker-253:1-998953-pool -> ../dm-2
lrwxrwxrwx 1 root root 7 Dec 15 06:08 rhel-root -> ../dm-1
lrwxrwxrwx 1 root root 7 Dec 15 06:08 rhel-swap -> ../dm-0
[root@rhell ~]# sanlun lun show -p

ONTAP Path: svmluns:/vol/linluns/linux.lun
LUN: 0
LUN Size: 10g
Product: cDOT
Host Device: 3600a0980774f6a36625d4d4f7775564d
Multipath Policy: round-robin 0
Multipath Provider: Native

-----
host      vserver
path      path
state     type      /dev/  host      vserver
          node   node   adapter  LIF
-----
up        primary  sde     host6     cluster1-01_iscsi_lif_2
up        primary  sdb     host3     cluster1-01_iscsi_lif_1
up        secondary sdc     host4     cluster1-02_iscsi_lif_2
up        secondary sdd     host5     cluster1-02_iscsi_lif_1
[root@rhell ~]#
```



**Tip:** You can see even more detail about the configuration of multipath and the LUN as a whole by issuing the `multipath -v3 -d -ll` or `iscsiadm -m session -P 3` commands. Because the output of these commands is rather lengthy, it is omitted here, but you are welcome to run these commands in your lab.

- The LUN is now fully configured for multipath access, so the only steps that remain before you can use the LUN is to create a filesystem and mount it.



**Important:** When you run the following commands in your lab you will need to substitute in the `/dev/mapper/...` string that identifies the specific LUN in your lab. You can get that string from the output of `ls -l /dev/mapper`. You can use bash tab completion when entering the multipath file name to save yourself some tedious typing.

```
[root@rhell ~]# mkfs.ext4 /dev/mapper/3600a0980774f6a36625d4d4f7775564d
mke2fs 1.42.9 (28-Dec-2013)
```

```

Discarding device blocks: done
Filesystem label=
OS type: Linux
Block size=4096 (log=2)
Fragment size=4096 (log=2)
Stride=0 blocks, Stripe width=16 blocks
655360 inodes, 2621440 blocks
131072 blocks (5.00%) reserved for the super user
First data block=0
Maximum filesystem blocks=2151677952
80 block groups
32768 blocks per group, 32768 fragments per group
8192 inodes per group
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736, 1605632

Allocating group tables: done
Writing inode tables: done
Creating journal (32768 blocks): done
Writing superblocks and filesystem accounting information: done

[root@rhell ~]# mkdir /linuxlun
[root@rhell ~]# mount -t ext4 -o discard /dev/mapper/3600a0980774f6a36625d4d4f7775564d /linuxlun
[root@rhell ~]# ls /linuxlun
lost+found
[root@rhell ~]# df
Filesystem                                1K-blocks    Used Available Use% Mounted on
/dev/mapper/rhel-root                    39265556 2811432  36454124   8% /
devtmpfs                                930956      0    930956   0% /dev
tmpfs                                   941996      0    941996   0% /dev/shm
tmpfs                                   941996    98884    843112  11% /run
tmpfs                                   941996      0    941996   0% /sys/fs/cgroup
/dev/sda1                             508588  129304    379284  26% /boot
tmpfs                                  188400      0    188400   0% /run/user/0
svml:/                                19456      96    18496   5% /svml
/dev/mapper/3600a0980774f6a36625d4d4f7775564d 10190100  36888   9612540  1% /linuxlun
[root@rhell ~]# ls /linuxlun
lost+found
[root@rhell ~]# echo "hello from rhell" > /linuxlun/test.txt
[root@rhell ~]# cat /linuxlun/test.txt
hello from rhell
[root@rhell iscsi]# ls -l /linuxlun/test.txt
-rw-r--r-- 1 root root 17 Apr  9 18:00 /linuxlun/test.txt
[root@rhell ~]#

```

The discard option for `mount` allows the Red Hat host to utilize space reclamation for the LUN.

12. To have RHEL automatically mount the LUN's filesystem at boot time, issue the following command (modified to reflect the multipath device path being used in your instance of the lab) to add the mount information to the `/etc/fstab` file. Issue the following command as a single line.

```

[root@rhell ~]# echo '/dev/mapper/3600a0980774f6a36625d4d4f7775564d
/linuxlun ext4 _netdev,discard,defaults 0 0' >> /etc/fstab
[root@rhell ~]#

```

## 4 References


The following references were used in writing this lab guide.

- TR-3982: [NetApp Clustered Data ONTAP 8.3.X and 8.2.X – an Introduction](#), November 2015
- TR-4100: [Nondisruptive Operations with SMB File Shares ONTAP 9.x](#), November 2016
- TR-4129: [Namespaces in clustered Data ONTAP](#), July 2014
- TR-4523: [DNS Load Balancing in ONTAP - Configuration and Best Practices](#), May 2020
- NetApp Hardware Universe: <https://hwu.netapp.com/Home/Index>, NetApp Support site login required.

## 5 Version History

Version	Date	Document Version History
Version 1.0	October 2014	Initial Release for Hands On Labs
Version 1.0.1	December 2014	Updates for Lab on Demand
Version 1.1	April 2015	Updated for Data ONTAP 8.3GA and other application software. NDO section spun out into a separate lab guide.
Version 1.2	October 2015	Updated for Data ONTAP 8.3.1GA and other application software.
Version 1.3	September 2016	Updated for ONTAP 9.0RC1 and other application software.
Version 1.3 Rev 1	November 2016	Updated for ONTAP 9.0P1, various errata.
Version 2.0	October 2018	Updated for ONTAP 9.4, created separate NAS and SAN lab guides.
Version 2.1	April 2019	Updated for ONTAP 9.5
Version 2.2	December 2020	Updated for ONTAP 9.7
Version 2.2.1	May 2023	Fixed typo

## Legal Notices

 **Note:** This document was created using machine translation. If you have any issues with this translated content please consult the English version.

Legal notices provide access to copyright statements, trademarks, patents, and more.

- **Copyright**

<http://www.netapp.com/us/legal/copyright.aspx>

- **Trademarks**

NETAPP, the NETAPP logo, and the marks listed on the NetApp Trademarks page are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.

<http://www.netapp.com/us/legal/netapptmlist.aspx>

- **Patents**

A current list of NetApp owned patents can be found at

<https://www.netapp.com/us/media/patents-page.pdf>

- **Privacy Policy**

<https://www.netapp.com/us/legal/privacypolicy/index.aspx>

- **Open Source**

Provides information about third-party copyright and licenses used in this product.

- [Notice for ONTAP 9.10.1](#)
- [Notice for ONTAP 9.9.1](#)
- [Notice for ONTAP 9.8](#)
- [Notice for ONTAP 9.7](#)
- [Notice for ONTAP 9.6](#)
- [Notice for ONTAP 9.5](#)
- [Notice for ONTAP 9.4](#)
- [Notice for ONTAP 9.3](#)
- [Notice for ONTAP 9.2](#)
- [Notice for ONTAP 9.1](#)



